

Schedule of the course(s)

- 2 Masters in parallel:
 - Master WET (Polytech Nantes), OSEmb (7h) and LogEmb (15h).
 - Master CORO (Ecole Centrale de Nantes), ESS (30h)
- Courses:
 - Embedded/Real-Time Operating Systems (CORO/ESS – WET/OsEmb)
 - Embedded Software (CORO/ESS – WET/LogEmb)
 - Additional courses, tutorials and labs for CORO/ESS

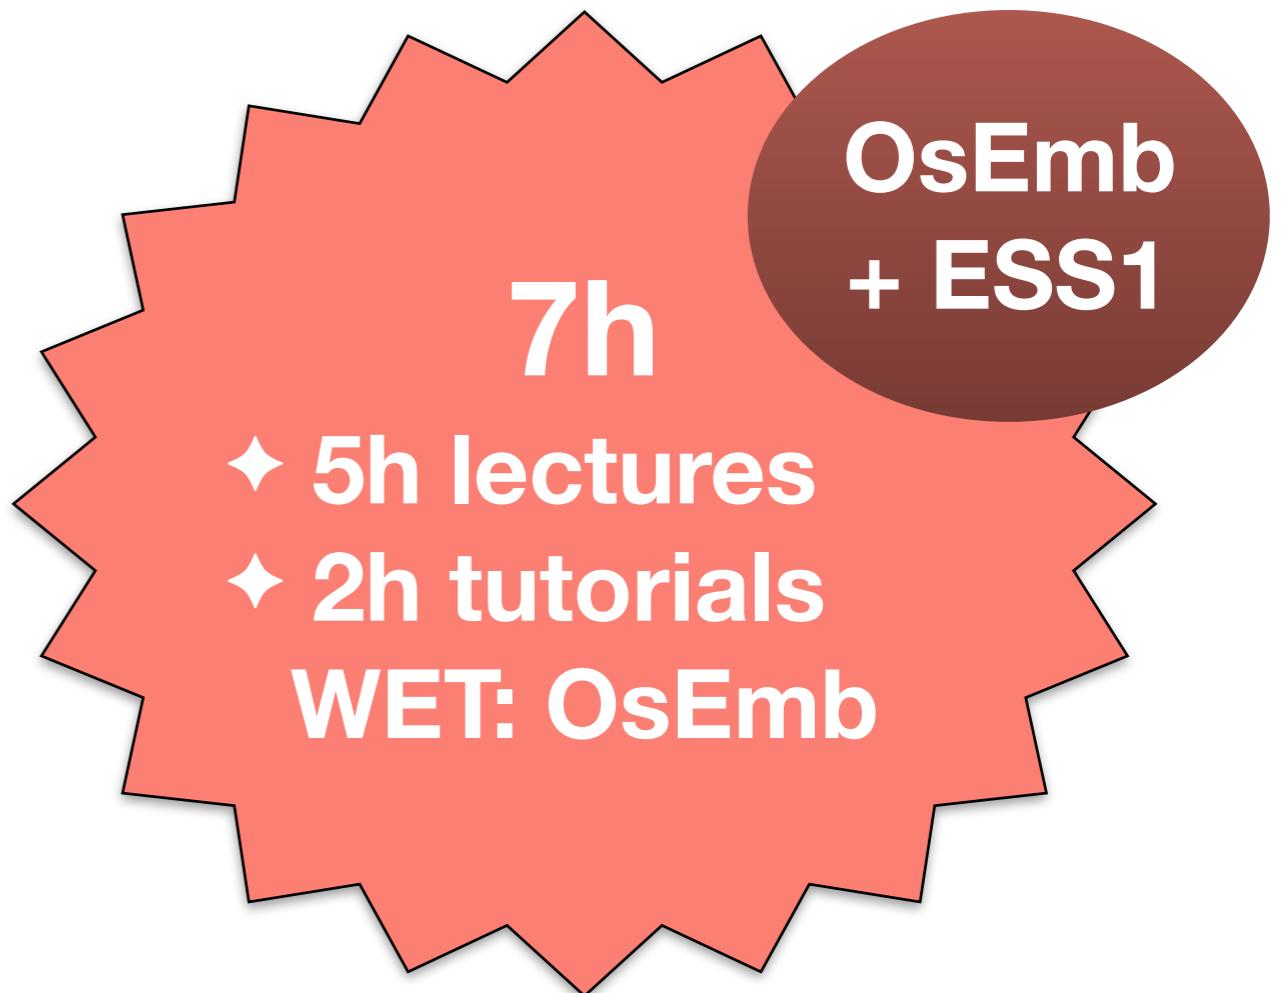
Documents : <http://wet.rts-software.org/>

Part 1

Teacher

Jean-Luc Béchennec

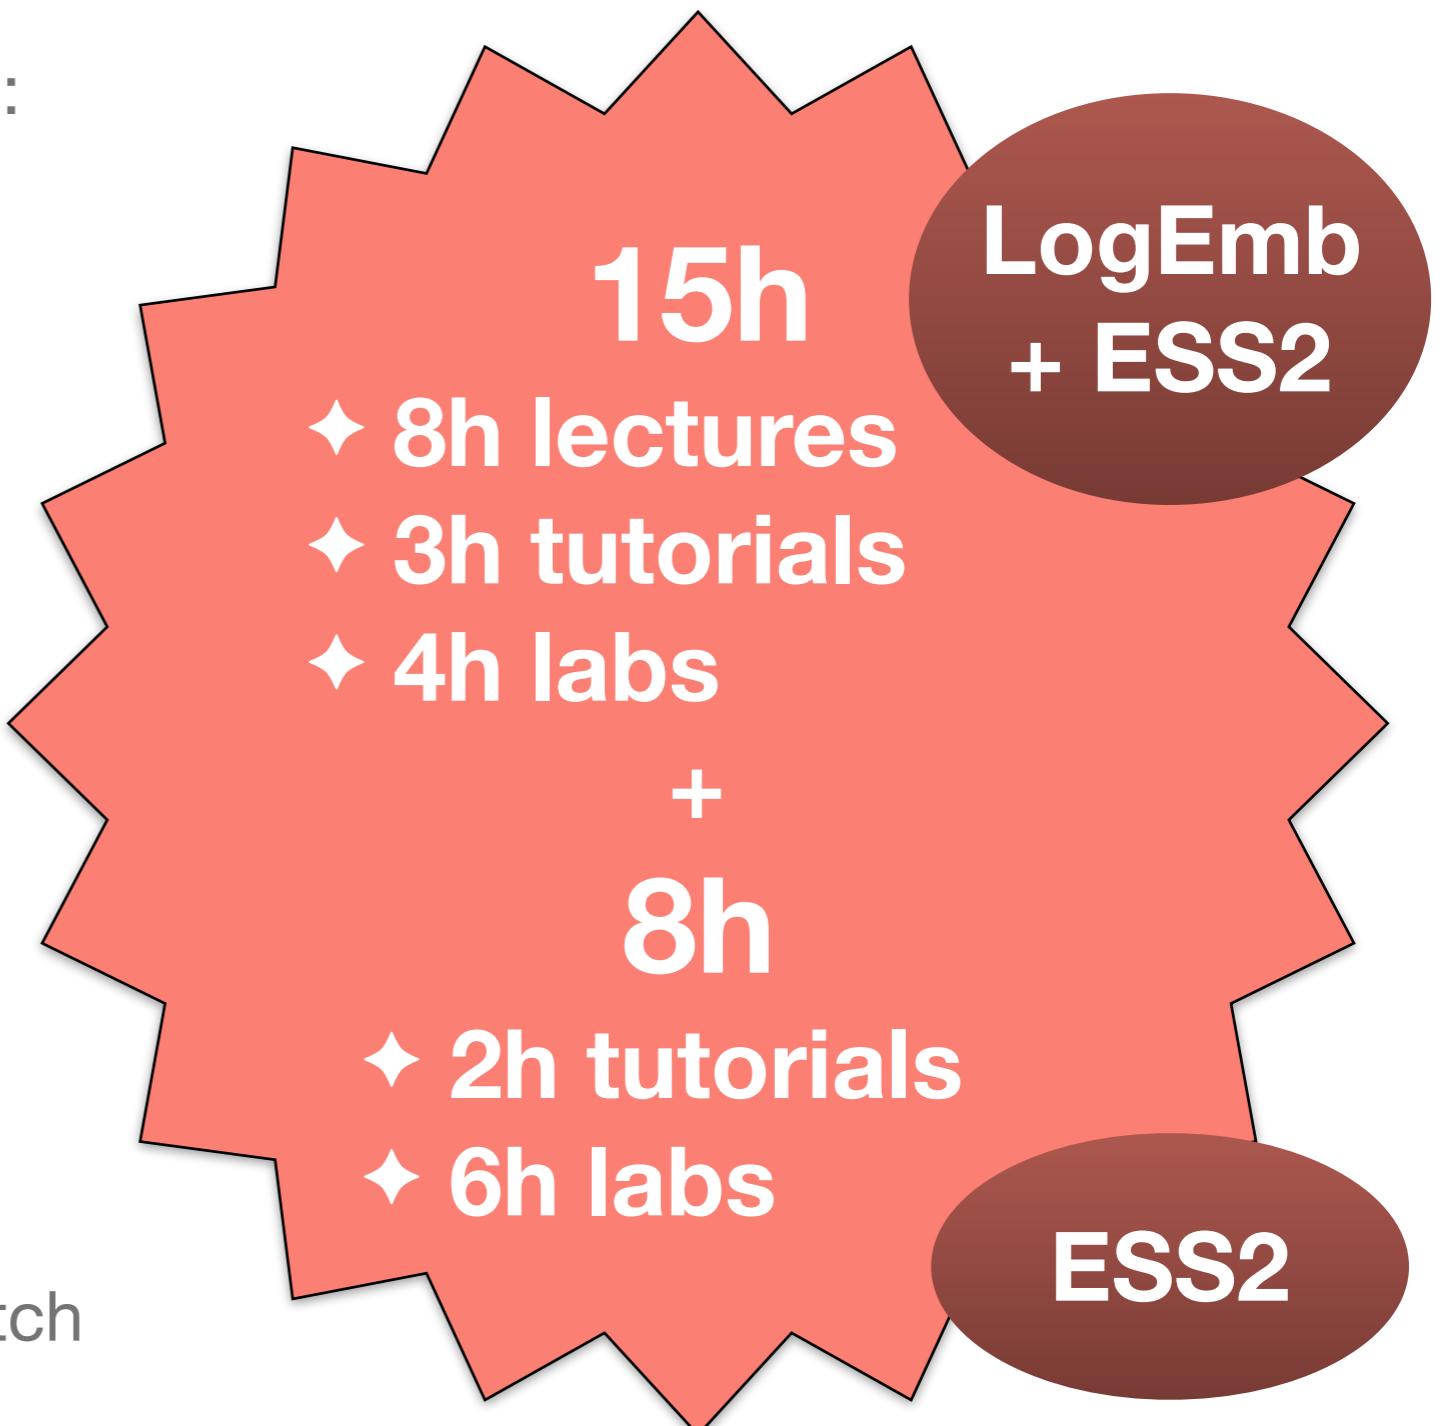
- Multiplexing task on a platform
 - specificities of embedded (RT)OS
 - interrupts
 - sequencer
 - operating system
- Scheduling
 - RM and EDF
- Synchronization
 - Hardware and data sharing: Mutex, semaphores, PCP
- Communication
 - Messaging systems vs shared variables



Part 2

Teachers **Jean-Luc Béchennec**
Sébastien Faucou
Mikaël Briday

- The C language is not your friend:
 - pitfalls
 - programming rules
- Good practices
 - Development process
 - Requirements / Test
- Application design
- Models and model checking
 - Verifying design
- C / Assembly / ABI
- Syscall / µP modes / context switch



Embedded Operating Systems / ESS1

Master CORO / WET

Jean-Luc Béchennec - jean-luc.bechennec@ls2n.fr

Generalities

Computer systems

- Hardware
 - Processor / Microprocessor / Microcontroller
 - Devices to interface with the real world
 - Data acquisition (sensors connected to input devices)
 - Action (actuators connected to output devices)
- Software
 - One or many applications
 - « Runs » on the hardware
 - Get data
 - Perform computation
 - Set action

Embedded systems

- An embedded system is:
 - A computing system doing a **dedicated function** or a set of dedicated function **within an electrical/mechanical system** to control (a part of) the electrical/mechanical system.
 - Low power consumption
 - Low cost (usually)
 - Size (not always)
 - Often **real-time**
- Can be discussed but:
 - does not encompass mobile devices (smartphones, tablets) which are very close to general purpose handheld computers.

Real-time systems

- A real-time system is a computing system that must produce a response within a specified (and known) time.
 - Implies a deadline.
 - Response after the deadline is a failure.
- Usually the input is an event or a data in the physical world.
 - Output is related to that event or data.
 - Output has to be produced quick enough to be compatible with the controlled process.
- Correctness of the system depends on :
 - The functional correctness (as in any system)
 - The **timeliness** : time critical systems
- Periodic systems

Desired properties

- **Predictability** (adapted from Stankovic)
 - A system is predictable if it is possible to show/prove/demonstrate that requirements are met subject to any assumption made, for example, concerning failures and workloads
 - A RTS must be **predictable** both in term of logical results and **timing behaviour**
- **Dependability** (adapted from Avizienis *et al.*)
 - A system is dependable if it has the ability to avoid service failures that are more frequent and more severe than is acceptable.

There are different classes of RTS

- **Hard real-time**
 - **No timing failure allowed** — a missing reaction could have catastrophic consequences on the environment (human life losses, environmental pollution, ...)
- **Soft real-time**
 - Occasional failure allowed — the environment is inherently robust to occasional late reactions.
- **Firm real-time**
 - Same as soft real-time but a late result has no value

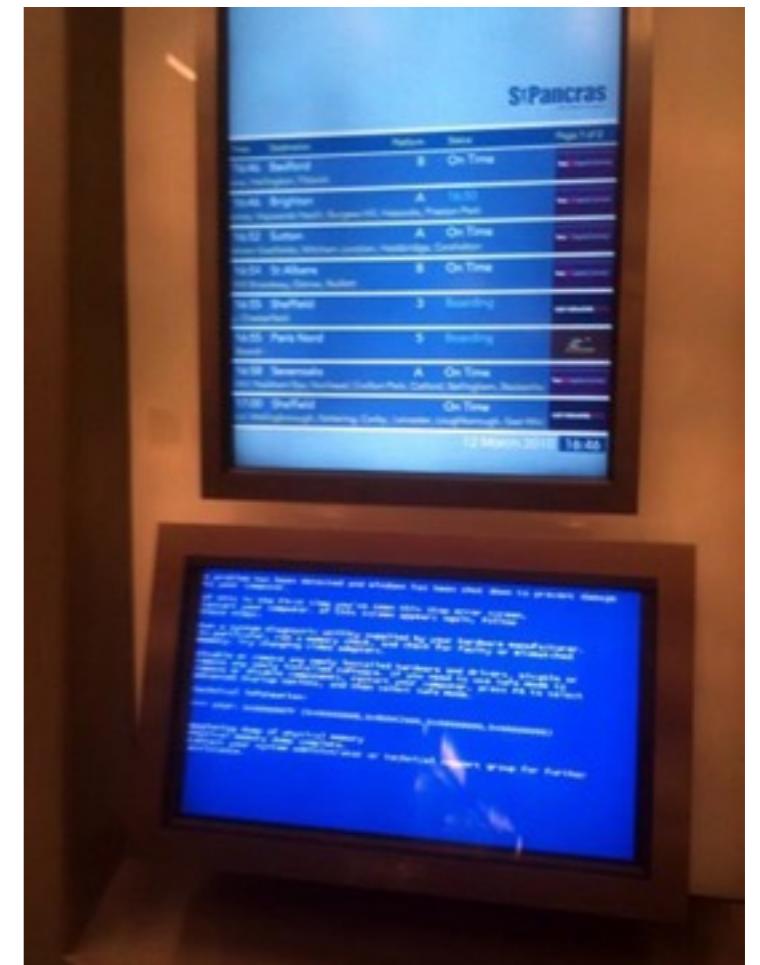
A single system can include hard real-time, firm real-time and soft real-time subsystems.

Summary

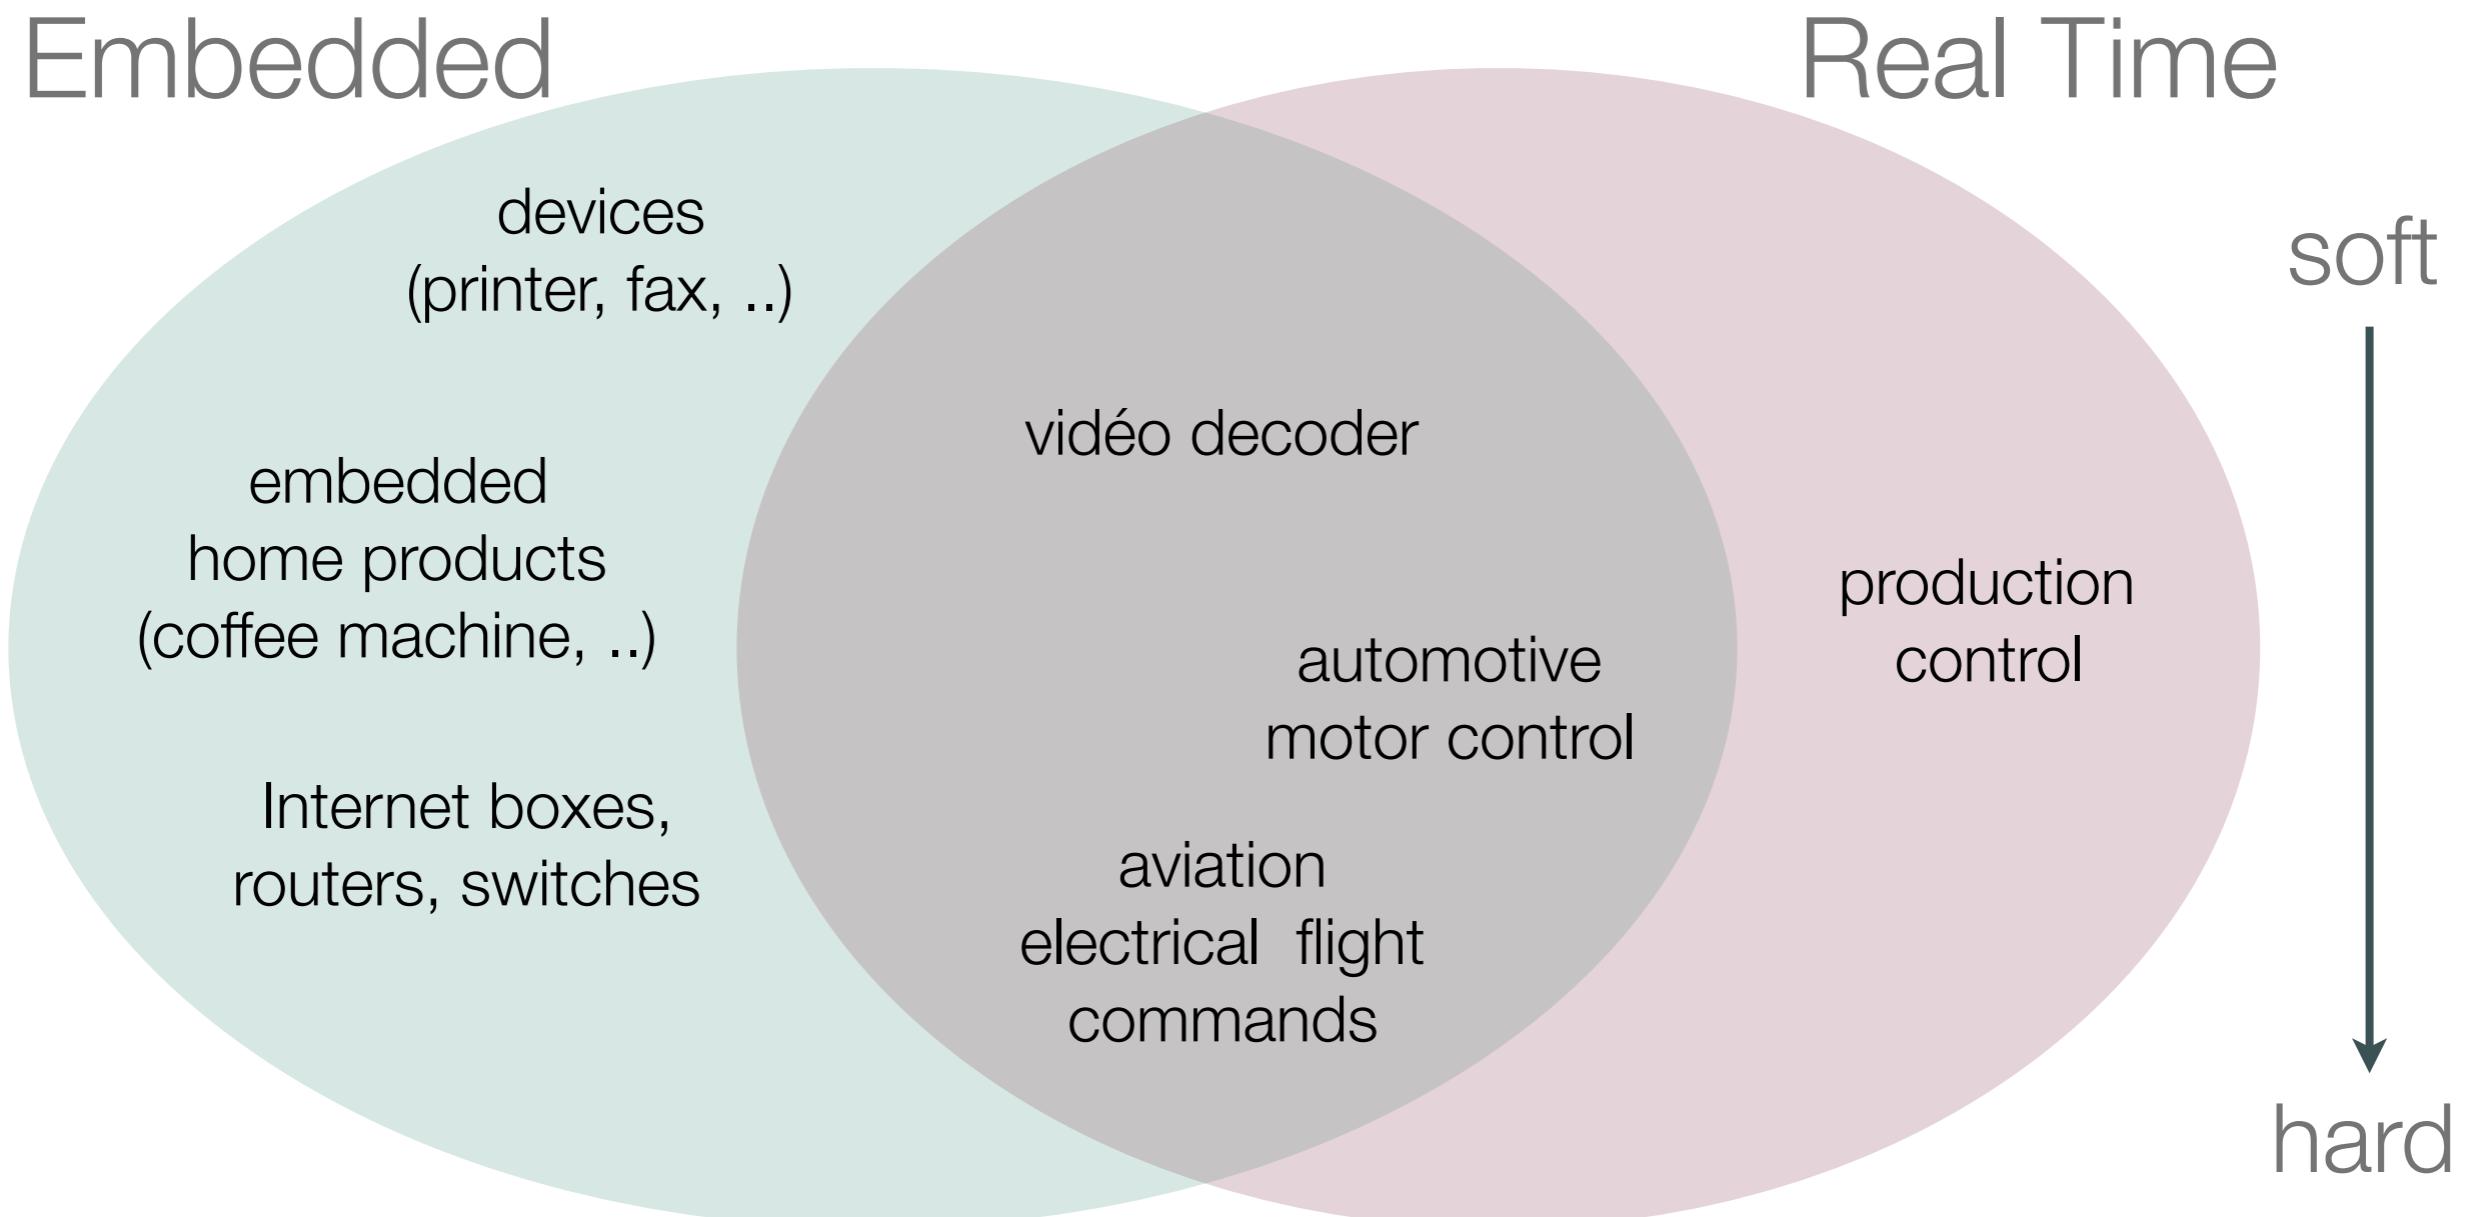


Real-time is not about being fast, it is about being on time. Be the tortoise, not the hare.

Real-time system design requires specific model, method, tools, and components.
Blue screen of death is not an option when commanding a plane or a car.

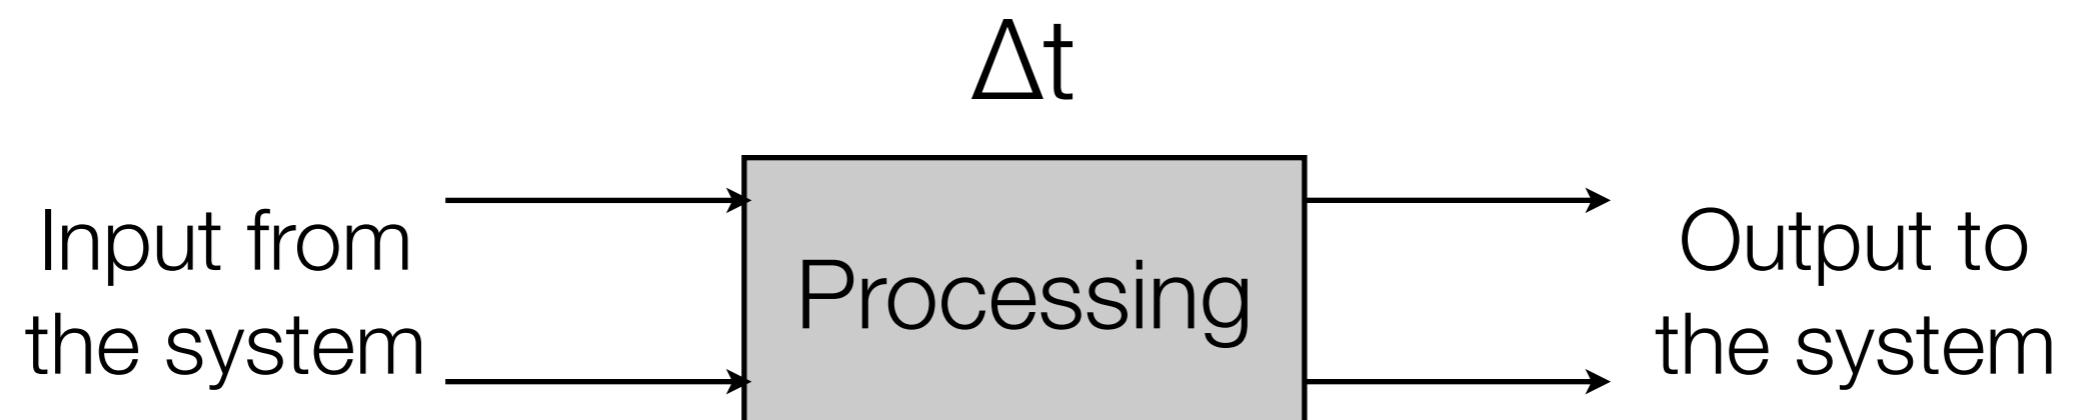


Embedded systems vs Real Time systems



Embedded / Real-time systems

- Simple system :
 - One task only

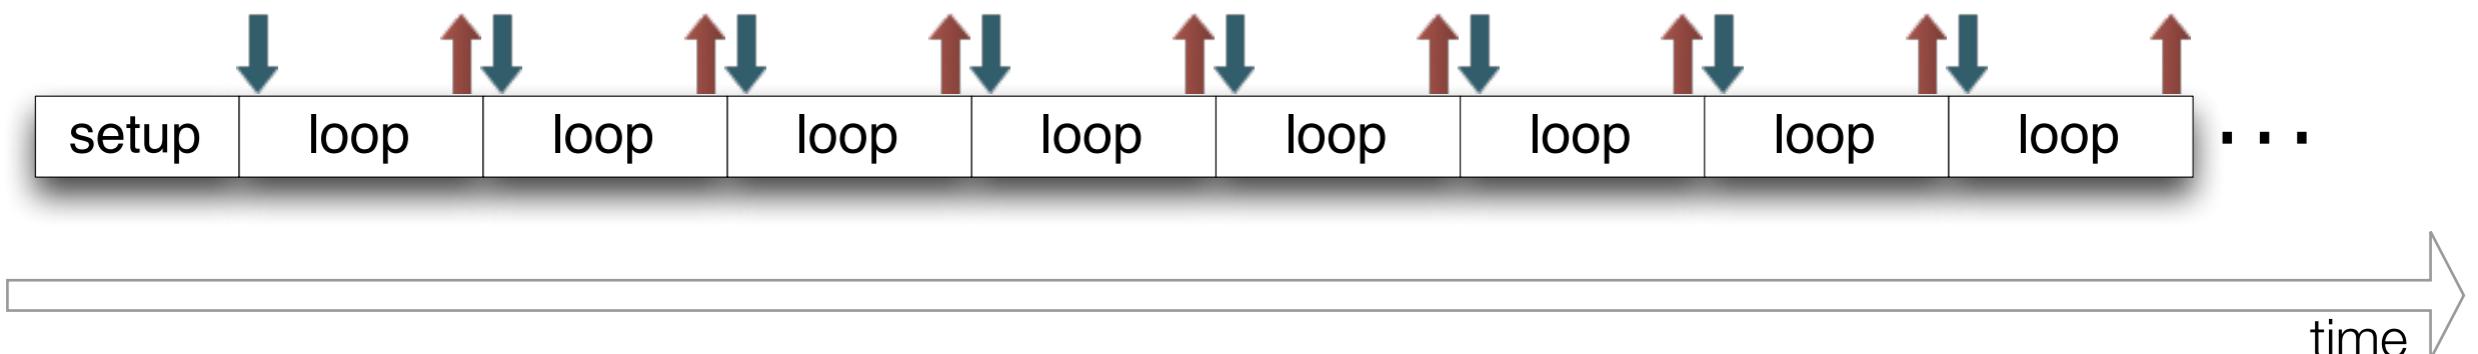


Example

```
void setup()
{
}

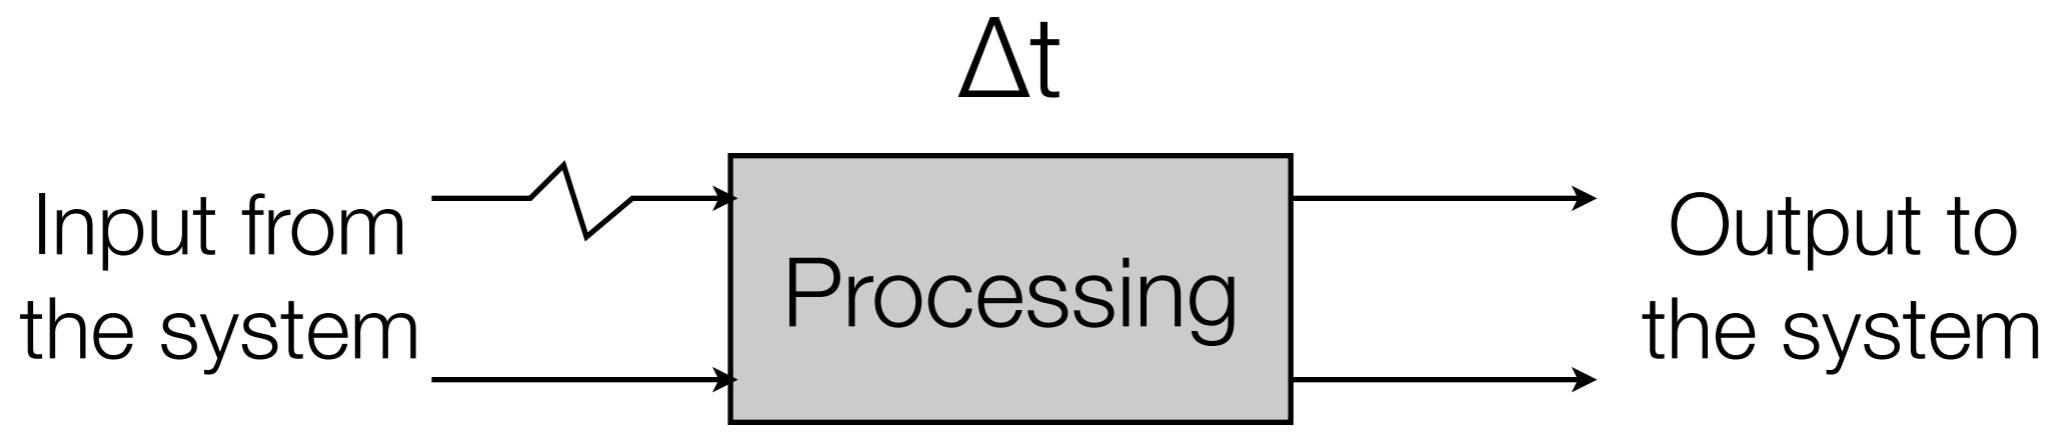
void loop()
{
    /* get the potentiometer value */
    uint16_t value = analogRead(A0); ←

    /* dim the LED */
    analogWrite(value/4); →
}
```



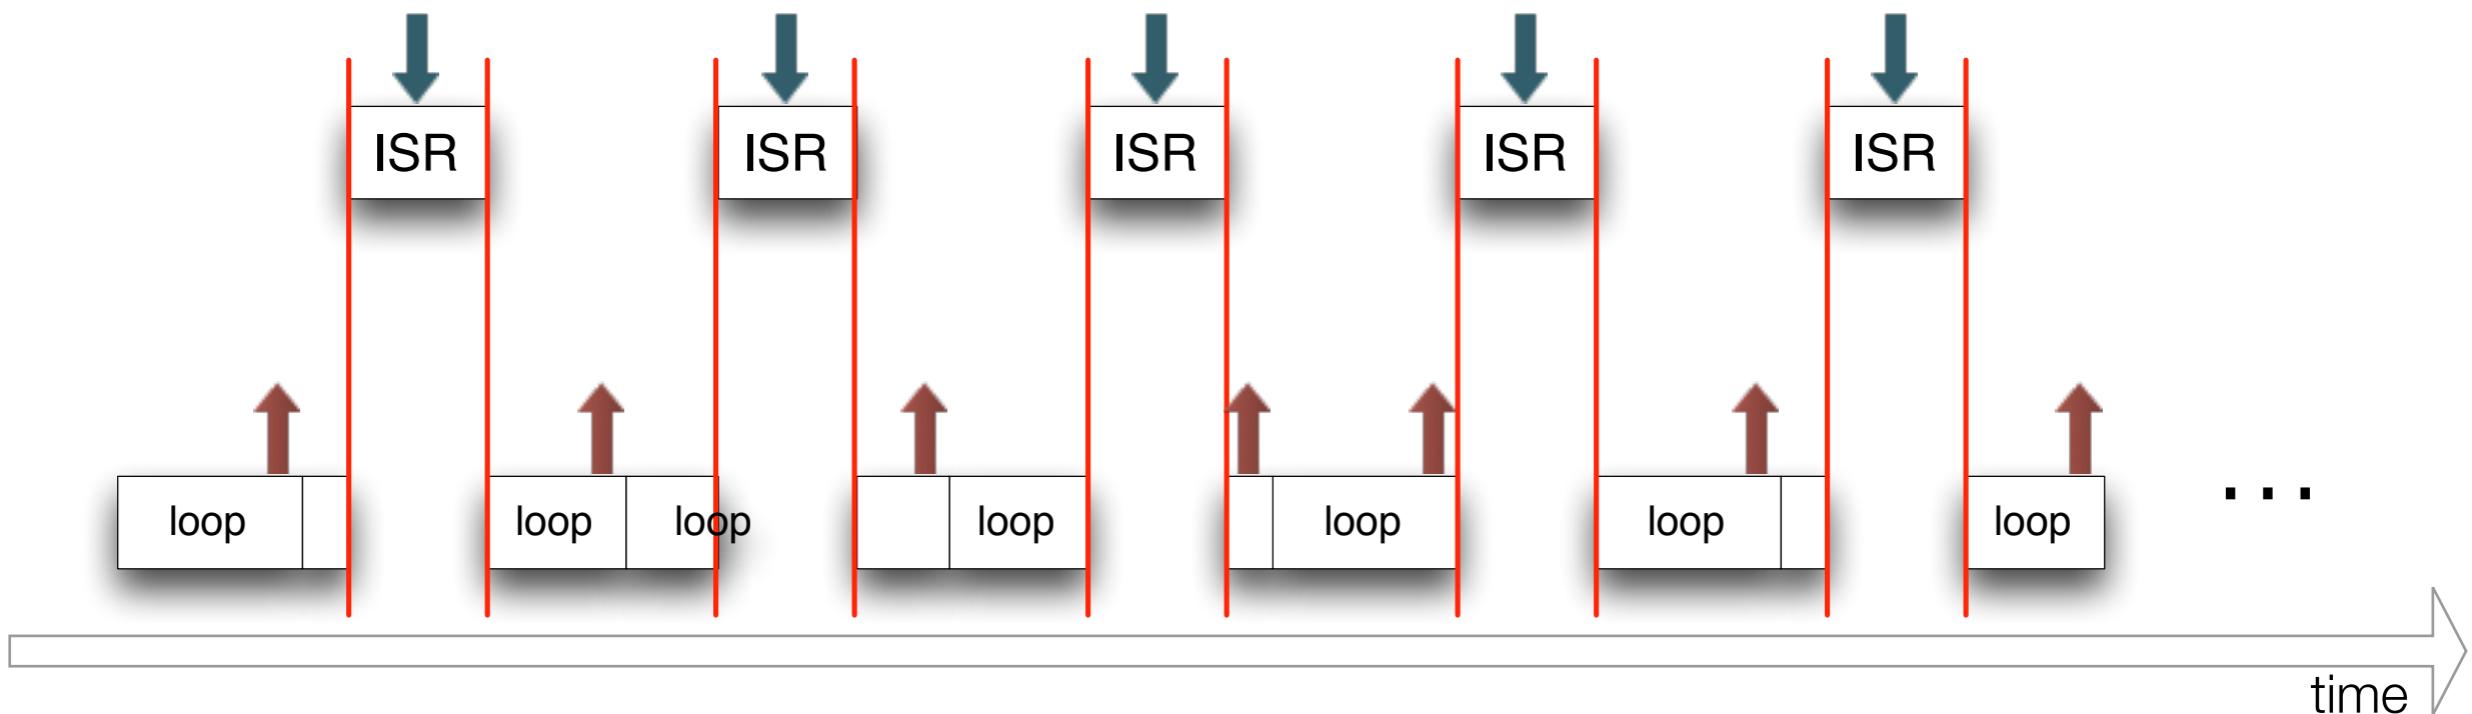
Embedded / Real-time systems

- Simple system :
 - Asynchronous



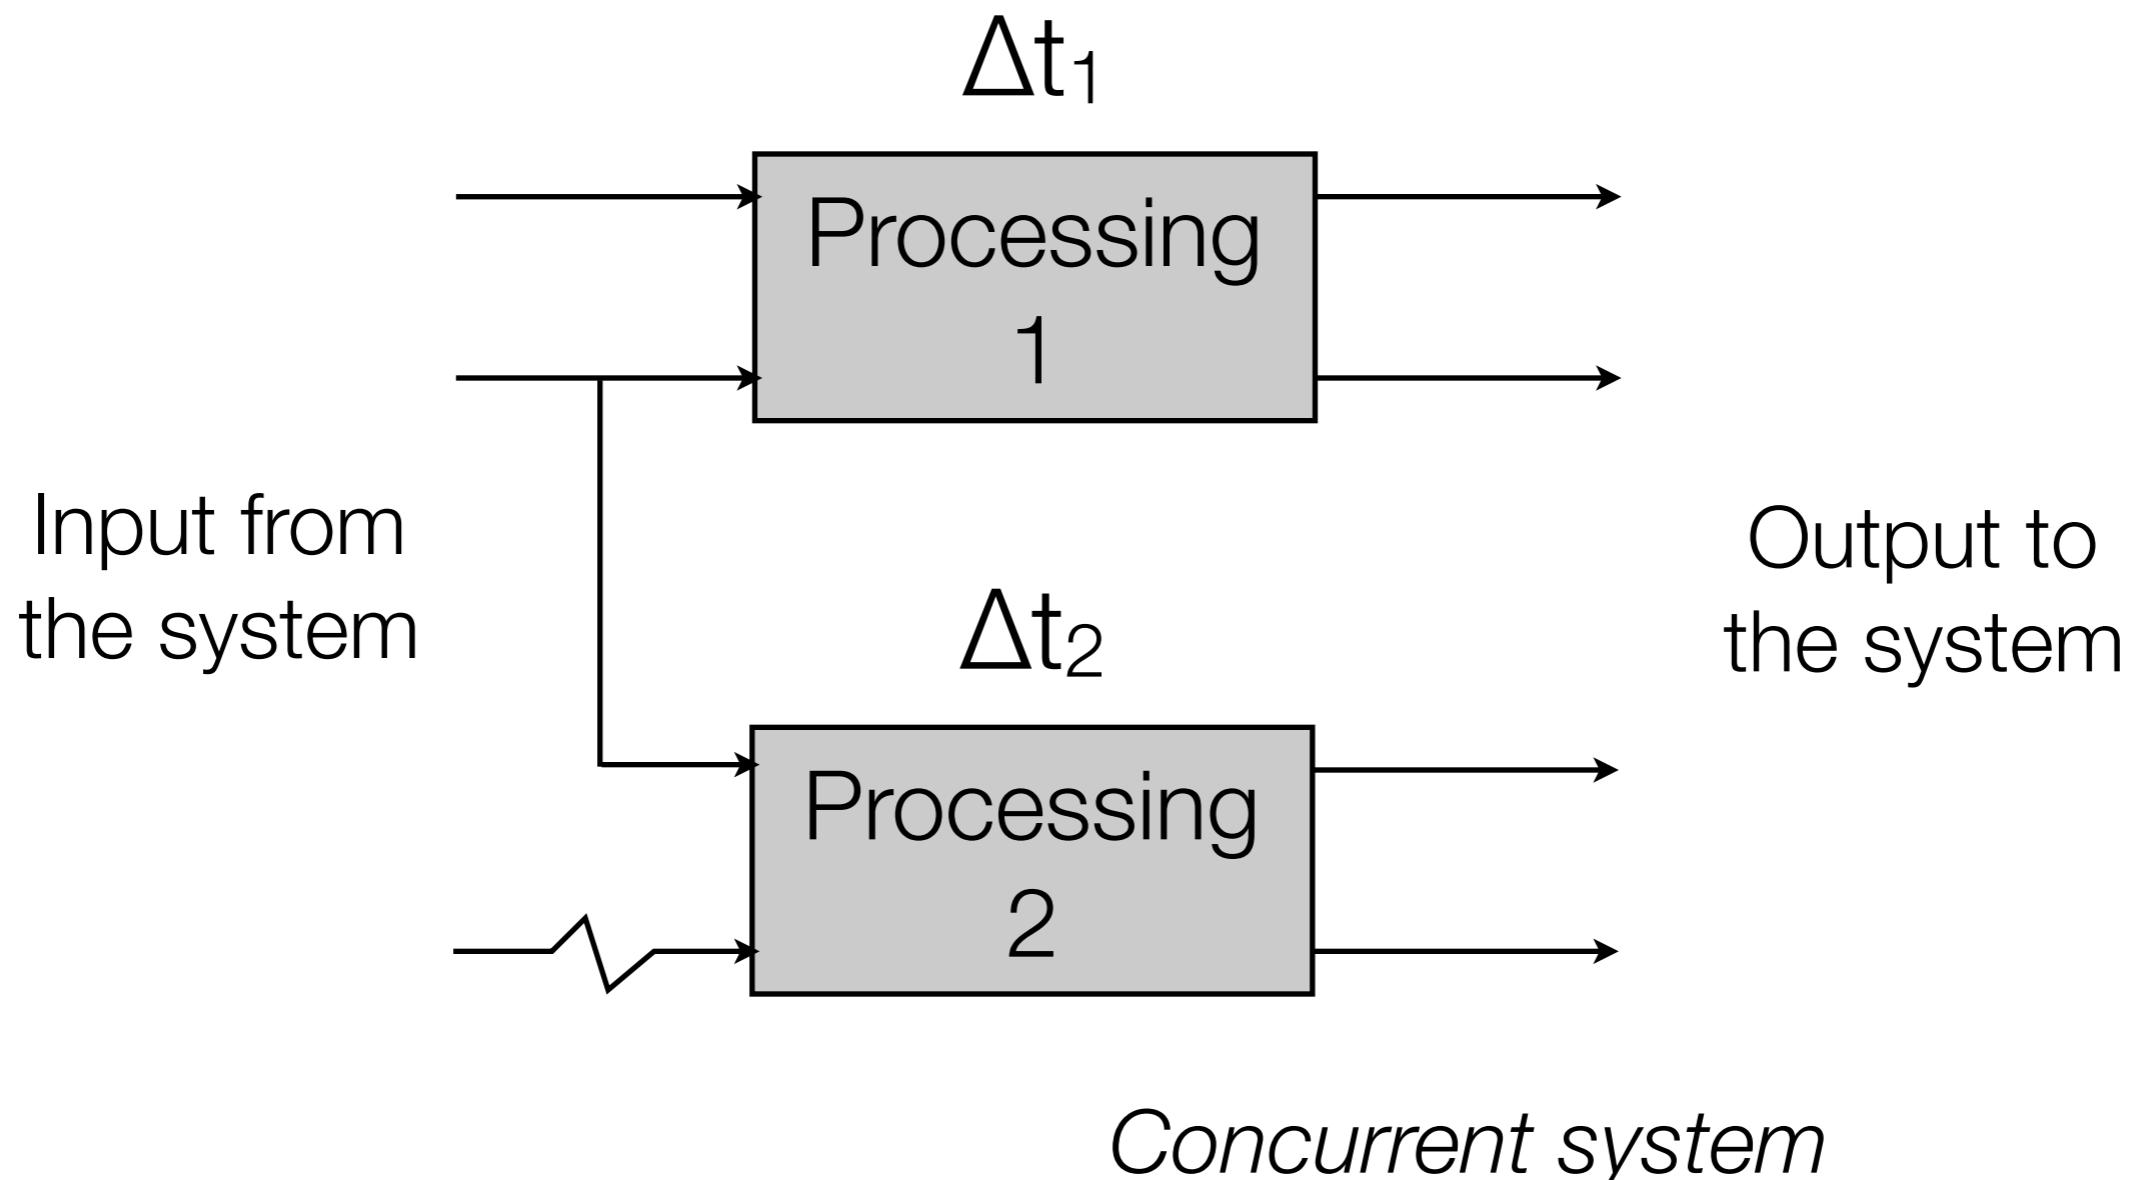
Example

```
volatile uint16_t value;  
  
ISR(ADC_vect)  
{  
    value = analogRead(A0); ←  
}  
  
void loop()  
{  
    analogWrite(value/4); →  
}
```



Embedded / Real-time systems

- More complex :
 - Multitask



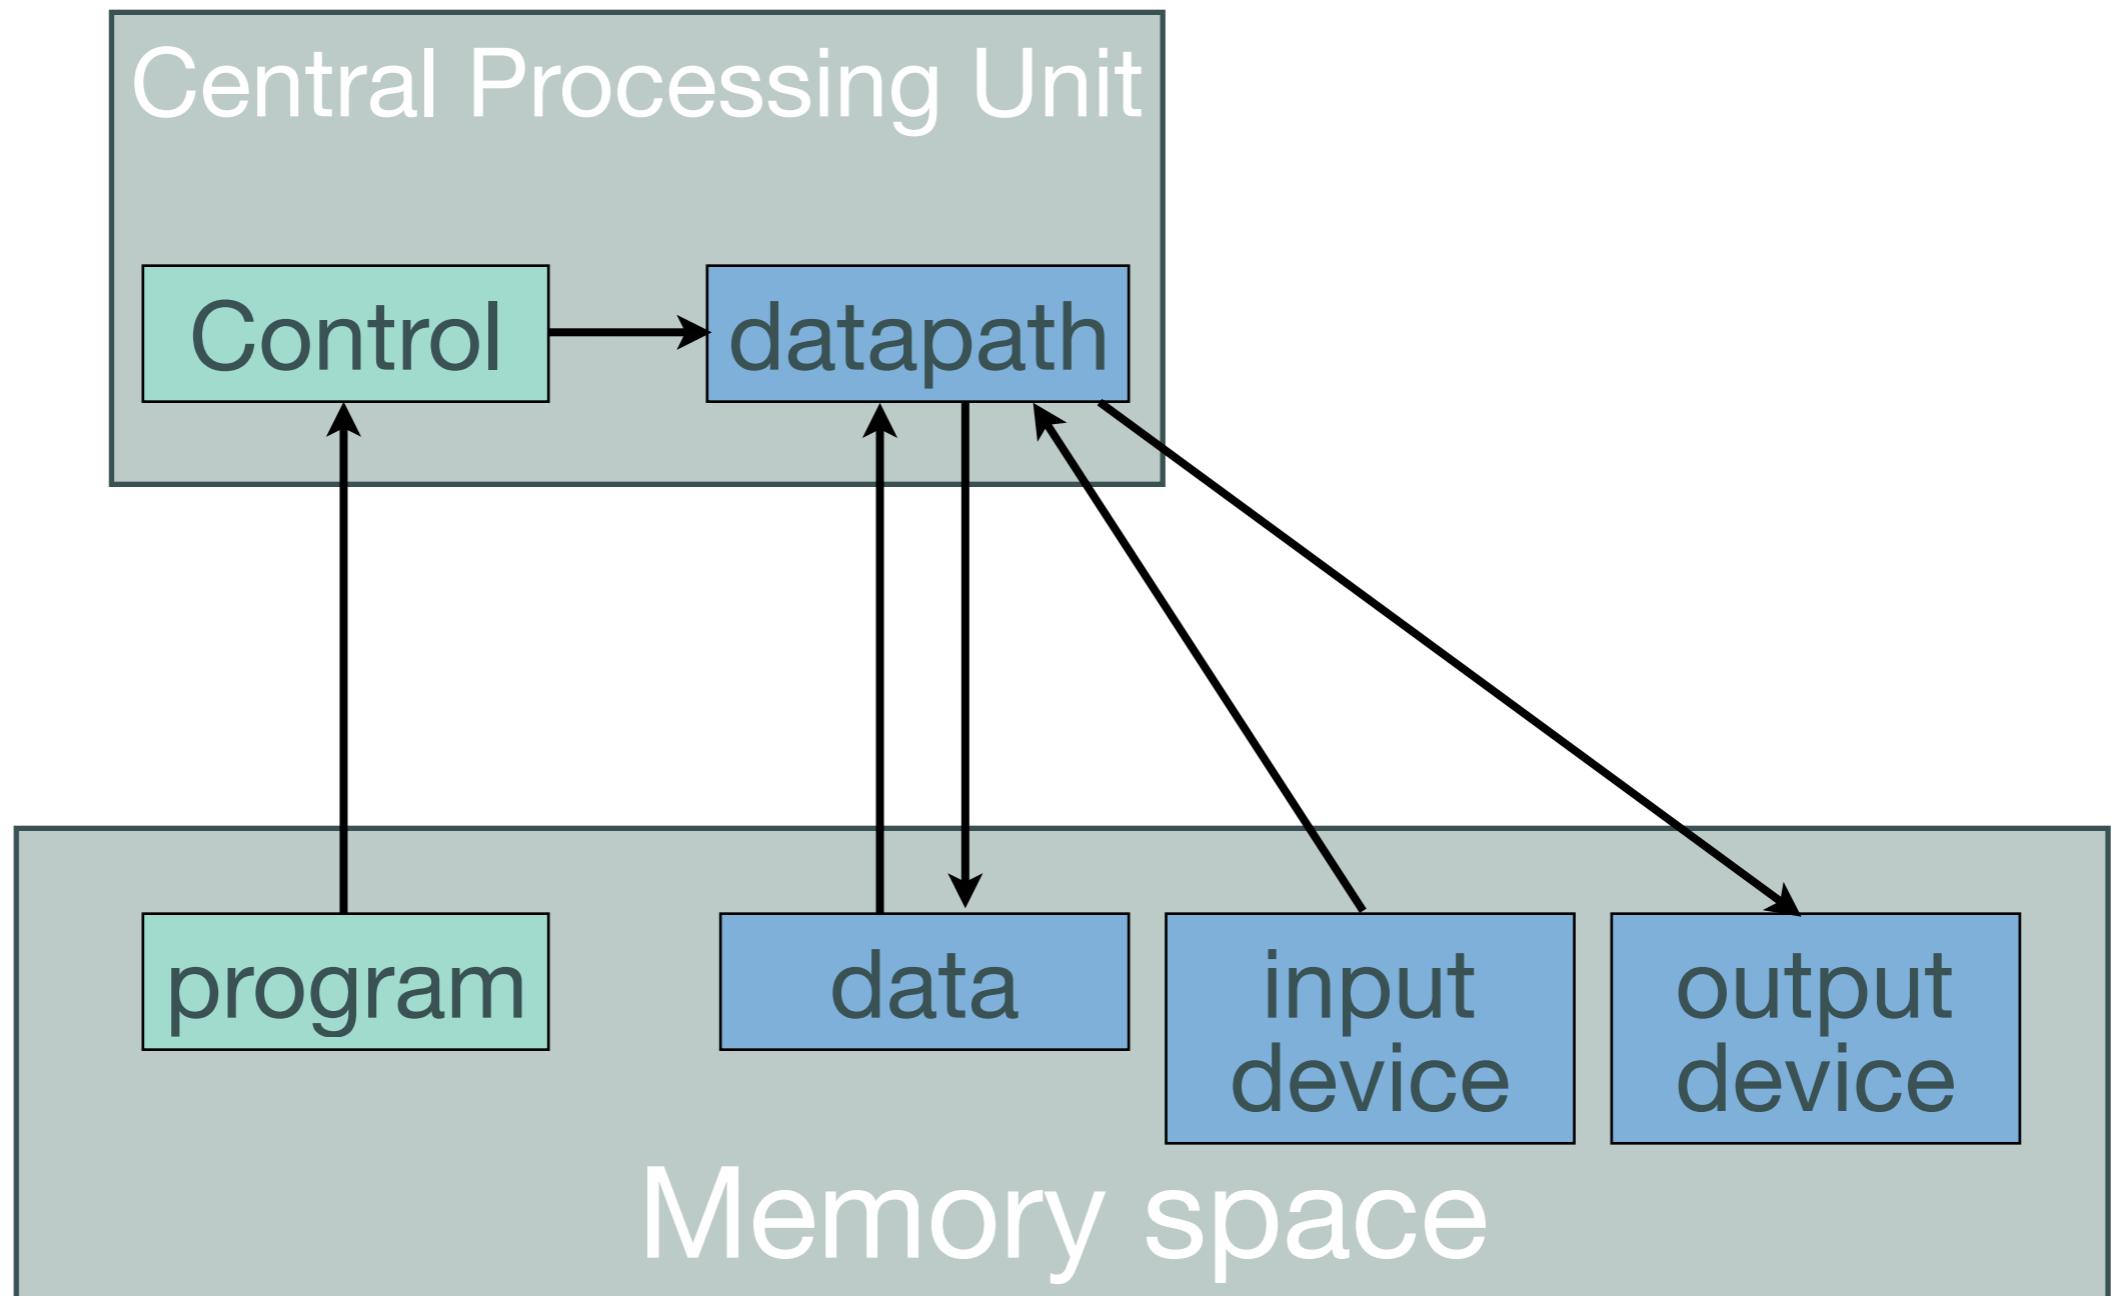
Software architecture

- When complexity of an application grows, it is necessary to **split the computing into separate tasks**
 - maintainability,
 - readability,
 - separation of concerns
 - programming abstraction

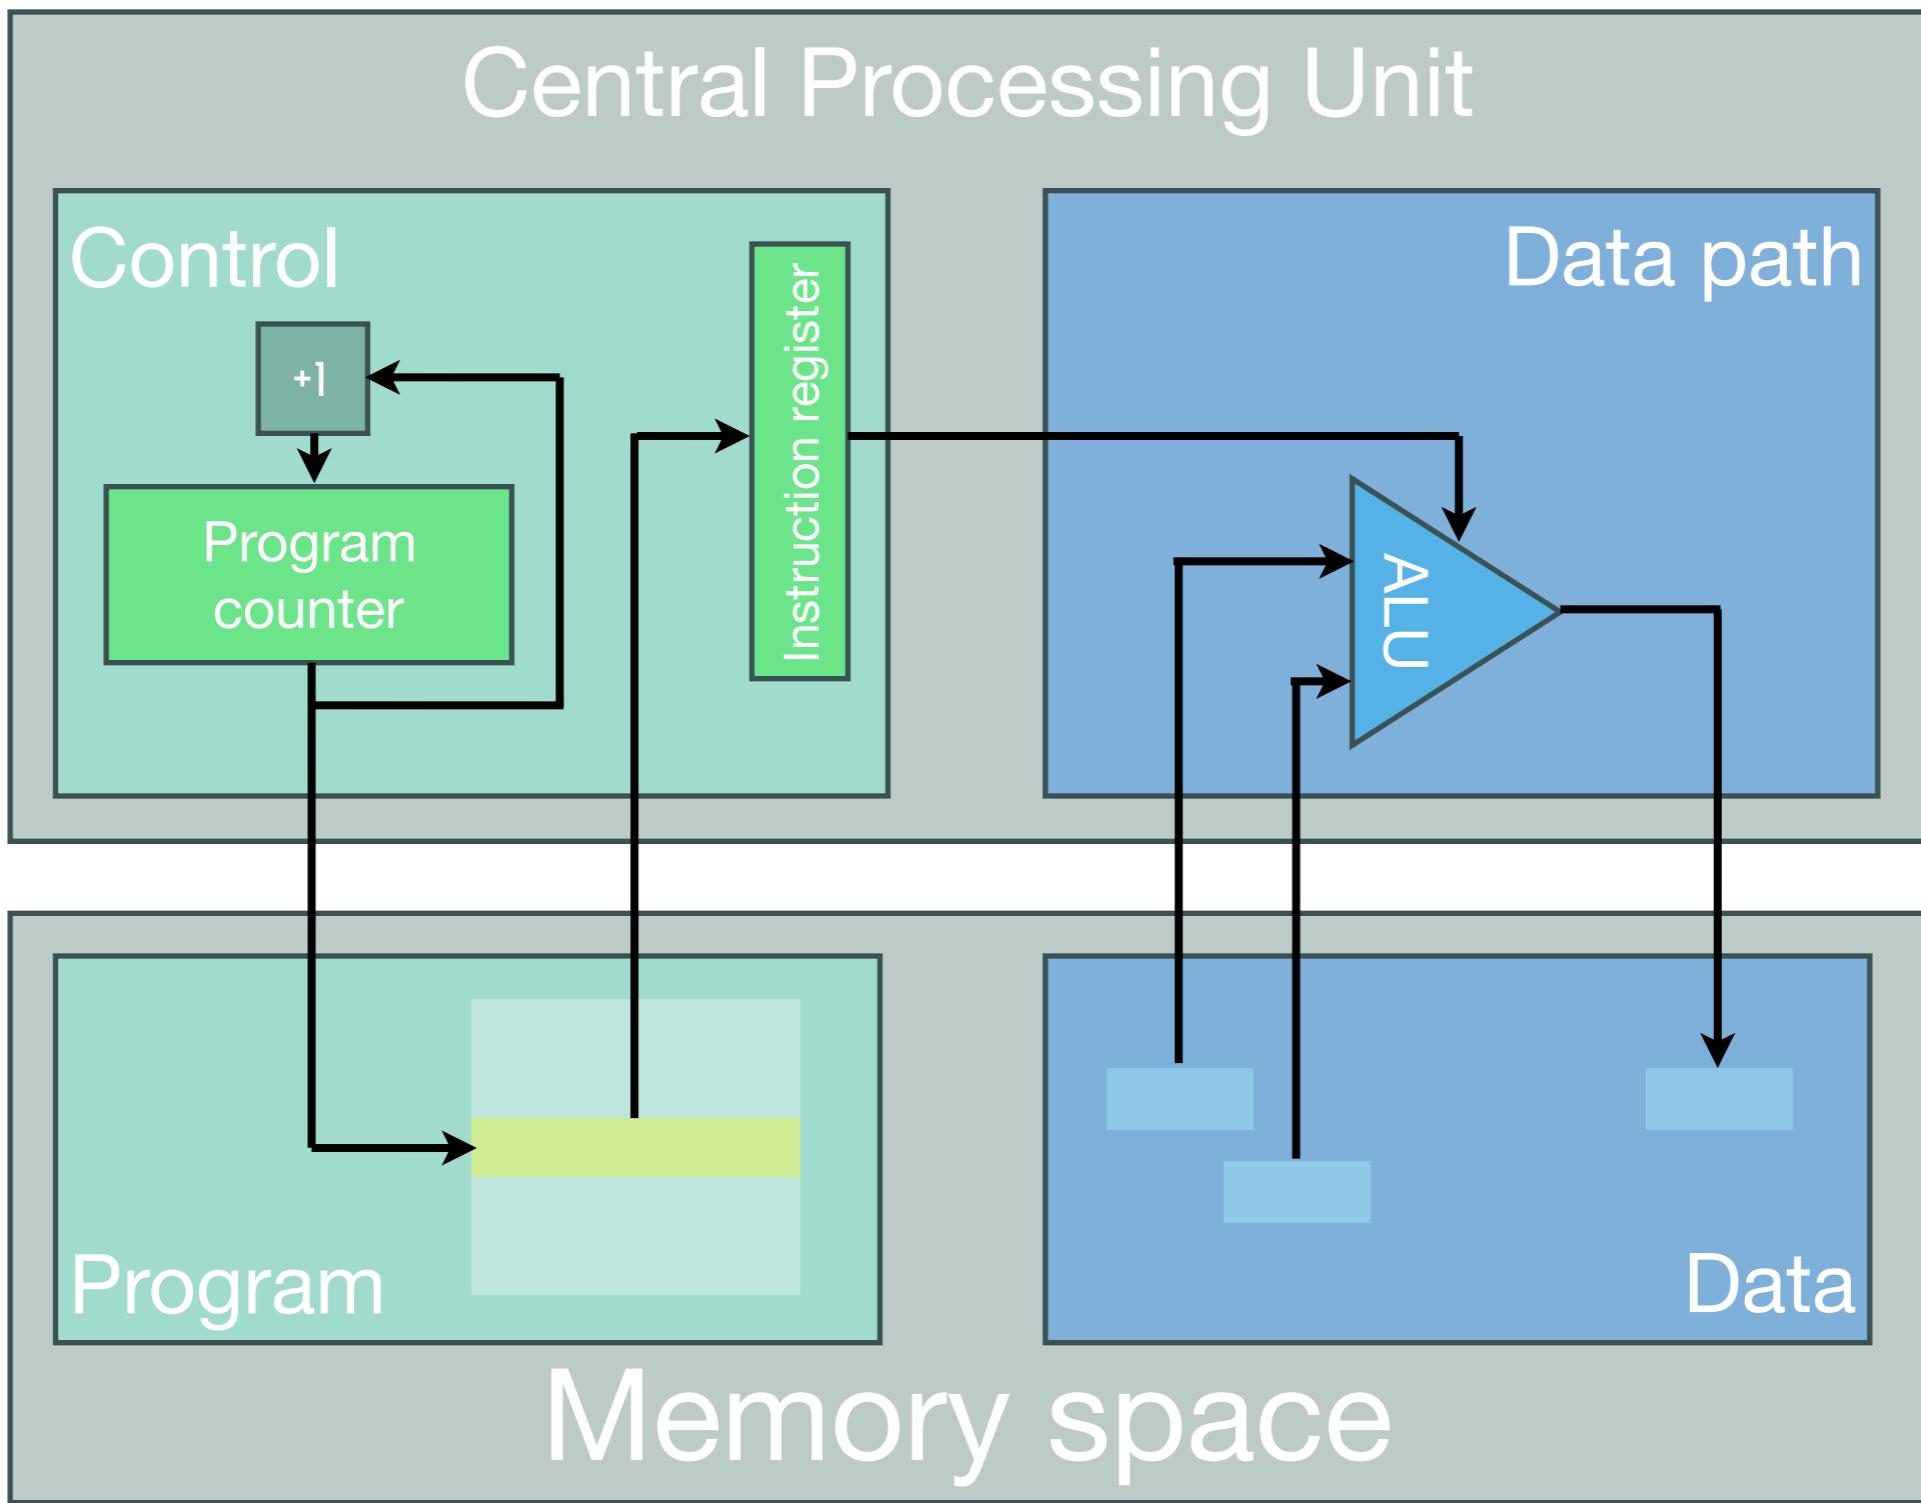
What means concurrent ?

- A computer is a sequential machine
 - A program is sequential (most of the time)
 - A program is a sequence of instructions
 - A computer executes the instructions one after the other
- Processing many tasks on the same computer means tasks are *concurrent* for their execution

Very simplified drawing



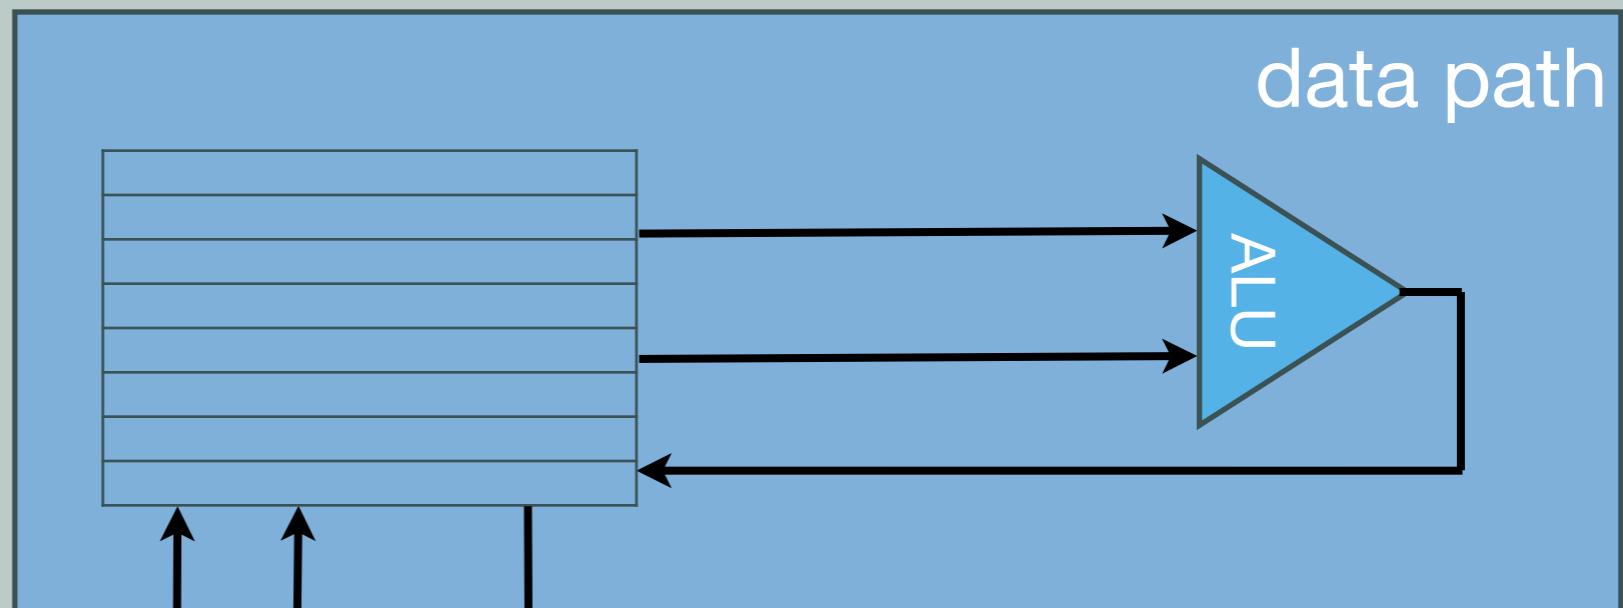
A little bit less simple



Focus on registers

Central processing unit

Registers



data

Memory

Example of the ARM Cortex (1)

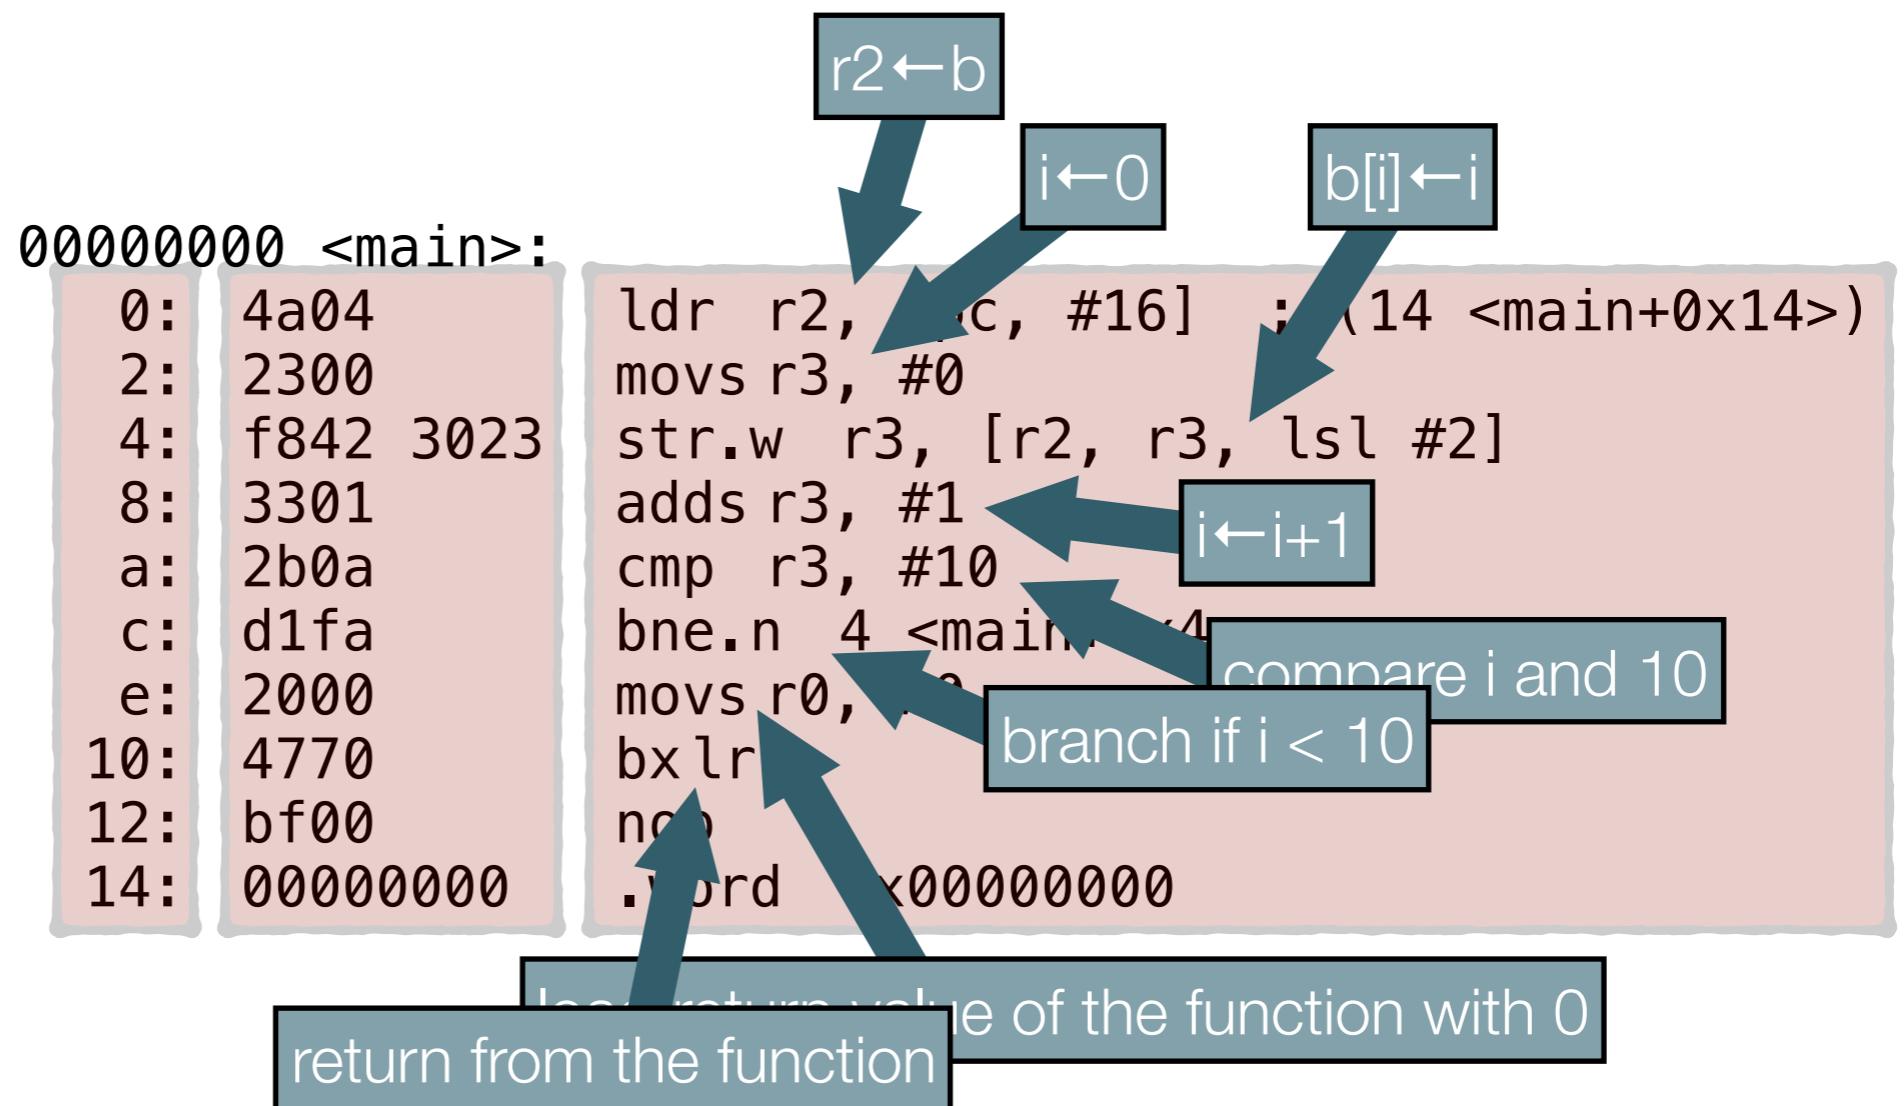
- The CPU uses 16 32 bits general purpose registers: **r0** to **r15**.
- Among them, 3 registers have a special meaning:
 - **r13** is the stack pointer, **sp**
 - **r14** is the link register, **lr**
 - **r15** is the program counter, **pc**
- At last, the PSR stores some flags, among others:
 - C, the Carry of an arithmetic/logic operation
 - N, Negative output of an arithmetic/logic operation
 - V, oVerflow in of an arithmetic/logic operation
 - Z, Zero output of an arithmetic/logic operation

This defines the context of execution

r0
r1
r2
r3
r4
r5
r6
r7
r8
r9
r10
r11
r12
sp
lr
pc
psr

A bit of assembly language

```
int b[10];  
  
int main()  
{  
    int i;  
  
    for (i = 0; i < 10; i++)  
    {  
        b[i] = i;  
    }  
  
    return 0;  
}
```



How to manage concurrency ?

- Bare metal approach
 - Program runs directly on the hardware
 - Asynchronous events are managed by interrupt handlers
 - Concurrency between interrupt handlers and normal execution
 - Limited to small and less complex systems.

How to manage concurrency ?

- Sequencer approach
 - The sequencer is a sequential program running in loop
 - The sequencer measures the time and calls functions (tasks) according to their activation period
 - Periodic system, predictable, deterministic
 - Not very flexible, all tasks are mixed, non preemptable.

Example of a sequencer

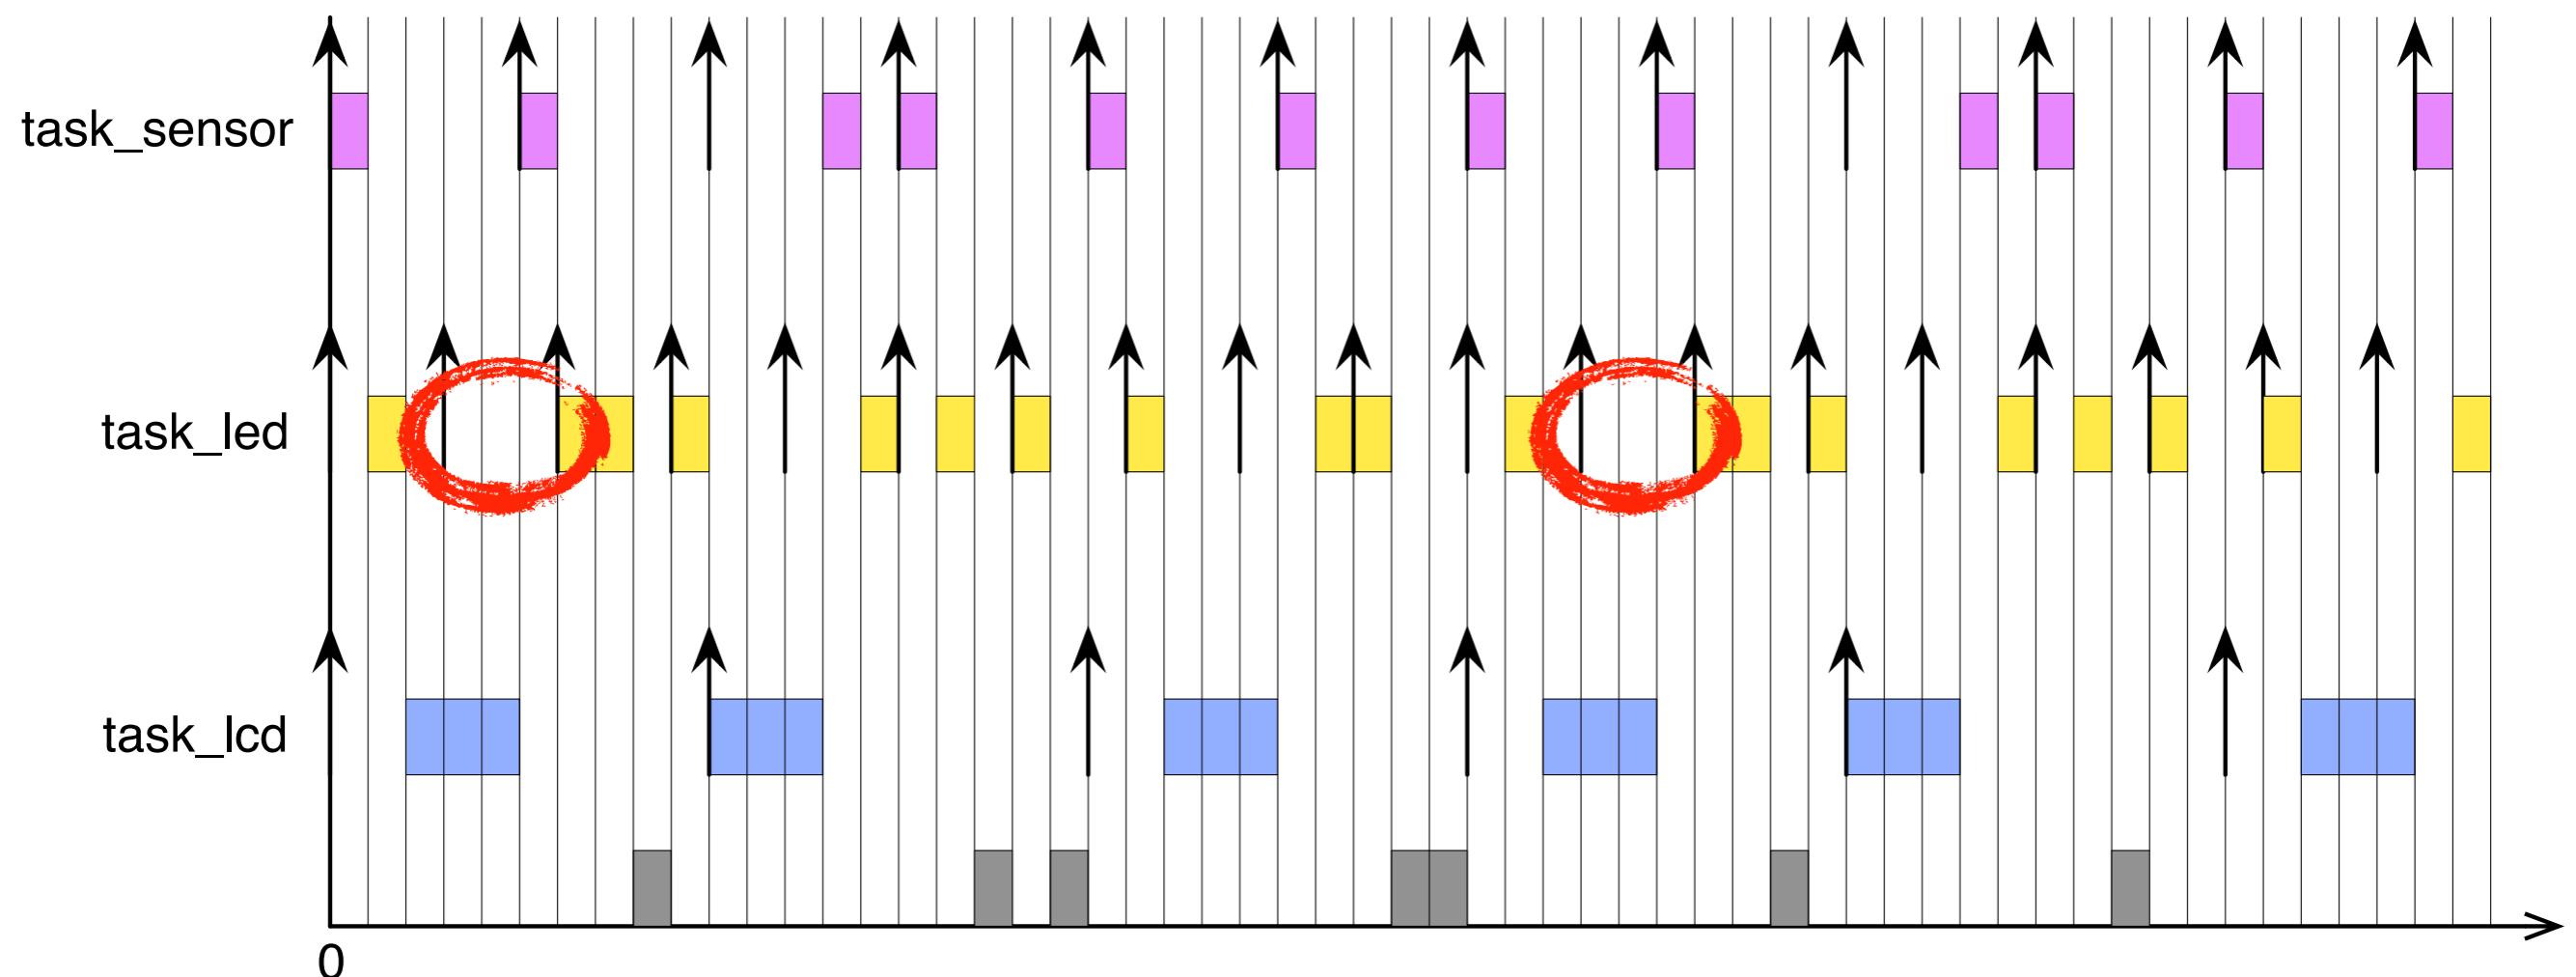
- Starting with previous analogRead / analogWrite example, we add a LCD.
 - Let's say we have 3 tasks each in a function:
 - task_sensor that does the analogRead to get the potentiometer value (execution time 1ms)
 - task_led that does the analogWrite to dim the led (execution time 1ms)
 - task_lcd that display the current value on a lcd (execution time 3ms)
 - First we decide of the period of each task:
 - task_sensor every 5ms
 - task_led every 3ms
 - task_lcd every 10ms

Example of a sequencer

```
uint16_t value;  
  
void task_sensor()  
{  
    value = analogRead(A0);  
}  
  
void task_led()  
{  
    analogWrite(value/4);  
}  
  
void task_lcd()  
{  
    lcd.write(value);  
}  
  
void loop()  
{  
    static uint32_t sensorTime = 0;  
    static uint32_t ledTime = 0;  
    static uint32_t lcdTime = 0;  
  
    uint32_t currentTime = millis();  
    if (currentTime >= sensorTime) {  
        sensorTime += 5;  
        task_sensor();  
    }  
    currentTime = millis();  
    if (currentTime >= ledTime) {  
        ledTime += 3;  
        task_led();  
    }  
    currentTime = millis();  
    if (currentTime >= lcdTime) {  
        lcdTime += 10;  
        task_lcd();  
    }  
}
```

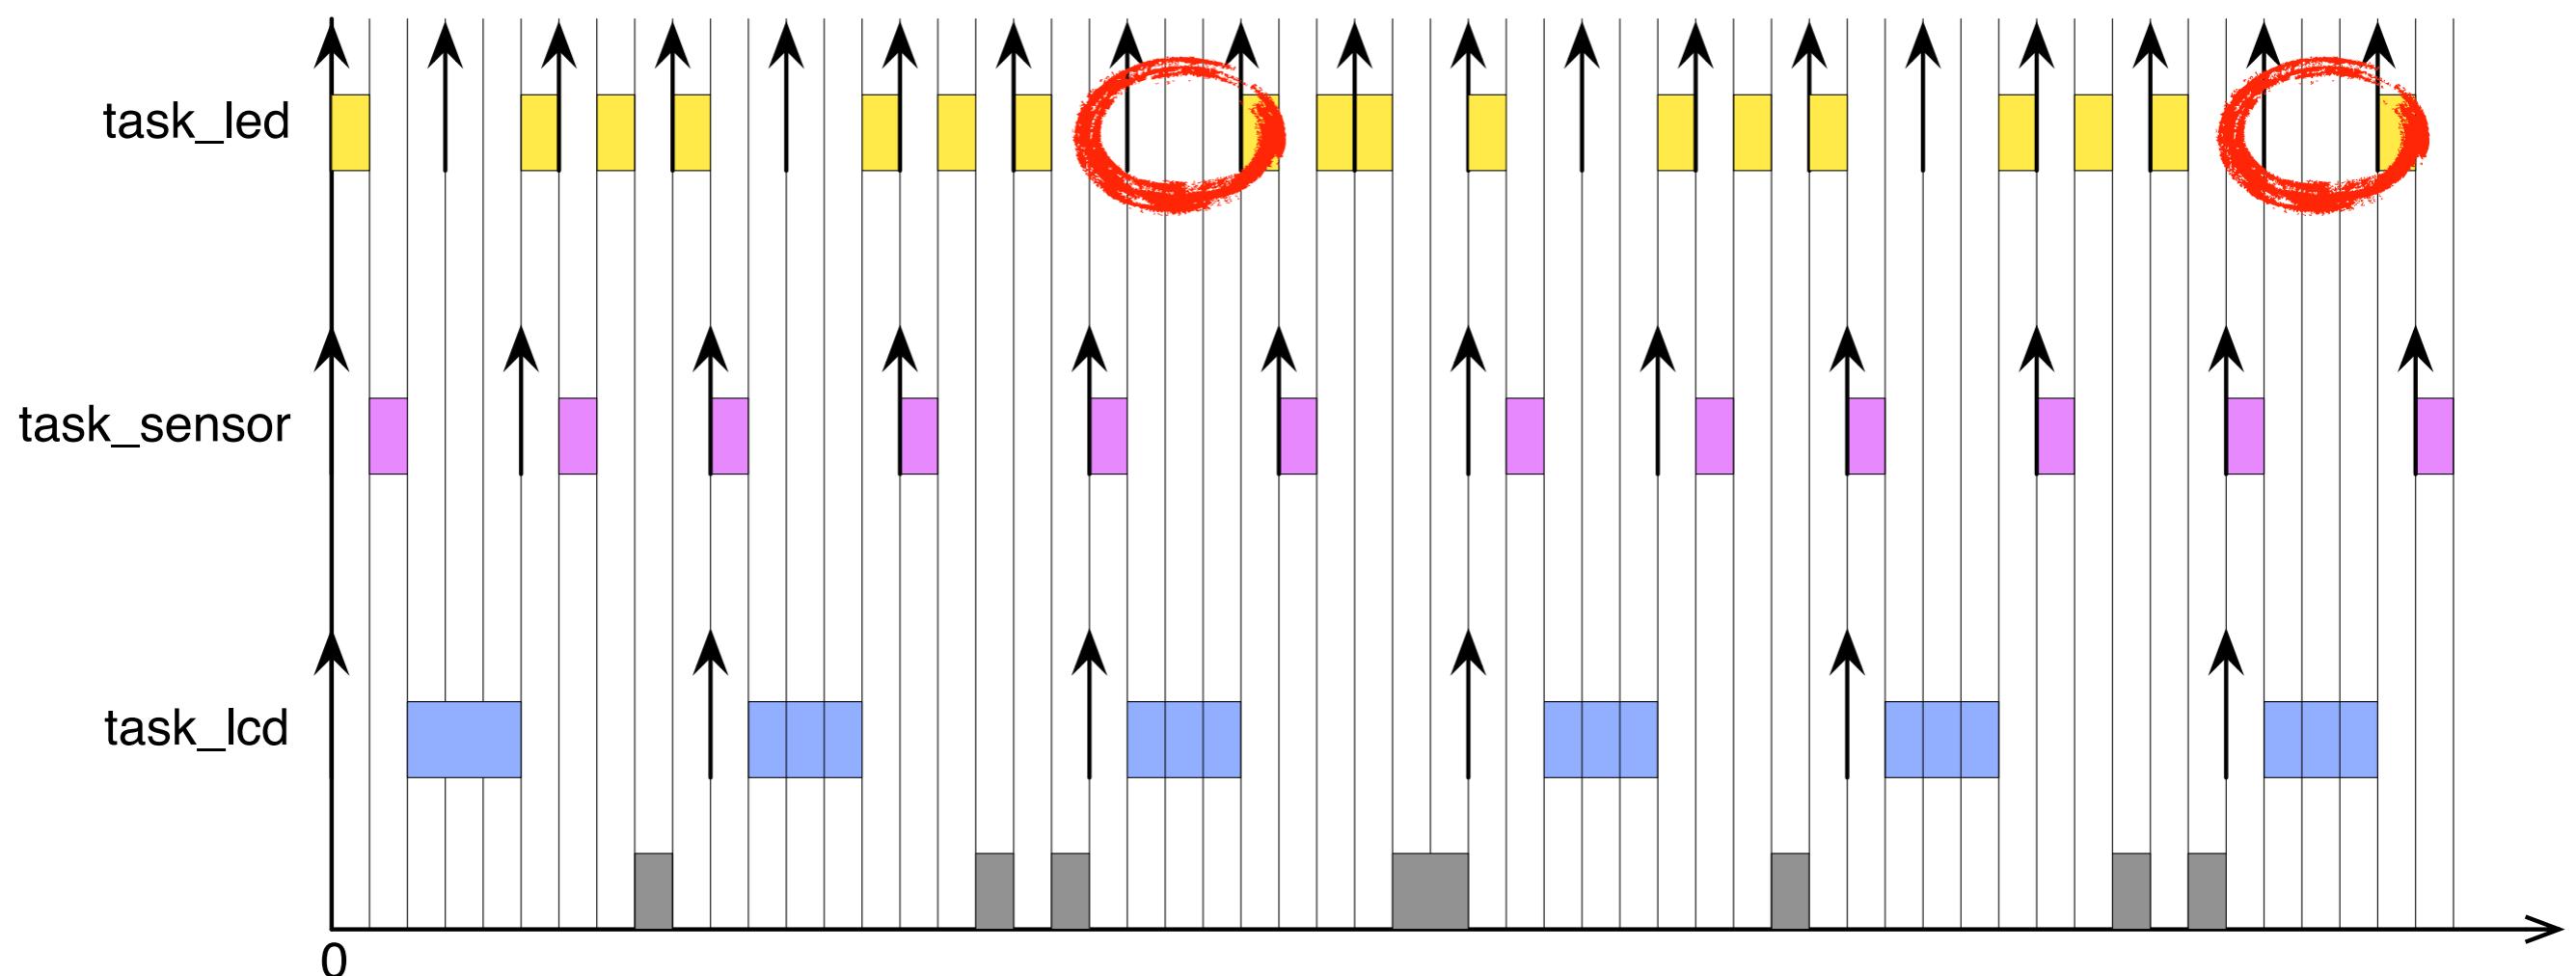
Example of a sequencer

task_sensor > task_led > task_lcd



Example of a sequencer

task_led > task_sensor > task_lcd



What does an operating system ?

- Running many task implies:
 - The computing power is shared among tasks
 - Tasks may need to synchronize and access shared resources in a safe way
 - Tasks may need to communicate
- An operating system:
 - decides what task run according to a scheduling policy
 - provides services to synchronize tasks and manage shared resources
 - provides services to communicate

Why a desktop OS is not suited ?

- Too much functionalities not needed in embedded systems
 - eats a lot of memory space
 - needs too much computing power
- Timing variability too high
- Not designed and not tested for safety / mission-critical applications
- Not designed to be modular and configurable.

What about device drivers ?

- Device drivers are not part of the OS (exception: system timer)
- They are implemented:
 - Using **tasks scheduled by the scheduler** as any other task of the application ⇒ in real-time context this **improve predictability**.
 - Target support is easier for the OS.



Scheduling

How to manage concurrency ?

- Real-time operating system (RTOS) preemptive
 - Uses real-time *scheduling*
 - Periodic system, predictable, deterministic
 - Preemptive system :
 - A running task may be stopped to allow a higher priority task to run and resumed later.
 - Tasks of the application are separated from the RTOS
 - Separation of concerns, reusability, abstraction layer.

Scheduling policies

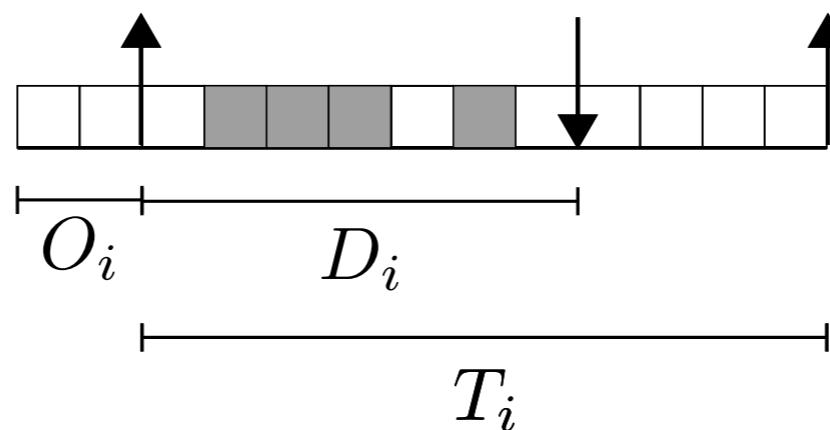
- A RTOS implements a scheduling algorithm.
 - a way to decide which task should run according to a priority
 - remark: the complete execution of a task is called a **job**
 - 2 categories : fixed priority and dynamic priority
 - fixed priority : the priority is assigned offline
 - dynamic priority : the priority is computed dynamically
 - A task may have a dynamic priority and the jobs a fixed one.
- We will see 3 scheduling policies :
 - RM : Rate monotonic (fixed)
 - DM : Deadline monotonic (fixed)
 - EDF : Earliest Deadline First (dynamic for tasks, fixed for jobs)

Definitions

- The **workload** (W) is composed of a set of tasks
- A **task** generates recurrently a potentially infinite set of jobs.
- Each **job** is a computation that must be executed on a processor. In a real-time system, a job has a WCET (Worst Case Execution Time) and a deadline.
 - as it is generally intractable to compute the WCET of a job, a conservative estimation is used
 - in critical systems, WCET estimations should not be considered as safe
- The scheduling algorithm is in charge of allocating jobs to processors so that each job receives enough computation time to finish before its deadline.

Task set, task

- A task set W is defined as a finite set of n independant tasks: $W = \{\tau_1, \dots, \tau_n\}$
- Each task τ_i is defined as a tuple $\tau_i = \langle O_i, T_i, C_i, D_i \rangle$ where:
 - $O_i \in \mathbb{N}$ is the offset of the task;
 - $T_i \in \mathbb{N}$ is the period of the task;
 - $C_i \in \mathbb{N}$ is the WCET of the task;
 - $D_i \in \mathbb{N}$ is the relative deadline of the task.



Jobs

- Each task τ_i generates a potentially infinite set of jobs:
 $J_i = \{\tau_{i,1}, \tau_{i,2}, \dots, \tau_{i,j}, \tau_{i,j+1}, \dots\}$
- The task set W generates: $J_W = \tau_i \in W J_i = \{\tau_{1,1}, \tau_{1,2}, \dots, \tau_{2,1}, \tau_{2,2}, \dots\}$
- A job $\tau_{i,j}$ is characterized by:
 - $a_{i,j} = O_i + (j - 1)T_i$ its arrival time;
 - $c_{i,j} = C_i$ its computation time;
 - $d_{i,j} = a_{i,j} + D_i$ its absolute deadline;

Rate Monotonic (RM)

- Priorities are assigned according to the periods of the tasks.
- The **lower** is the period, the **higher** is the priority.
- Exemple : 3 tasks
 - t1, period 10ms
 - t2, period 20ms
 - t3, period 7ms
- Priorities are assigned so that $P(t3) > P(t1) > P(t2)$

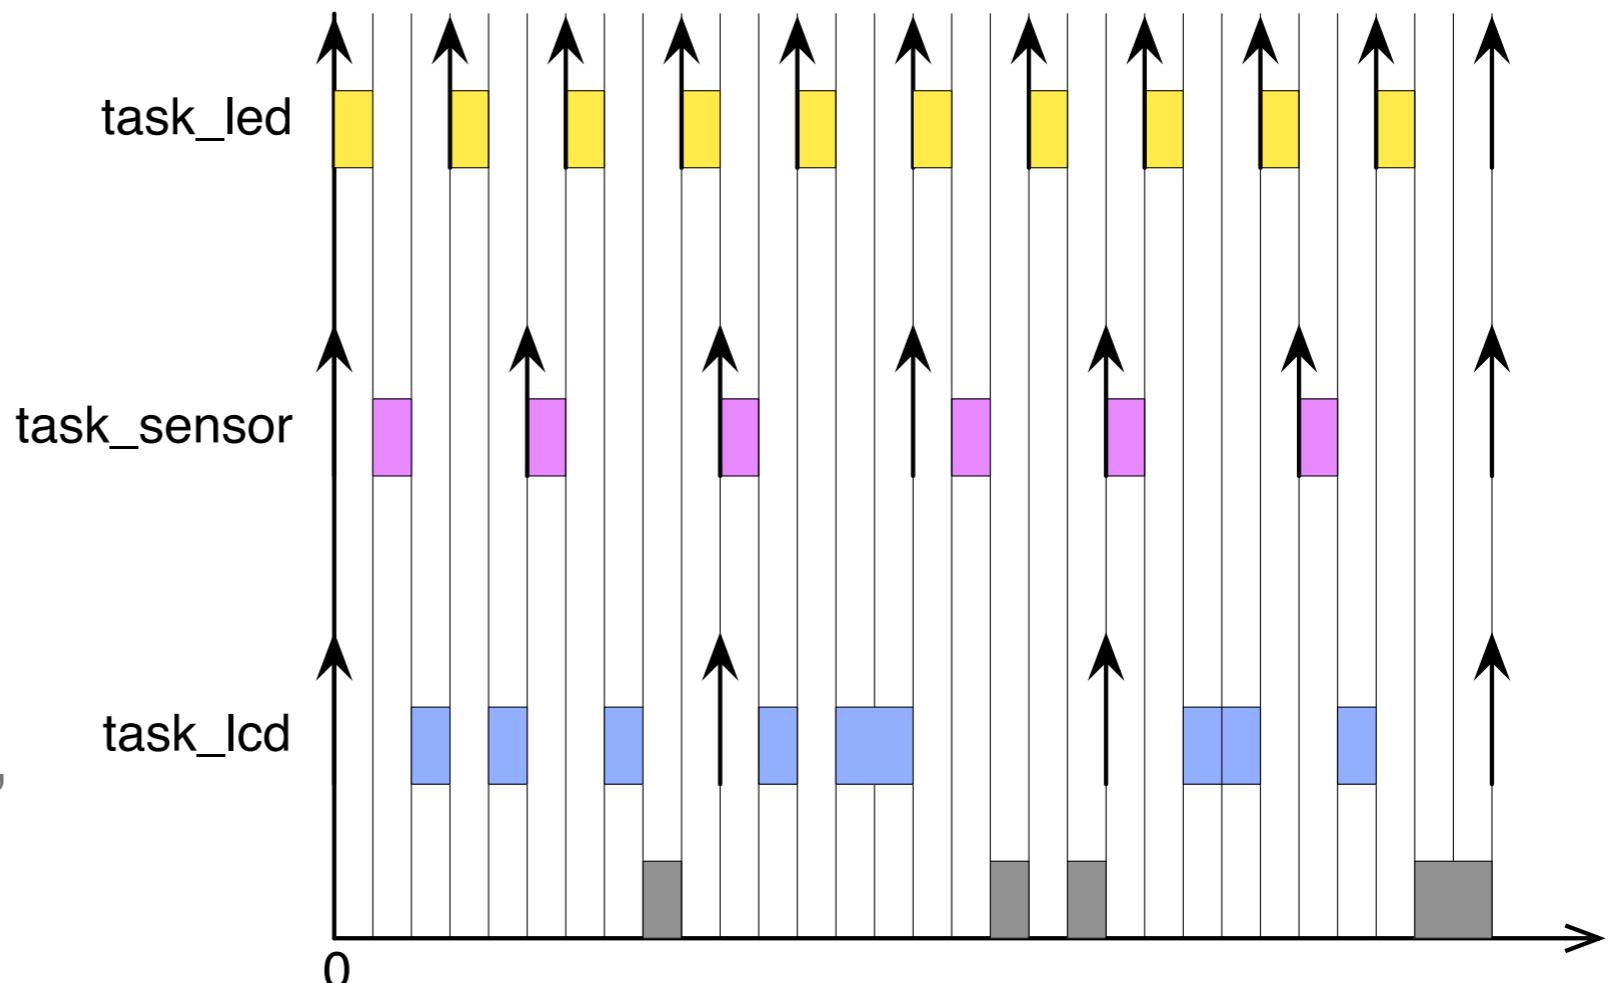
Optimality of RM

- Theorem (Optimality of RM, [LL73]) Preemptive Rate Monotonic is optimal when:
 - all tasks are independent
 - all tasks are synchronous: $\forall \tau_i O_i = a_{i,1} = 0$
 - all tasks are periodic: $\forall \tau_{i,j} a_{i,j} = (j - 1)T_i$
 - deadlines equal periods $\forall \tau_i D_i = T_i$
 - execution time of a task is constant: $\forall \tau_{i,j} c_{i,j} = C_i$
 - a job cannot voluntarily suspend itself
 - context switch time and other OS overheads are zero
 - inter-task precedence constraints are enforced
 - there is a single processing unit
 - scheduling algorithm is on-line, fixed-task priority, preemptive, work conserving

$$U = \sum_{\tau_i \in W} \frac{C_i}{T_i} \leq n(2^{1/n} - 1)$$

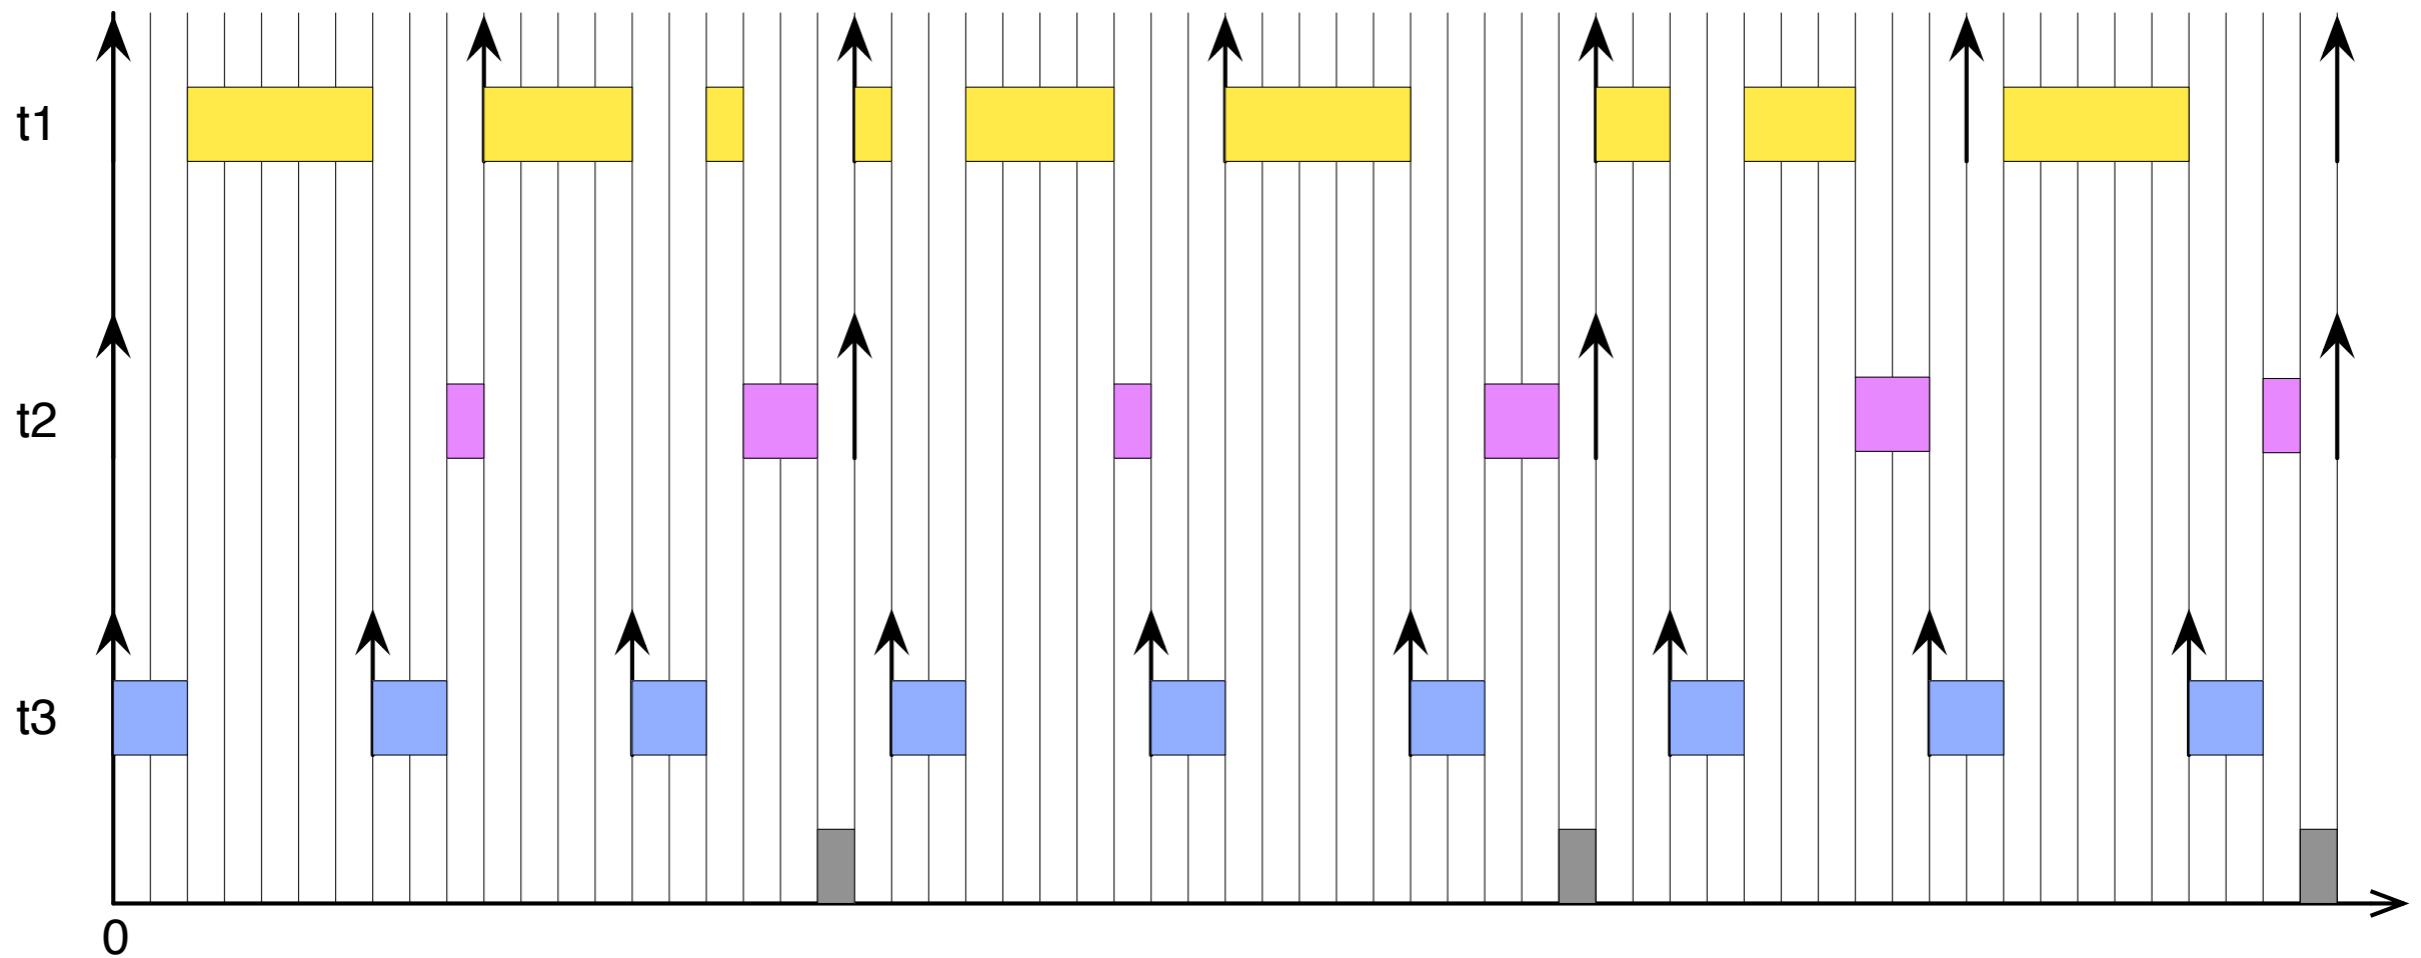
Example of rate monotonic

- From the previous example:
 - task_led:
period 3ms, execution time 1ms
 - task_sensor:
period 5ms, execution time 1ms
 - task_lcd: period 10ms, execution time 3ms



Another example

- t1, period 10ms, execution time 5ms
- t2, period 20ms, execution time 3ms
- t3, period 7ms, execution time 2ms



Deadline Monotonic (DM)

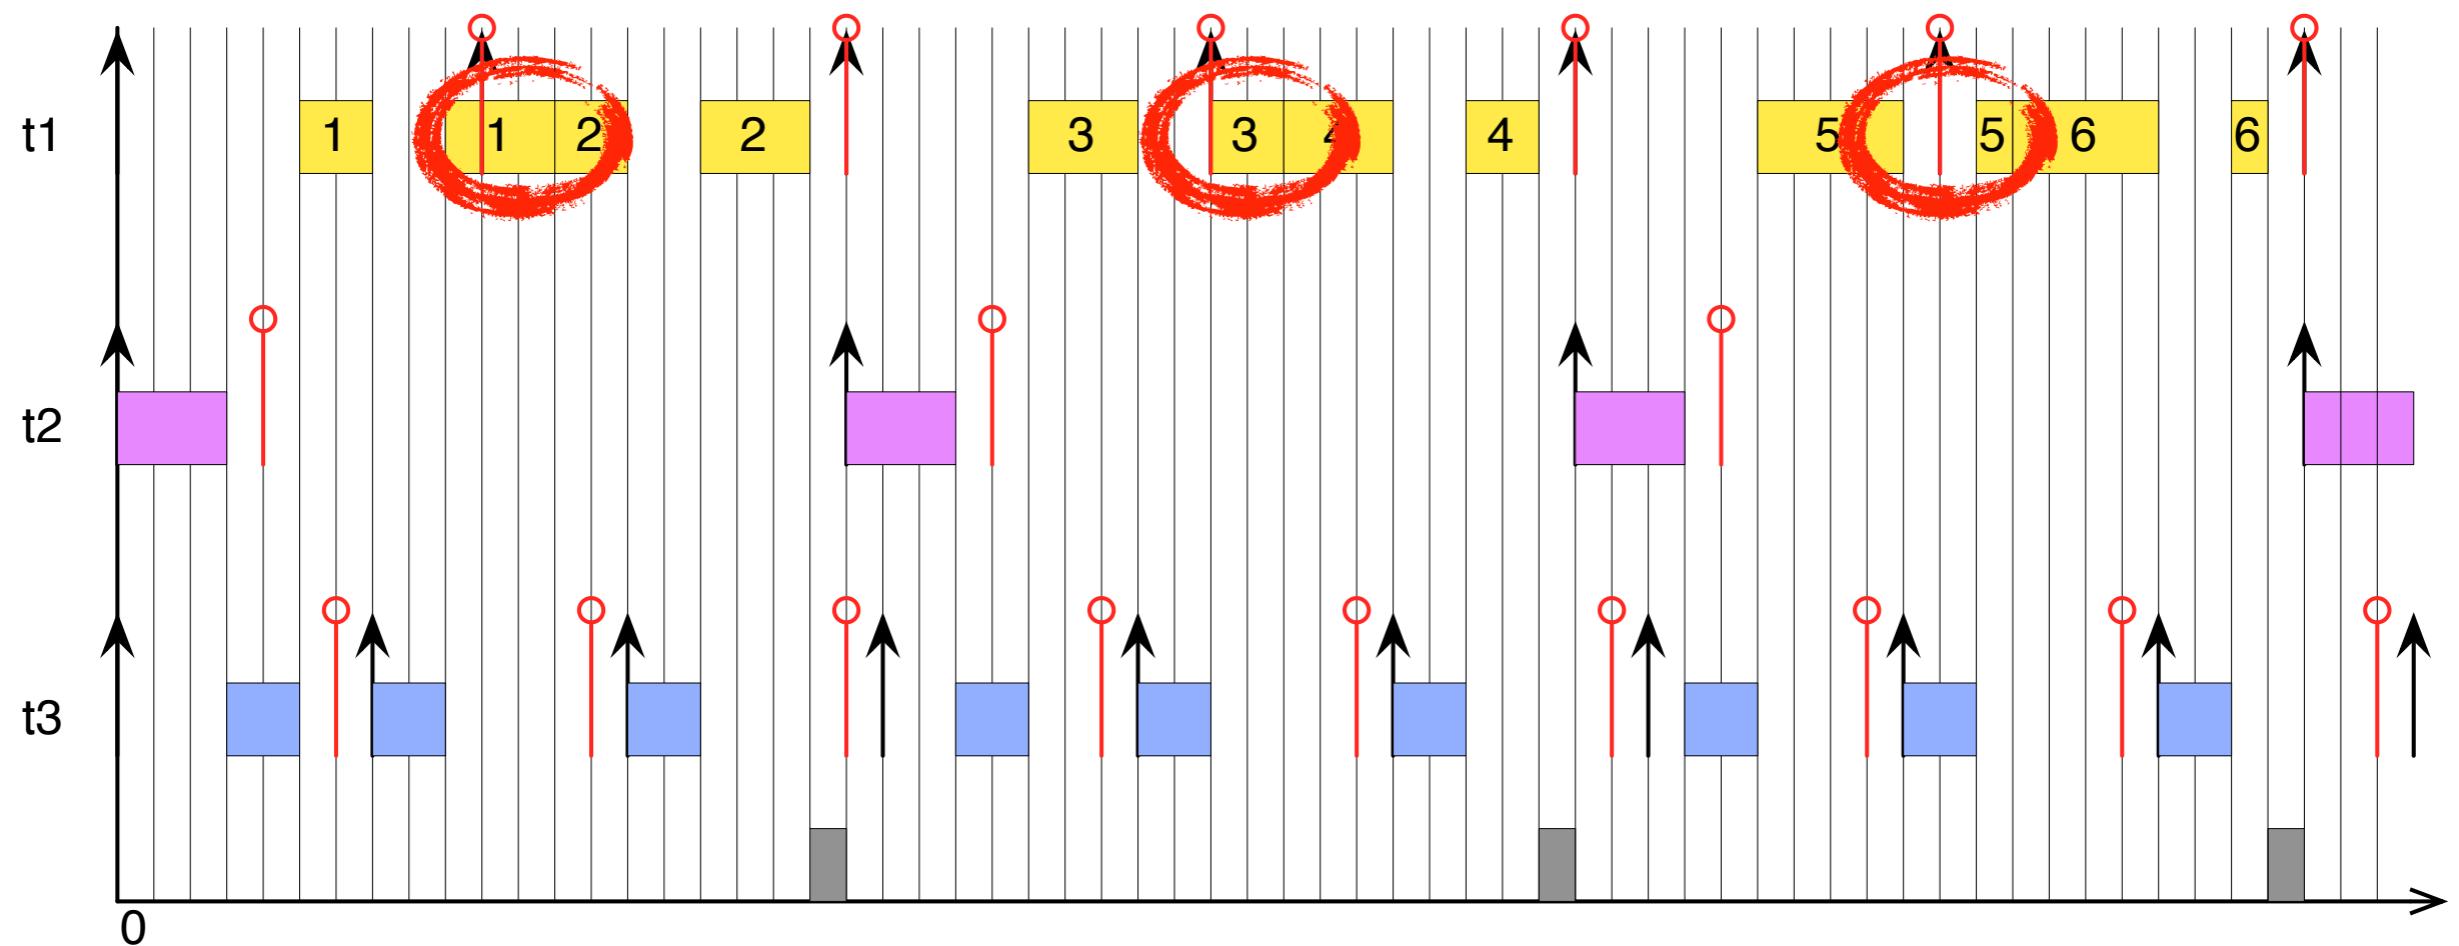
- Priorities are assigned according to the deadline of the tasks.
- The lower is the deadline, the higher is the priority.
- Exemple : 3 tasks
 - t1, period 10ms, execution time 5ms, deadline 10ms
 - t2, period 20ms, execution time 3ms, deadline 4ms
 - t3, period 7ms, execution time 2ms, deadline 6ms
 - Priorities are assigned so that $P(t2) > P(t3) > P(t1)$

$$\forall \tau_i \in W \quad C_i + \sum_{\tau_j \in hp(\tau_i)} \left\lceil \frac{D_j}{T_i} \right\rceil C_j \leq D_i$$

- where $hp(\tau_i)$ is the subset of W containing all tasks with priority higher than τ_i :

$$hp(\tau_i) = \{\tau_j \mid \tau_j \in W \wedge D_j < D_i\}$$

Deadline monotonic (DM)

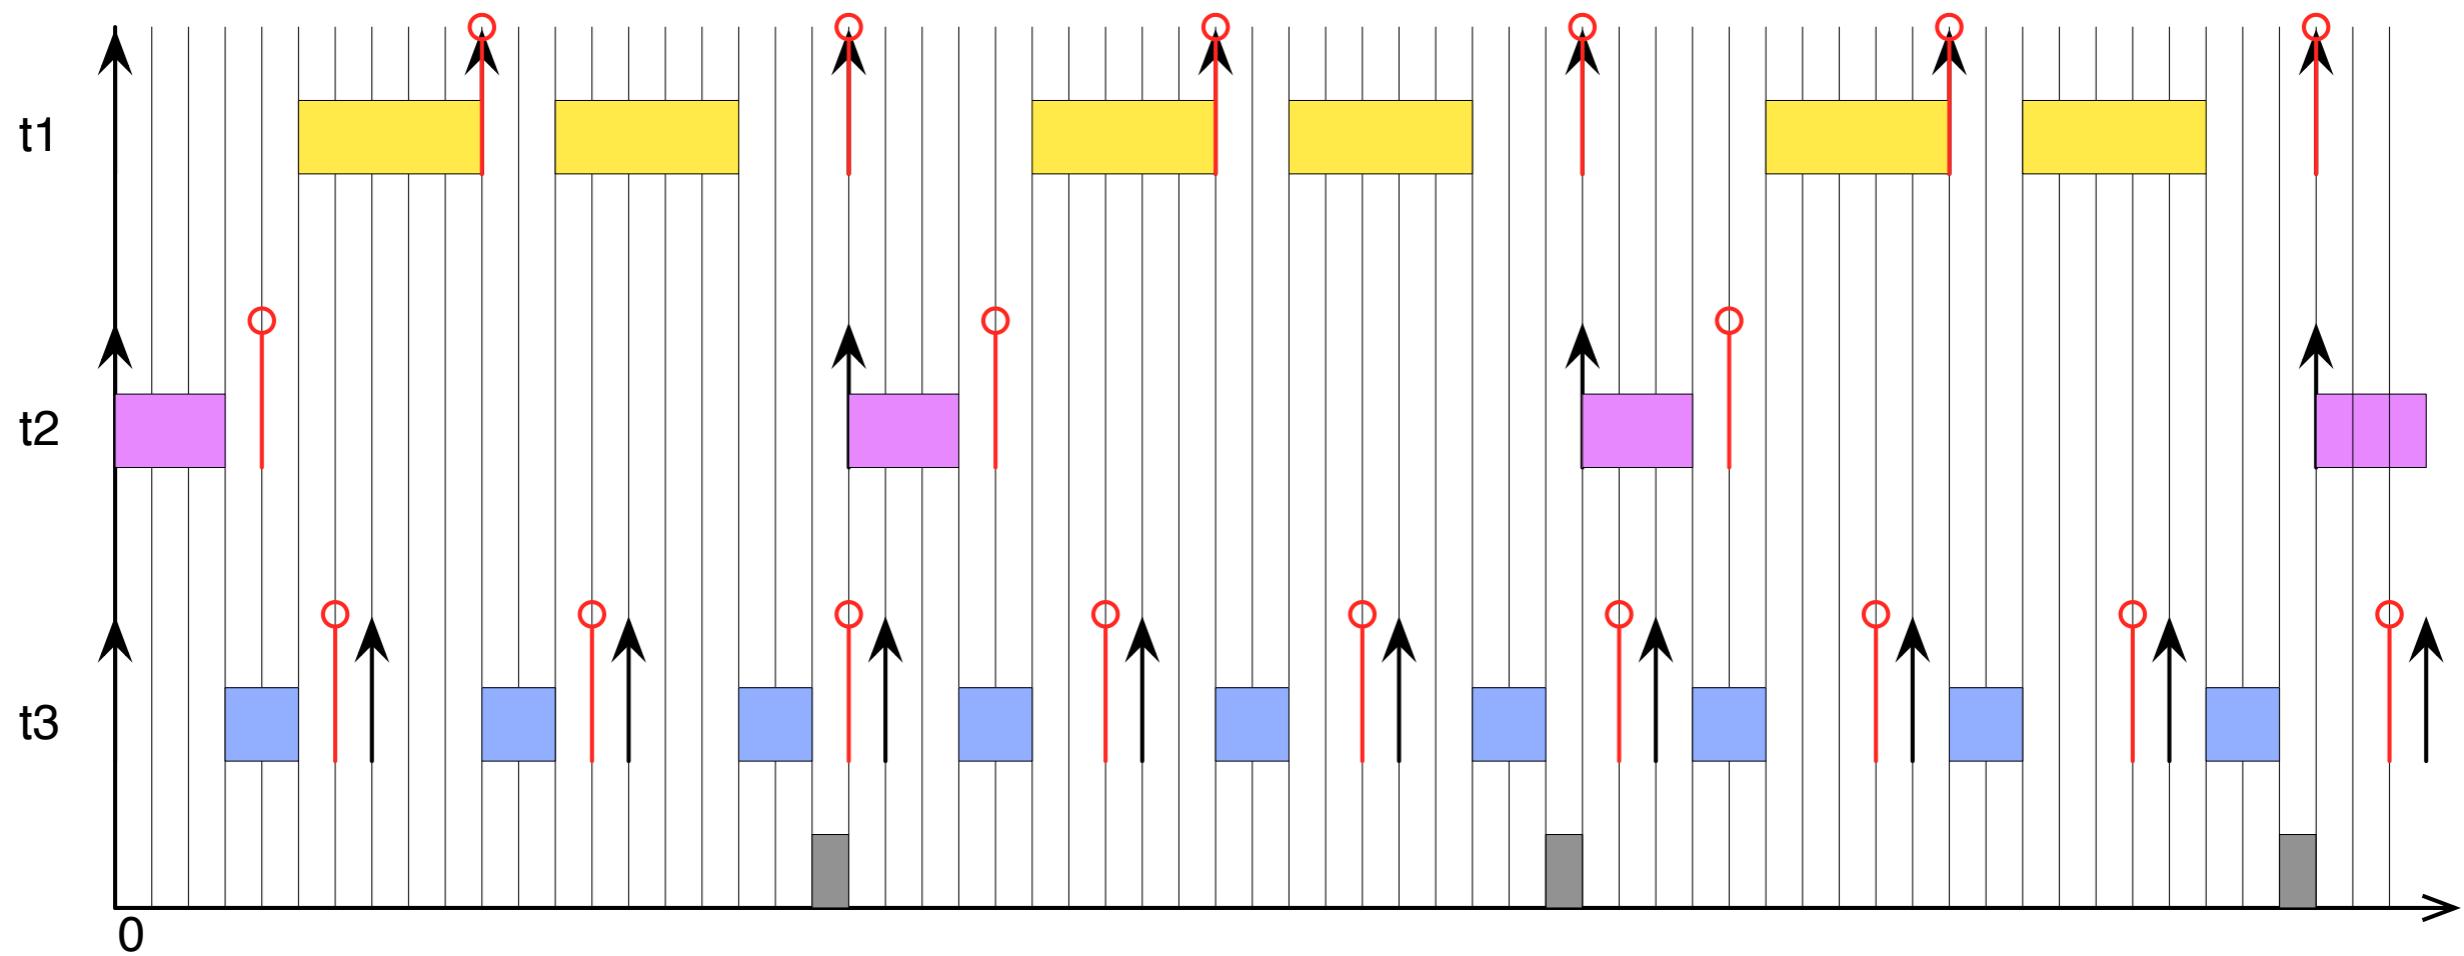


Earliest Deadline First (EDF)

- Priorities are assigned according to the deadline of the tasks at the time a rescheduling occurs.
- The lower is the deadline, the higher is the priority.
- Optimal algorithm.

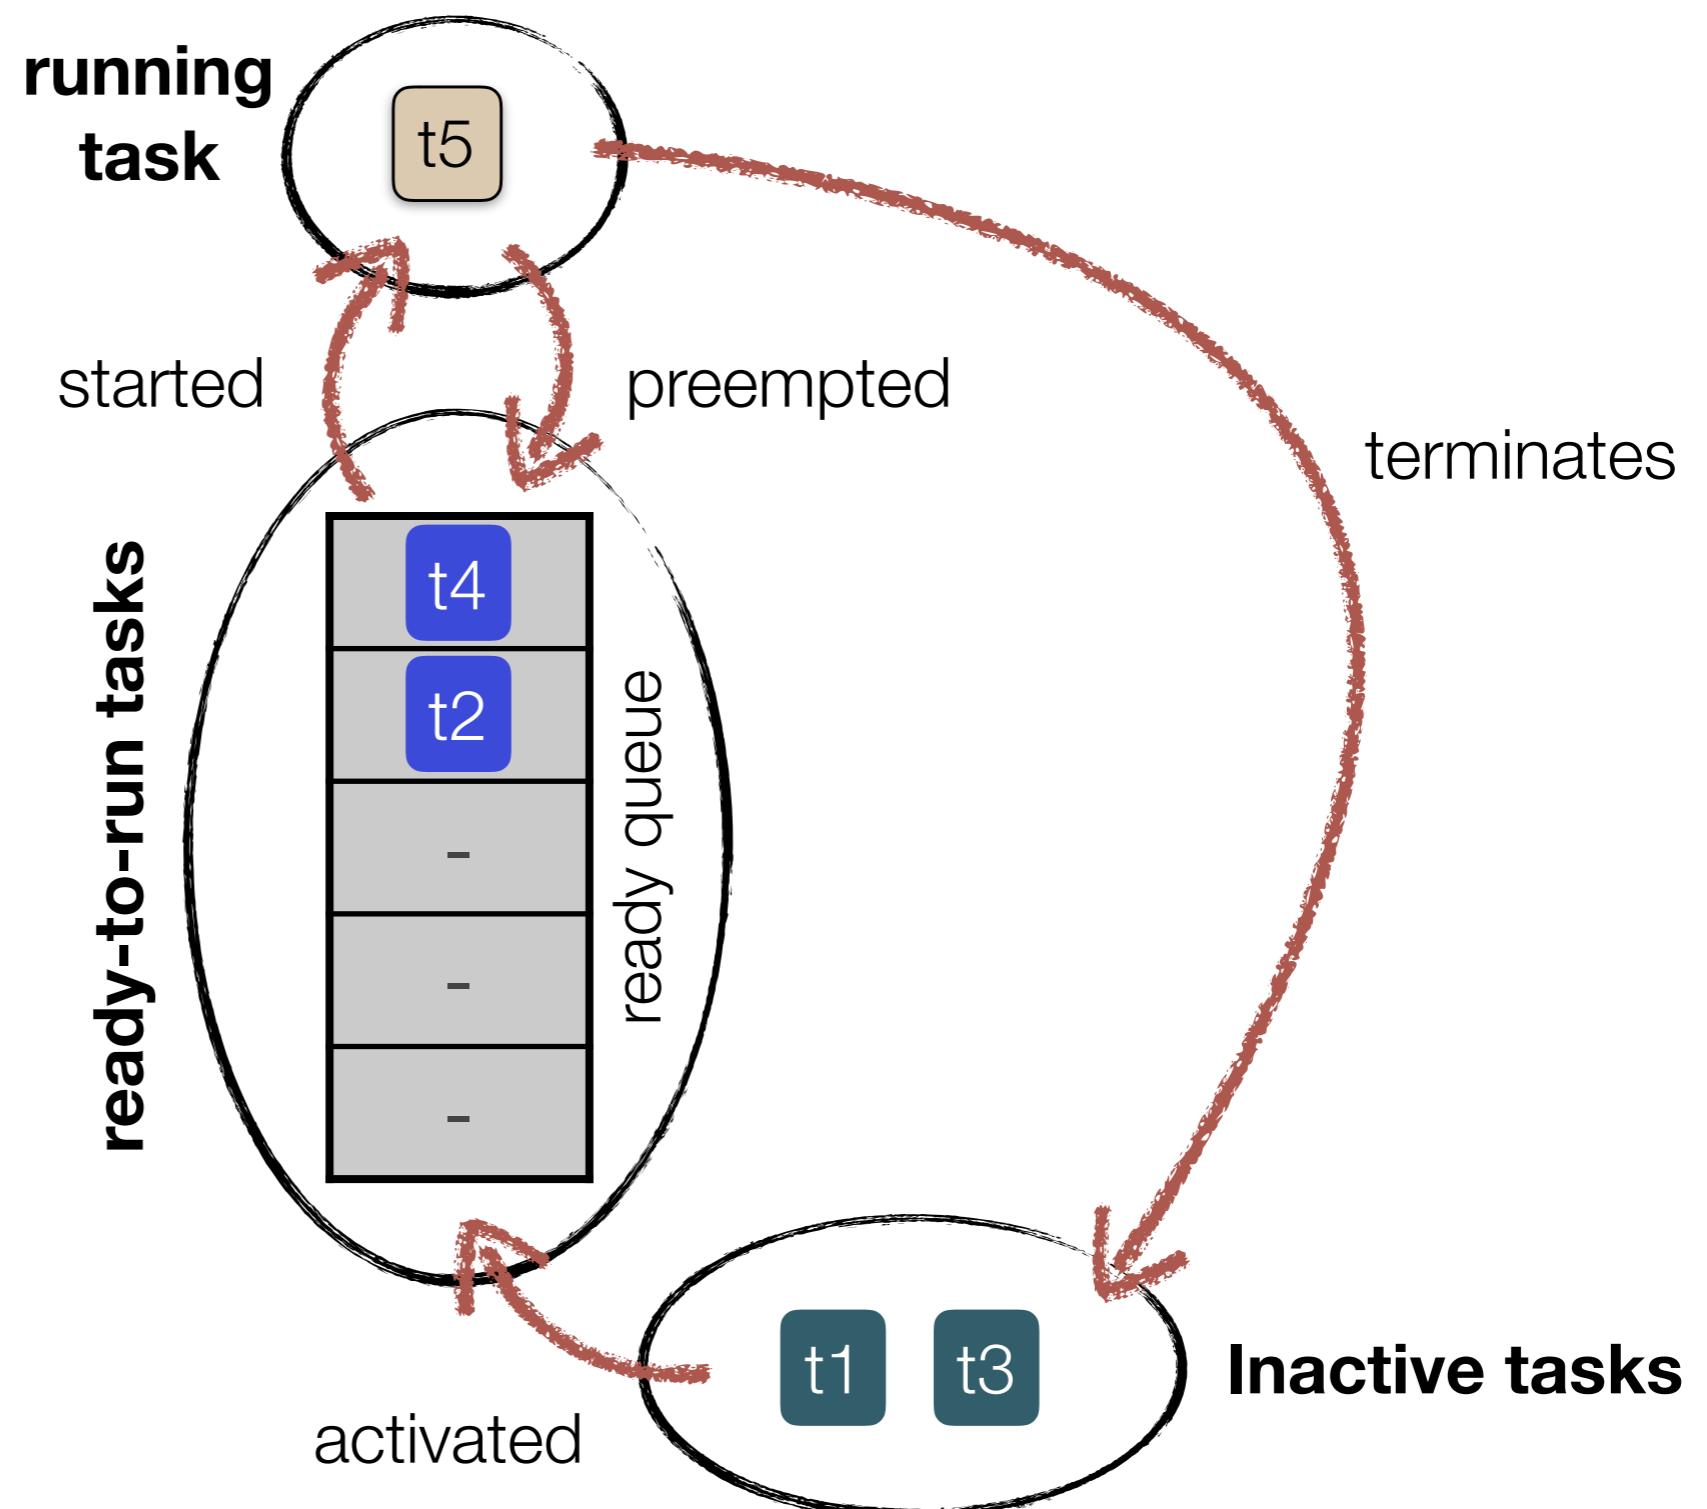
Earliest Deadline First (EDF)

- Exemple : 3 tasks
 - t1, period 10ms, execution time 5ms, deadline 10ms
 - t2, period 20ms, execution time 3ms, deadline 4ms
 - t3, period 7ms, execution time 2ms, deadline 6ms



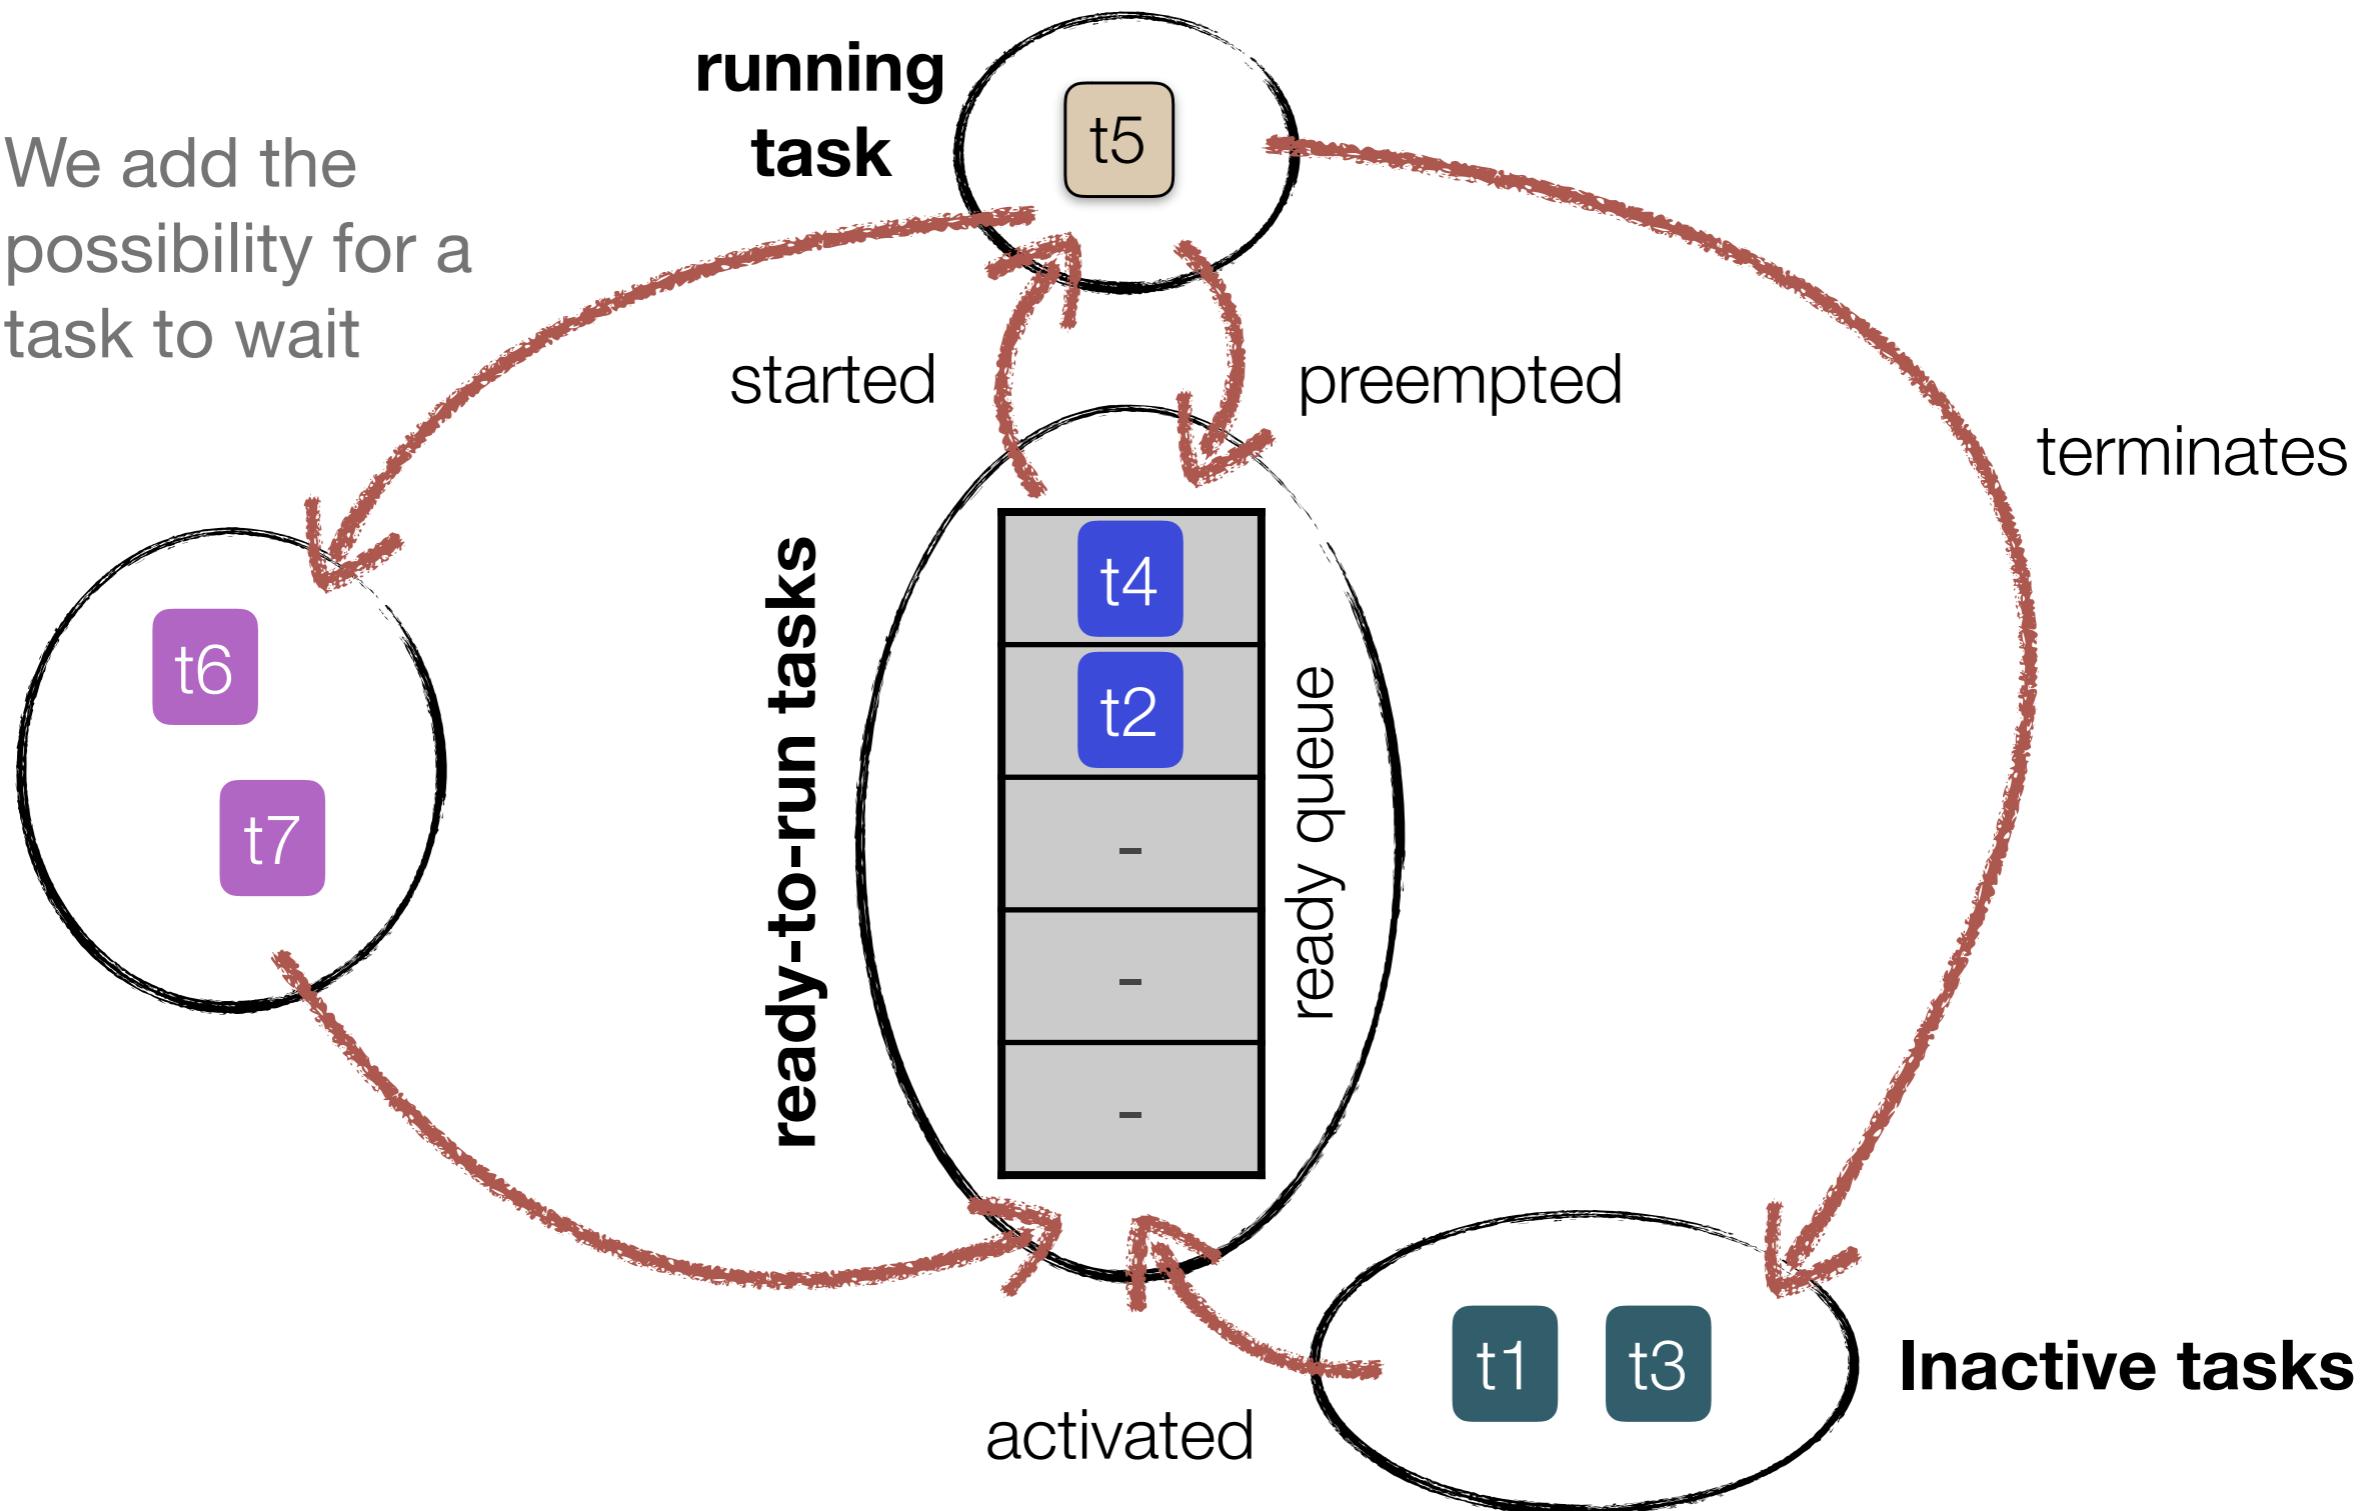
Managing tasks

- The scheduler manages a ready queue: a data structure where ready tasks are stored

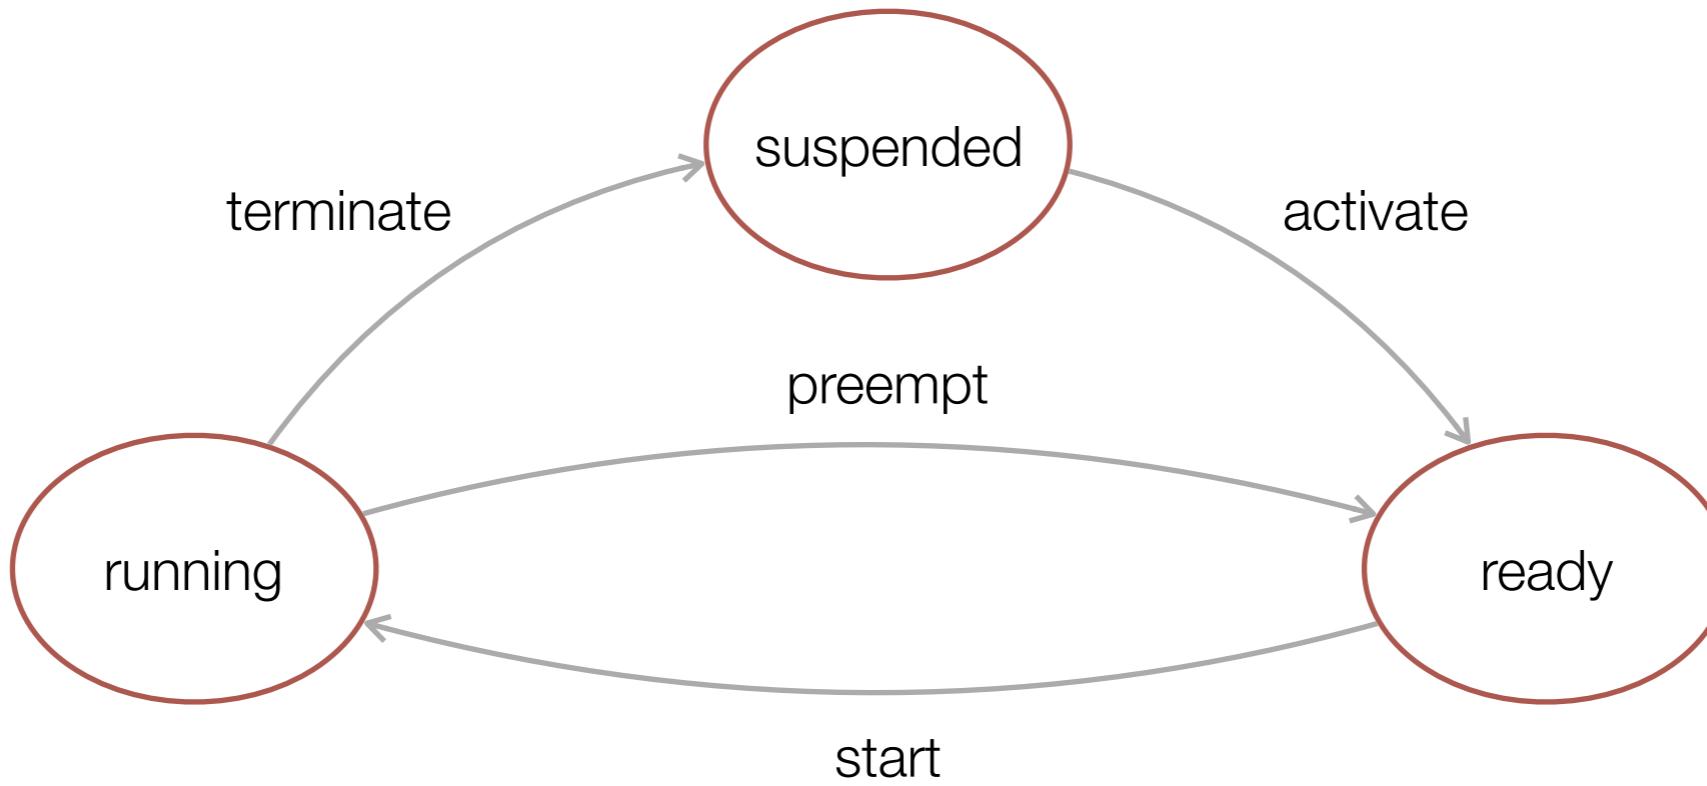


Managing tasks

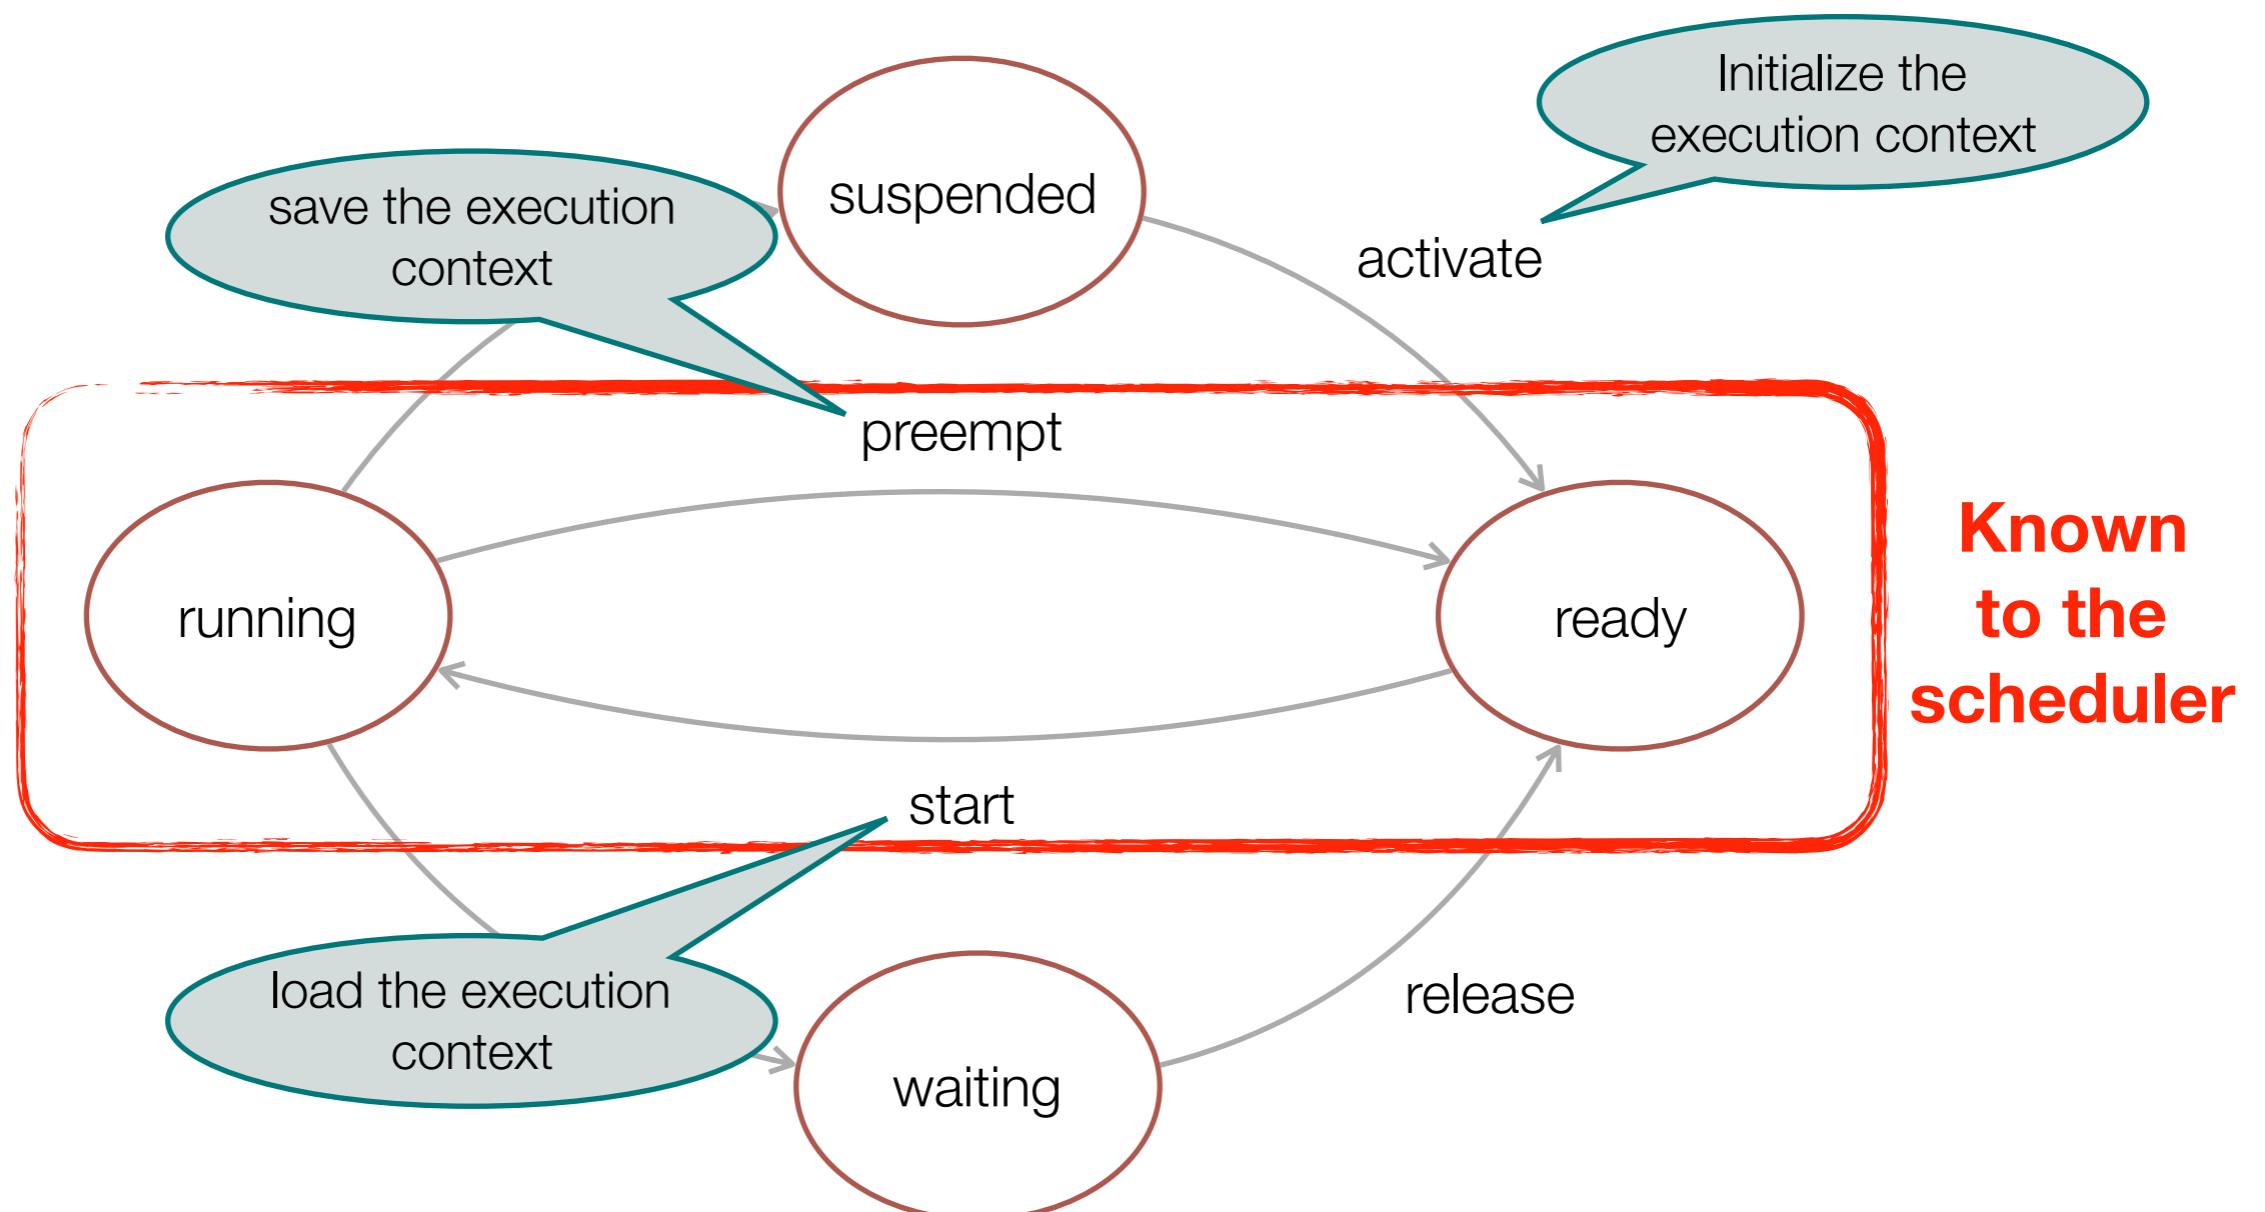
- We add the possibility for a task to wait



Task states



Task states



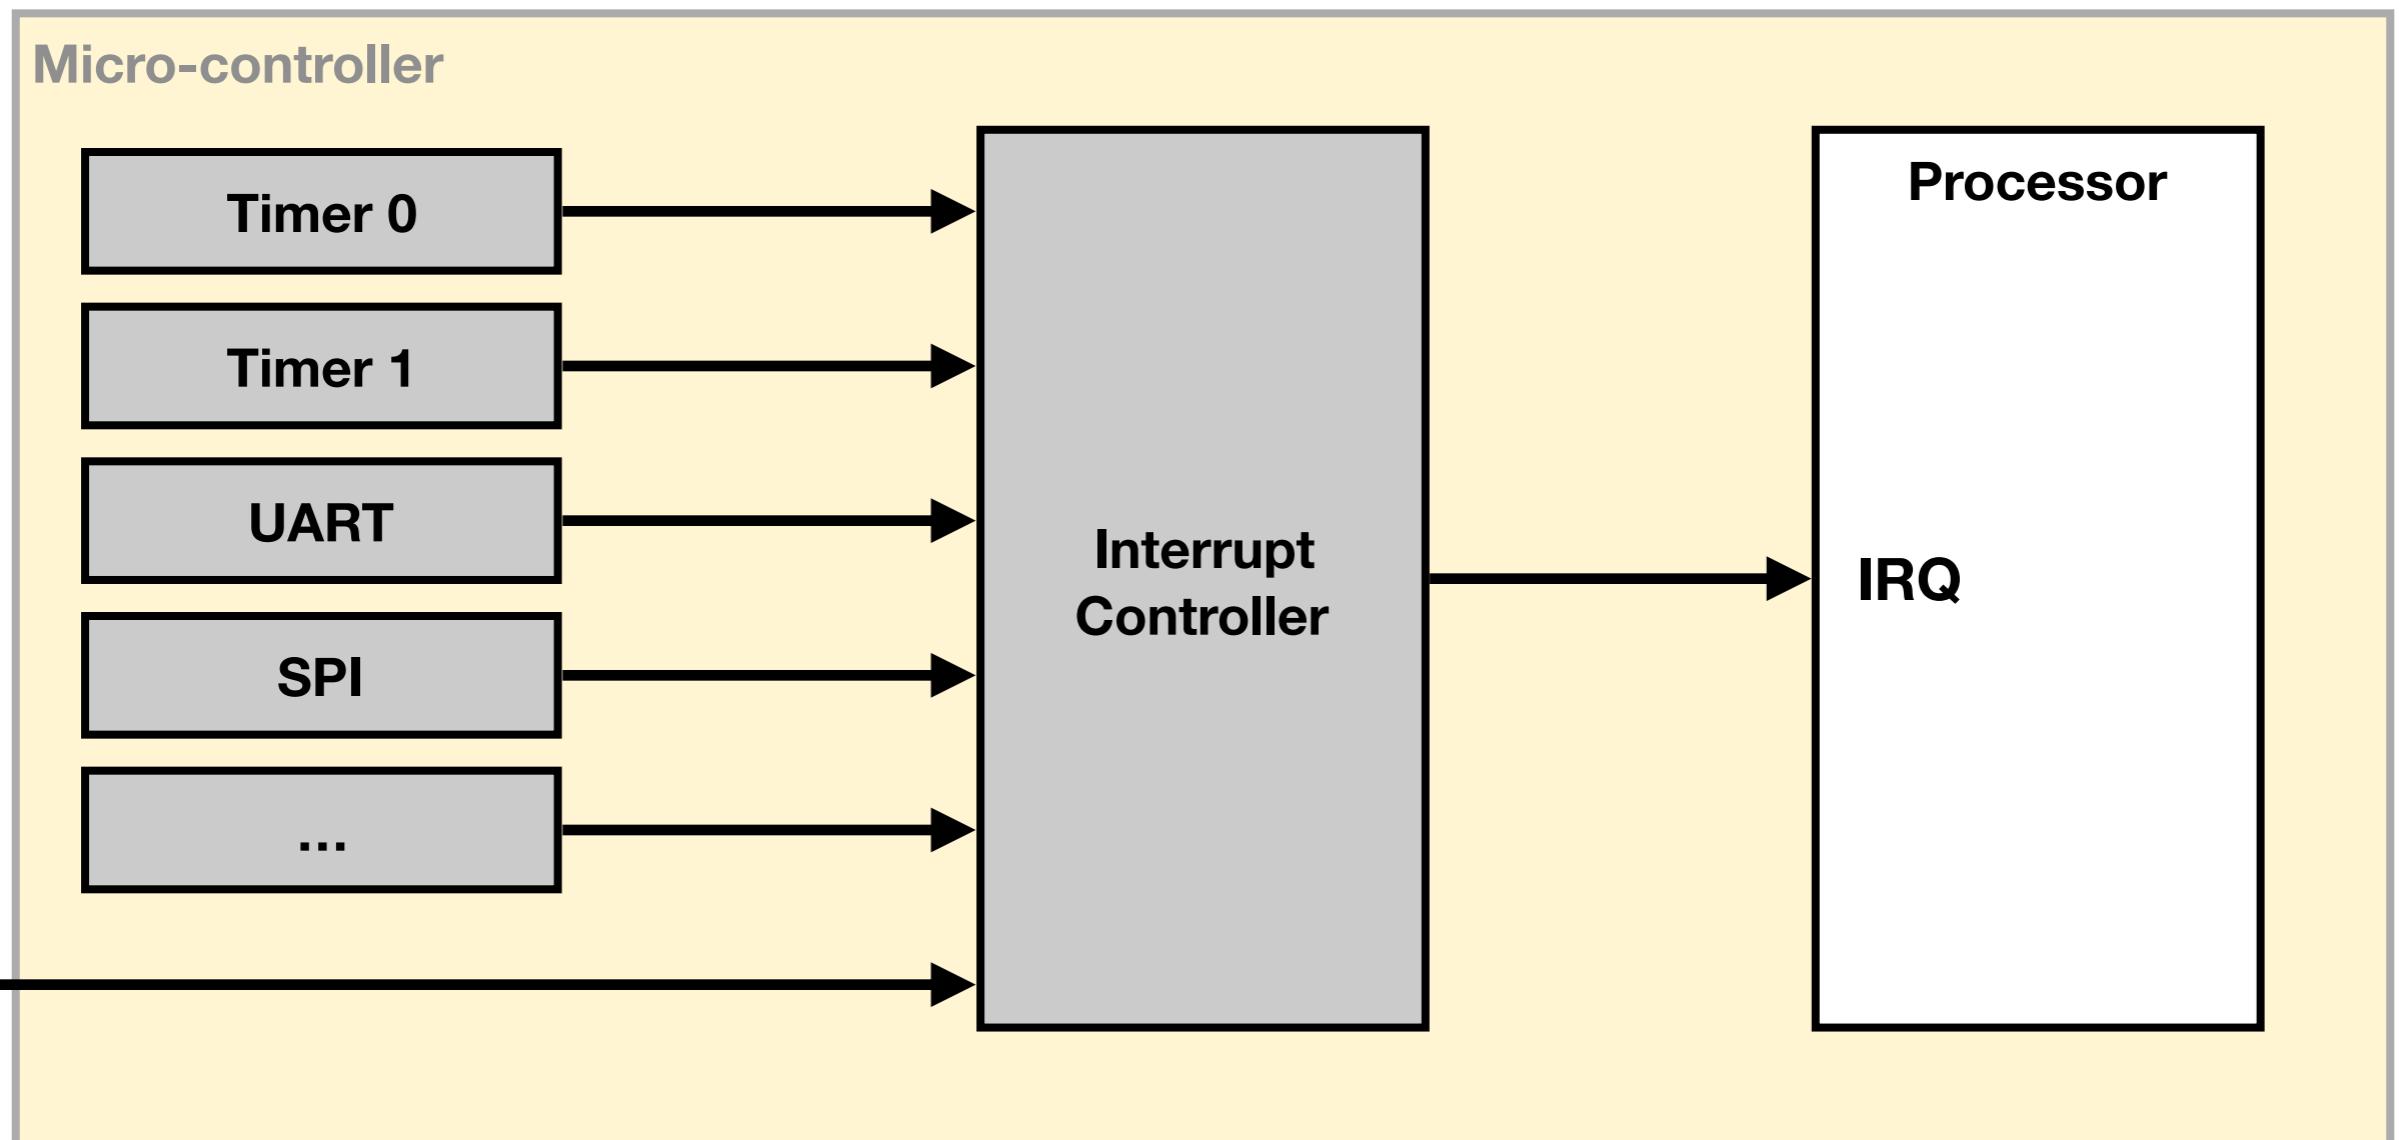
A journey in the real world

- When no task is ready something has to run
 - hidden **idle** task with the lowest possible priority
 - **always ready/running**
 - implements an infinite loop
 - can put the micro-controller in **idle/low power state**
- Periodic tasks (with precise time) / preemption
 - A hardware timer delivers an interrupt periodically (usually called the **systick**)
 - Usually a 1ms period is chosen for this interrupt but it can be adjusted to fit the needs of the application

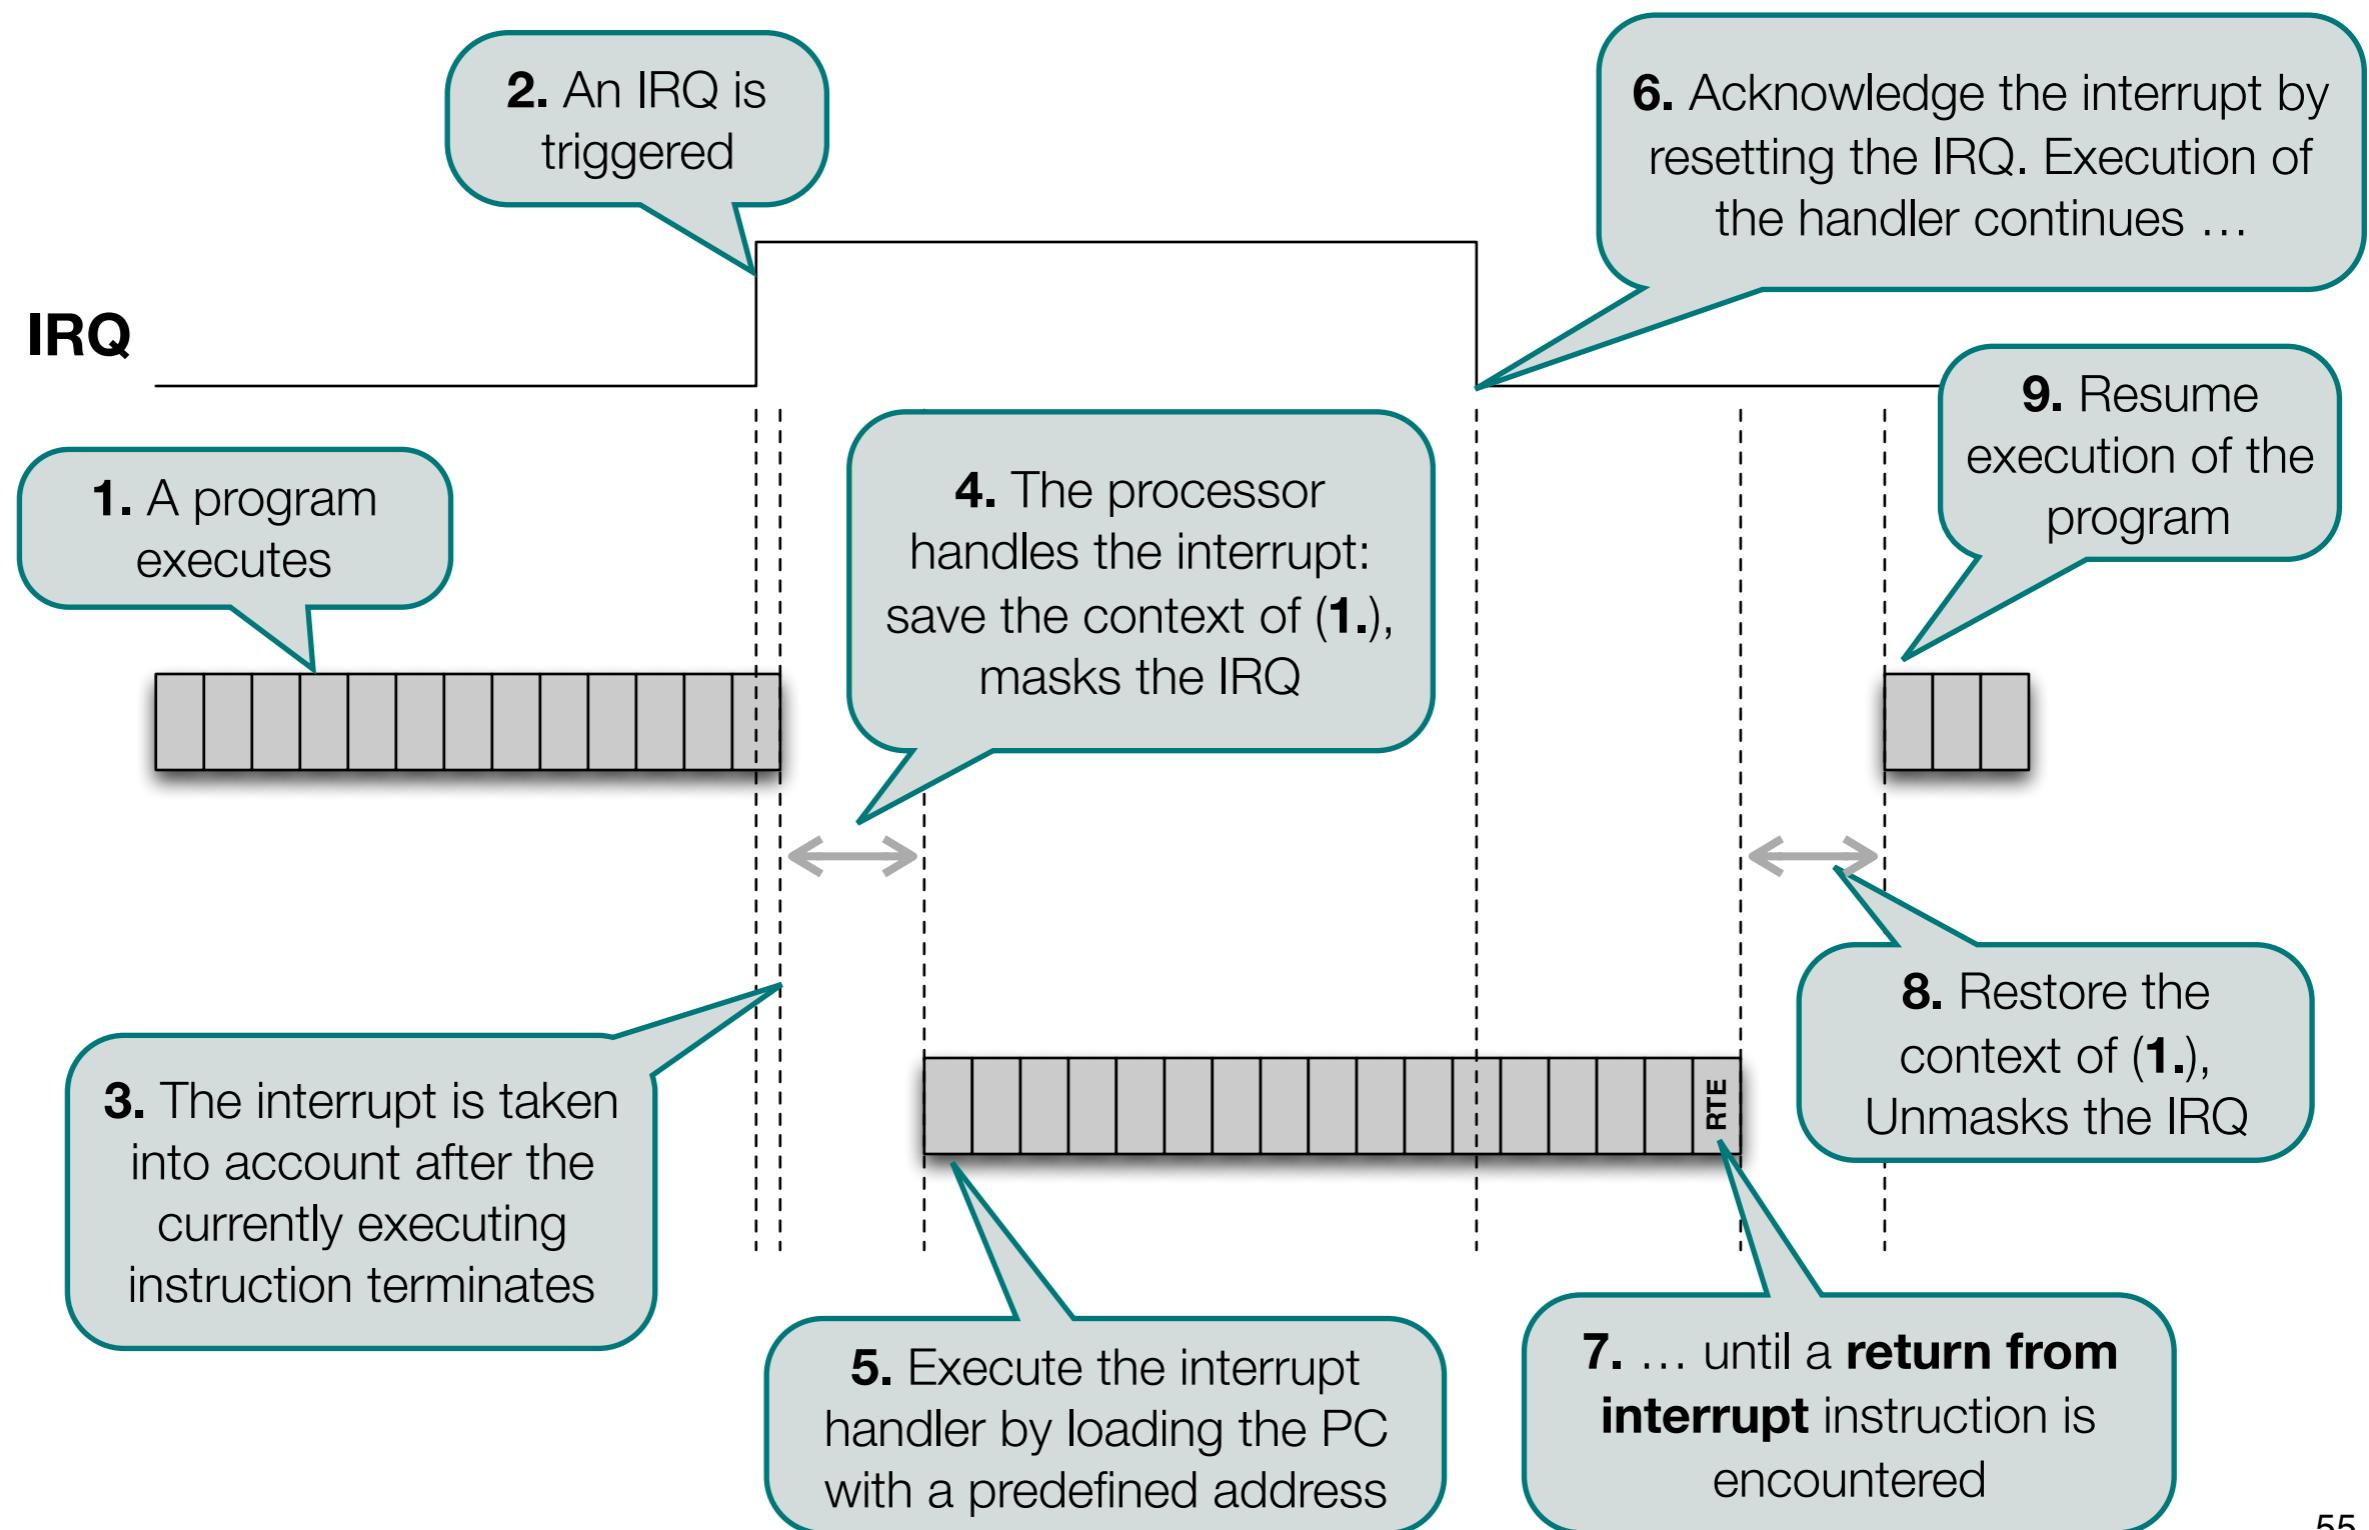
Interrupts (1)

- From hardware point of view, an interrupt line is a hardware input of the processor (**IRQ**, Interrupt ReQuest). The IRQ is a **state**
- The processor may be in states where:
 - the IRQ is ignored (aka masked) or
 - the IRQ is taken into account (unmasked)
- When the IRQ becomes **active** and if the interrupt is **unmasked**, the current executing program is halted and the processor executes another program called the **interrupt handler**.
- When the interrupt handler terminates, the processor goes back to the previously halted program.

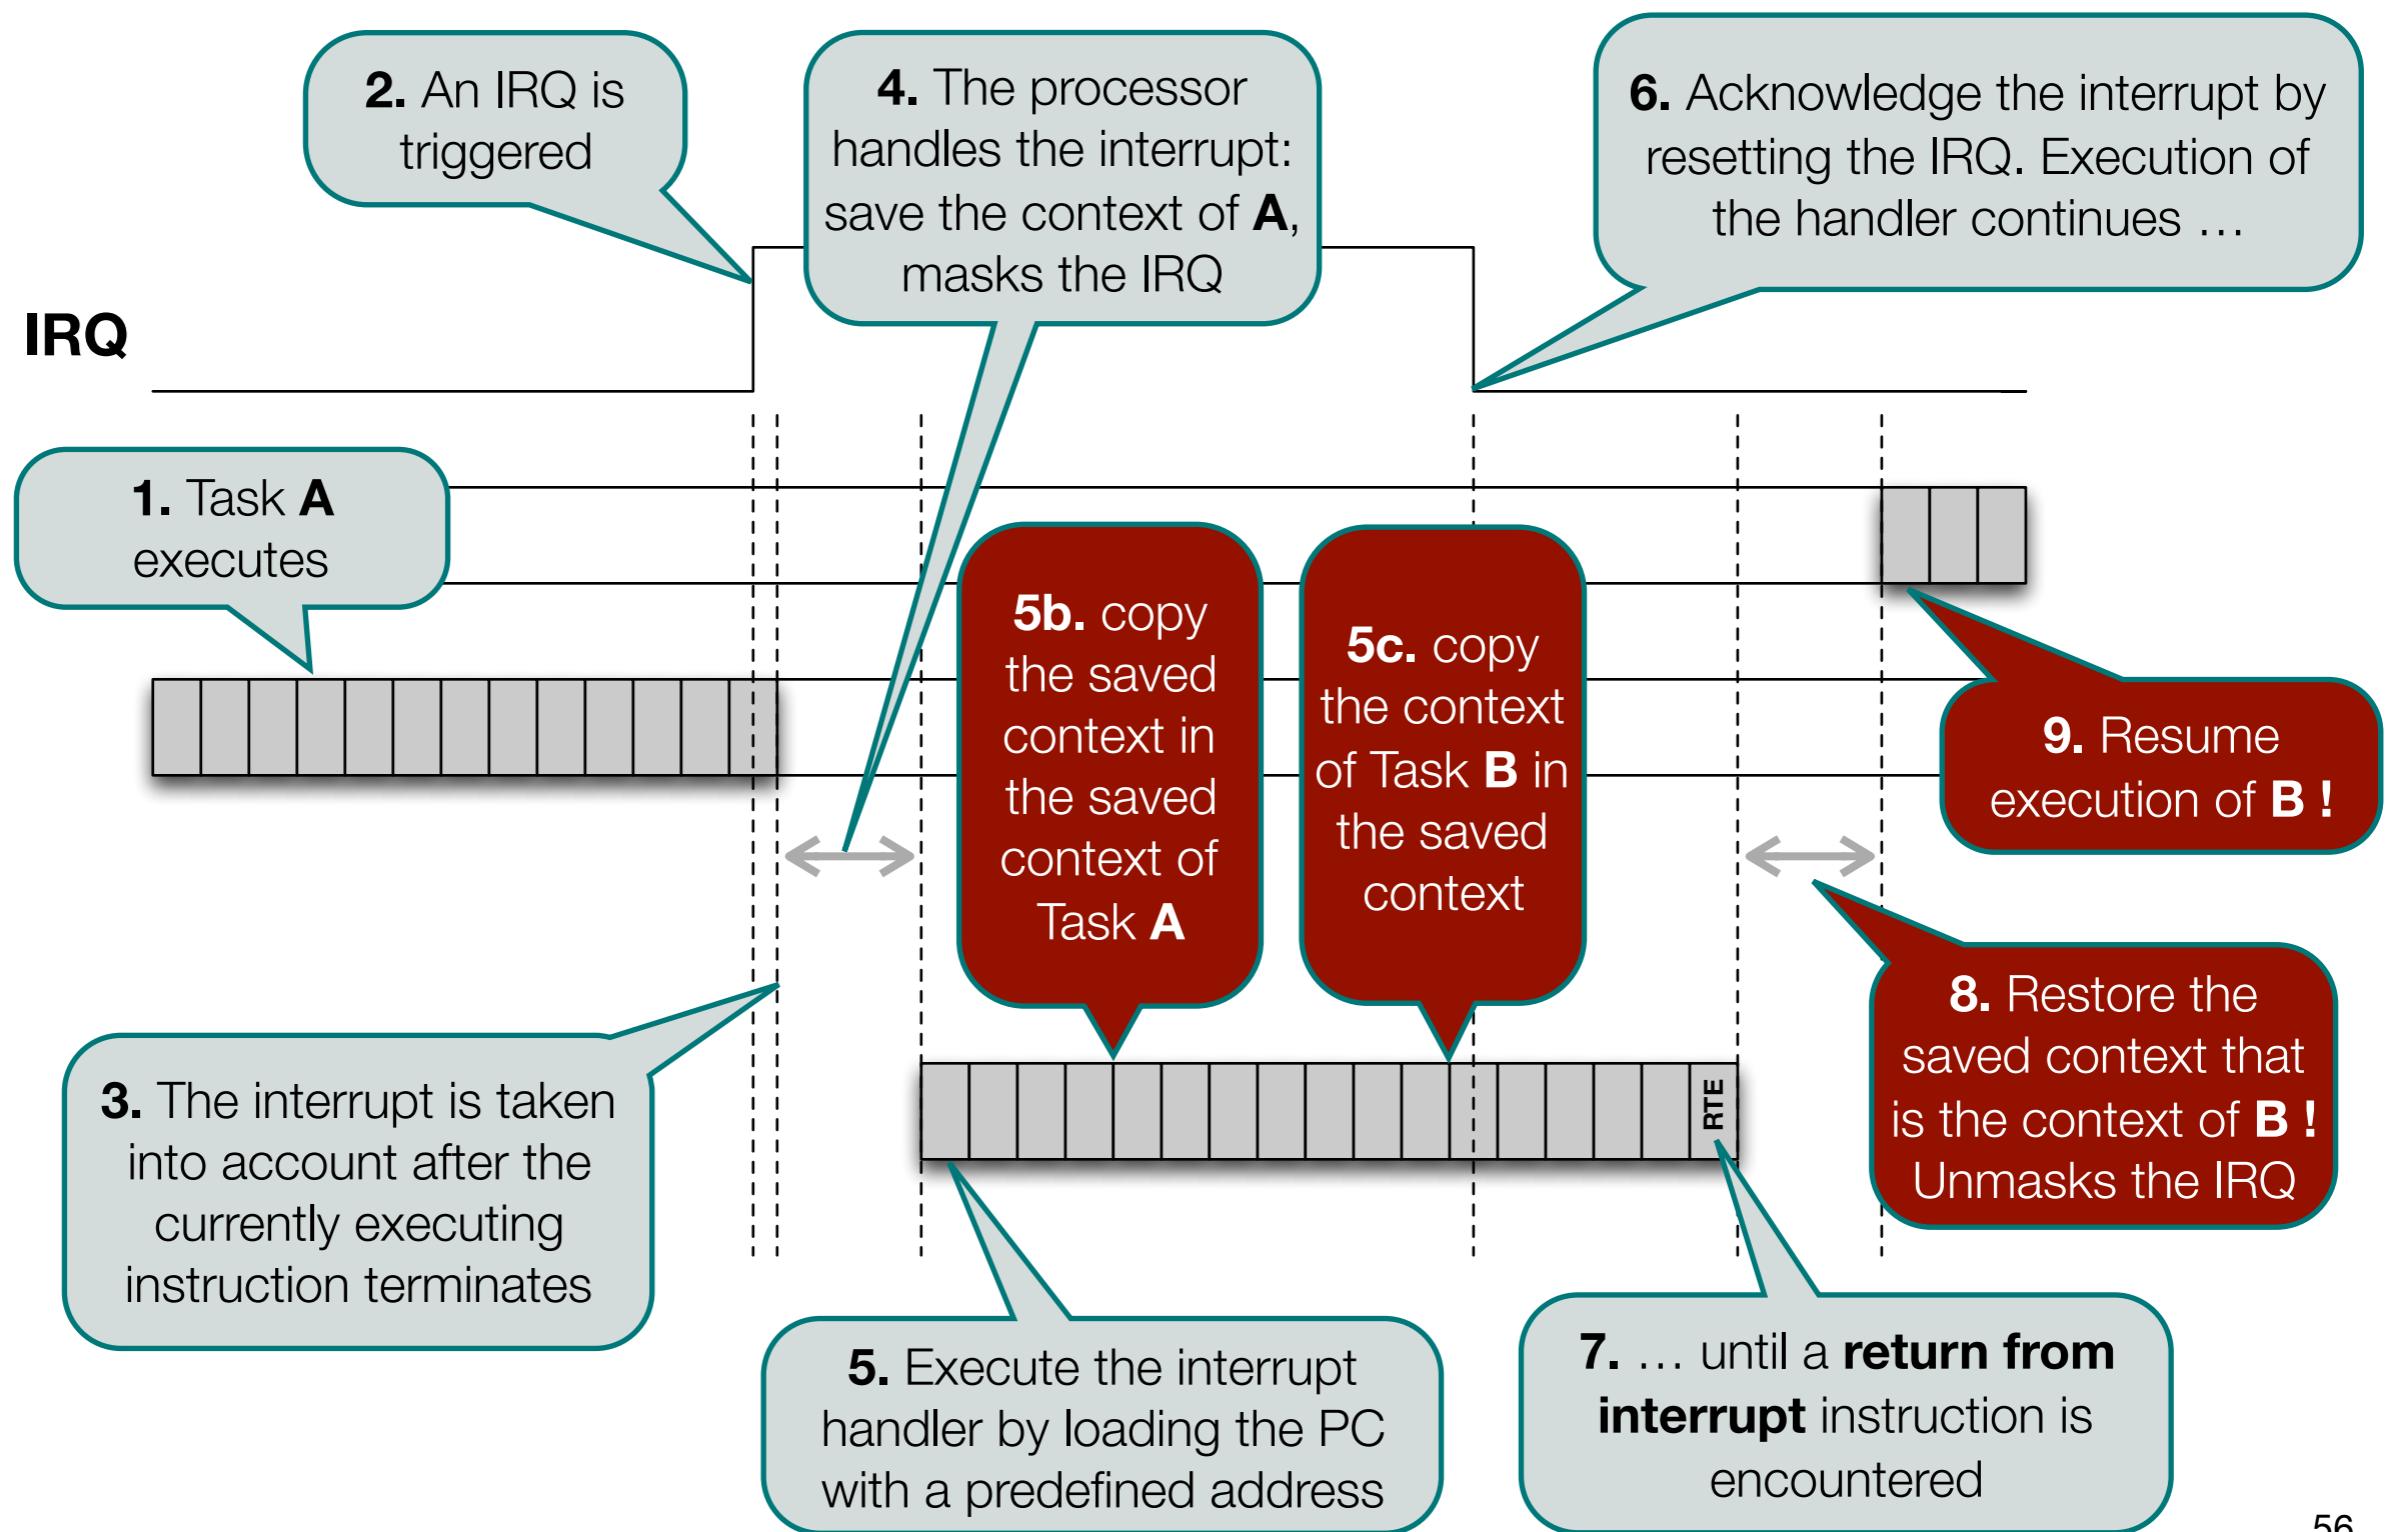
Interrupts (2)



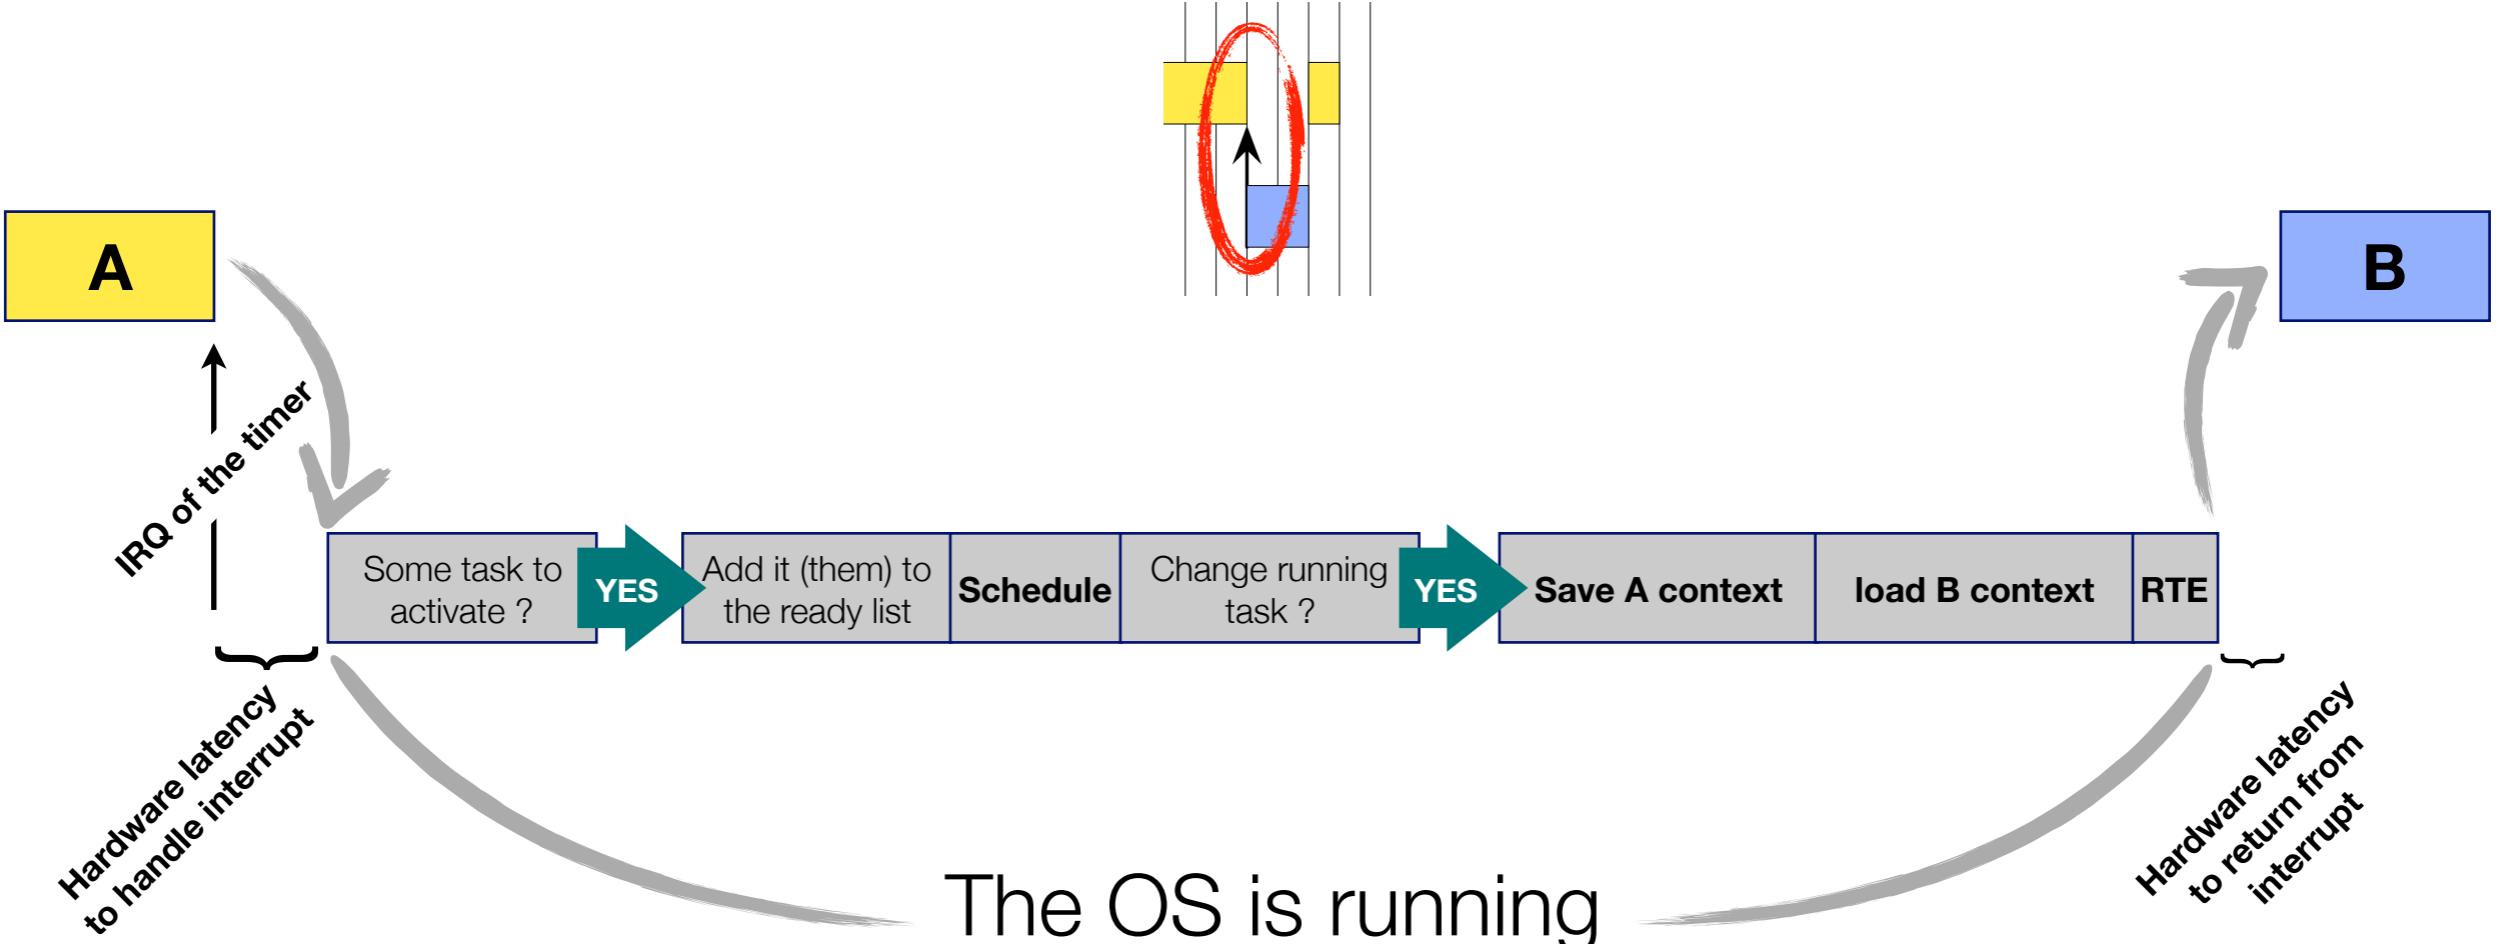
Interrupts (3)



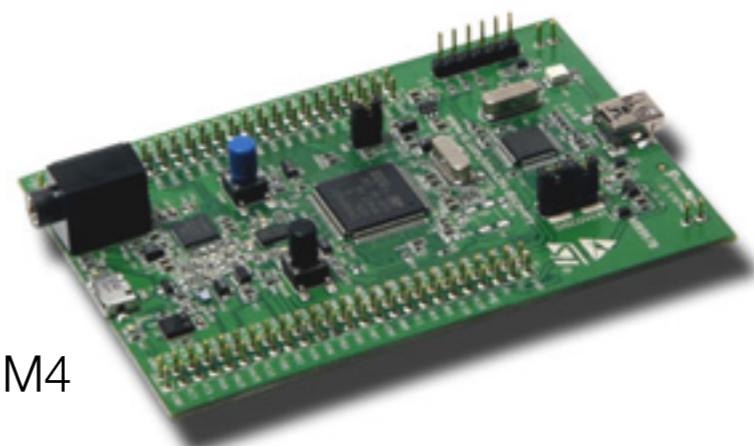
Interrupts and preemption



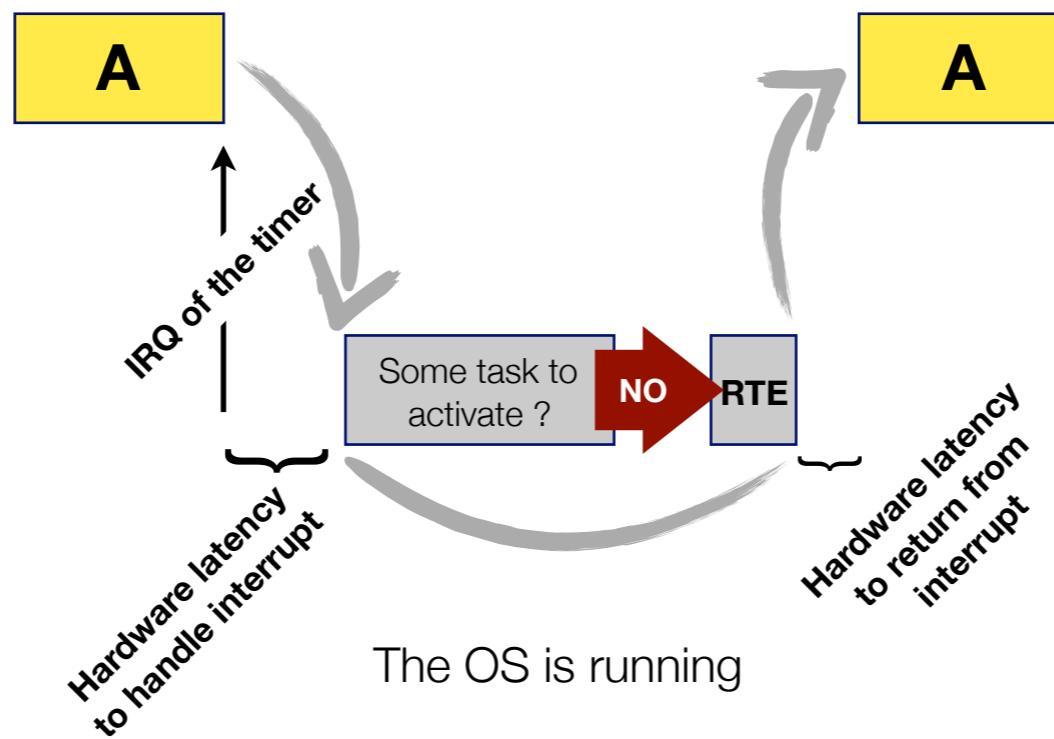
Preemption



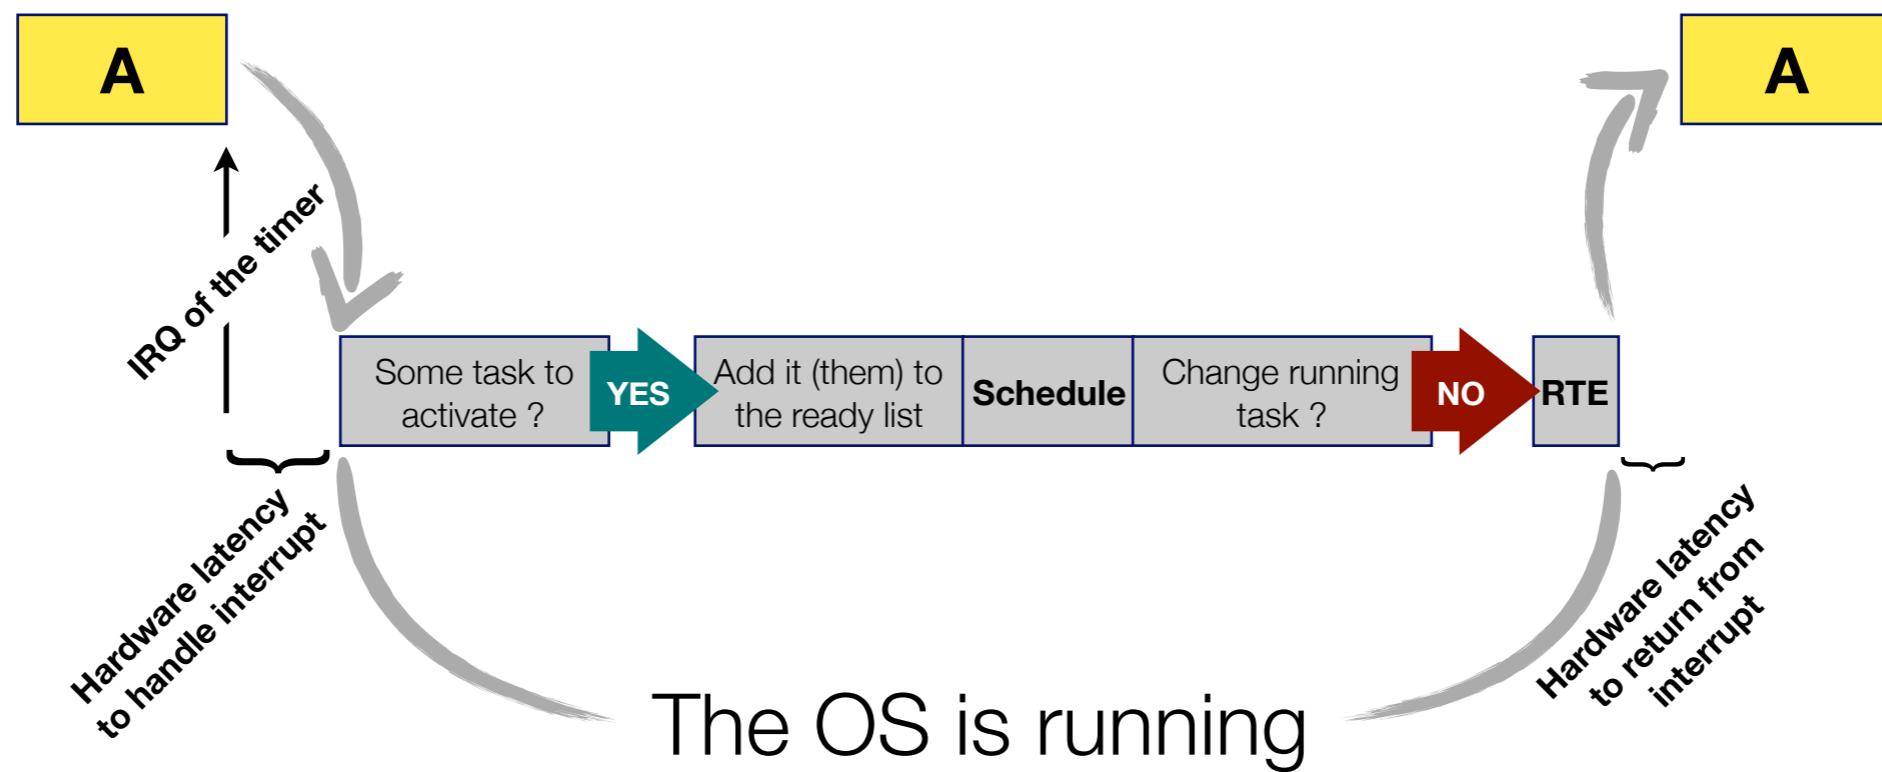
takes ~5µs on a 168MHz ARM Cortex M4



No preemption (1)



No preemption (2)



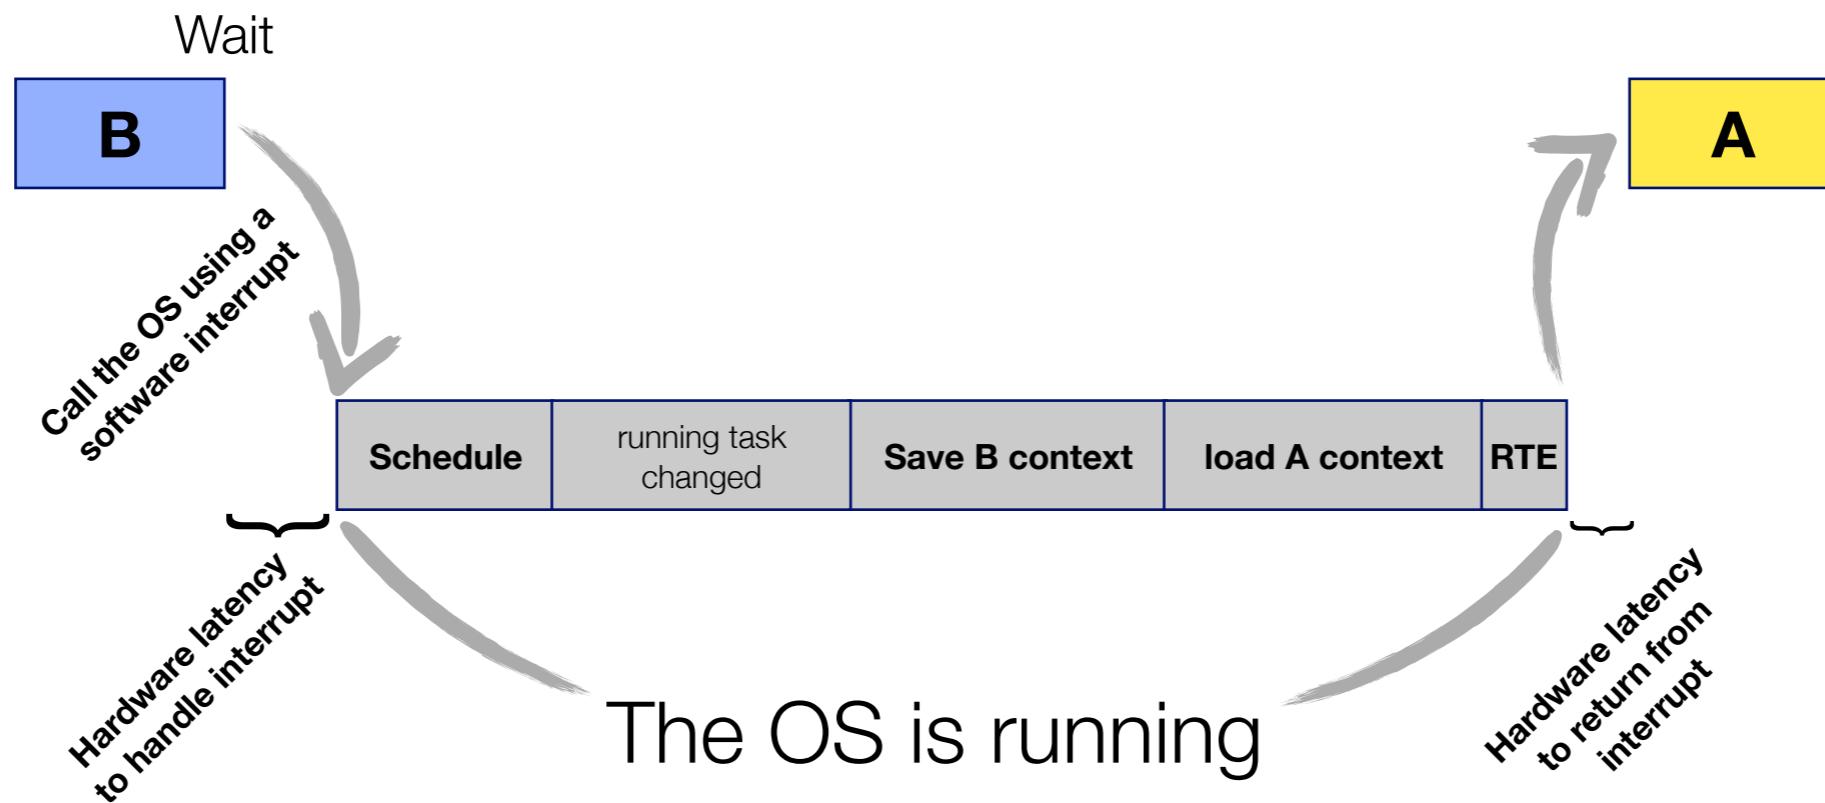
There are not only interrupts

- We saw the OS runs each time an interrupt dealing with tasks occurs but there are other occurrences that do not correspond to hardware interrupts: System Calls or System Services → functions of the OS that can be called by the tasks.

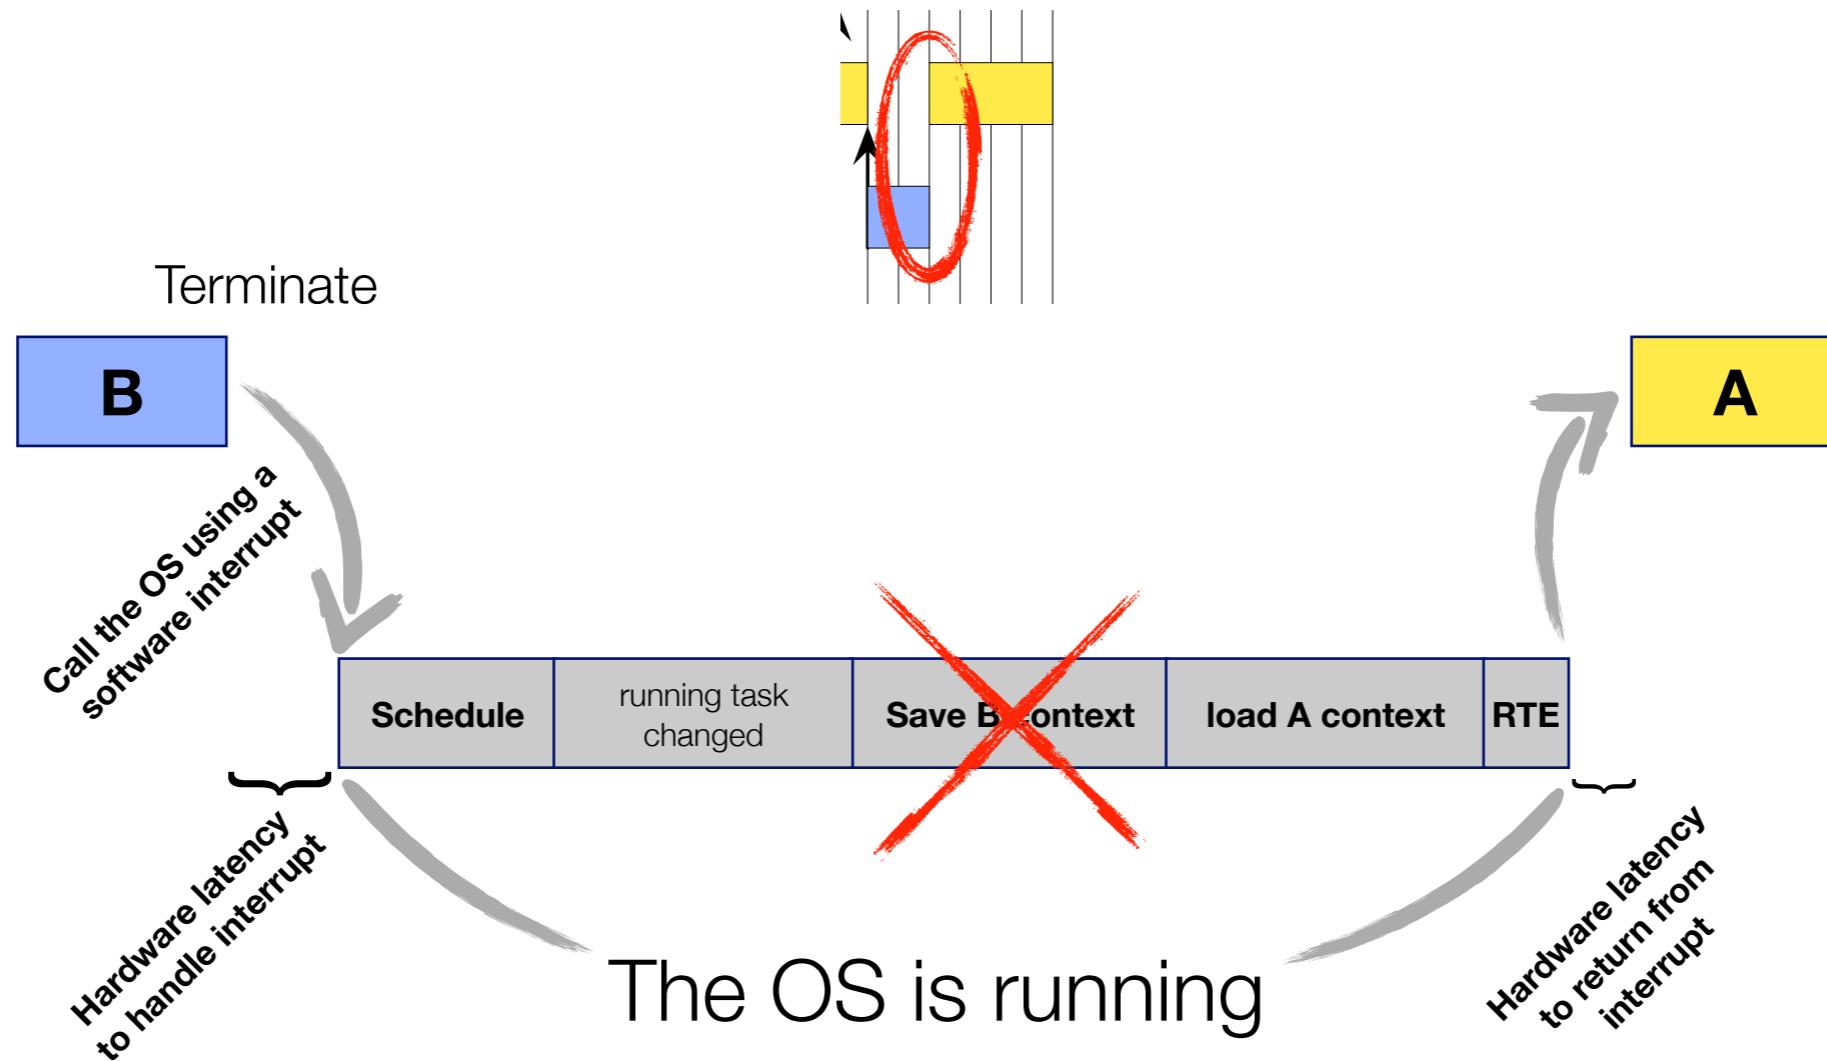
Examples:

- A task waits
- A task terminates
- A task activates another task (similar to activation by a hardware interrupt)
- ...

A task waits



A task terminates



Resource sharing and Synchronization

Synchronization purpose

- Tasks may need to synchronize to
 - synchronize the execution of (part of) code to a particular event
 - manage concurrent access to a shared resource
- Example: shared access to a liquid crystal display



Resource sharing / LCD example (1)

- Access to a LCD is a communication, usually a parallel 8 bits port (using the GPIO). Usually the chip which drives the screen is a Hitachi HD44780
 - Putting a character on the LCD is a single write
 - Clearing / setting the cursor position is a single write too
- 2 periodic tasks (t_1 and t_2 , $\text{prio}(t_2) > \text{prio}(t_1)$ so that t_2 can preempt t_1):
 - t_1 sets the cursor at column 0, line 0 and writes « MASTER »
 - t_2 sets the cursor at column 0, line 1 and writes « WET »

Does it work ?

Resource sharing / LCD example (2)

Arduino like code using LiquidCrystal

```
TASK(t1)
{
    lcd.setCursor(0,0);
    lcd.print("MASTER");
    TerminateTask();
}
```

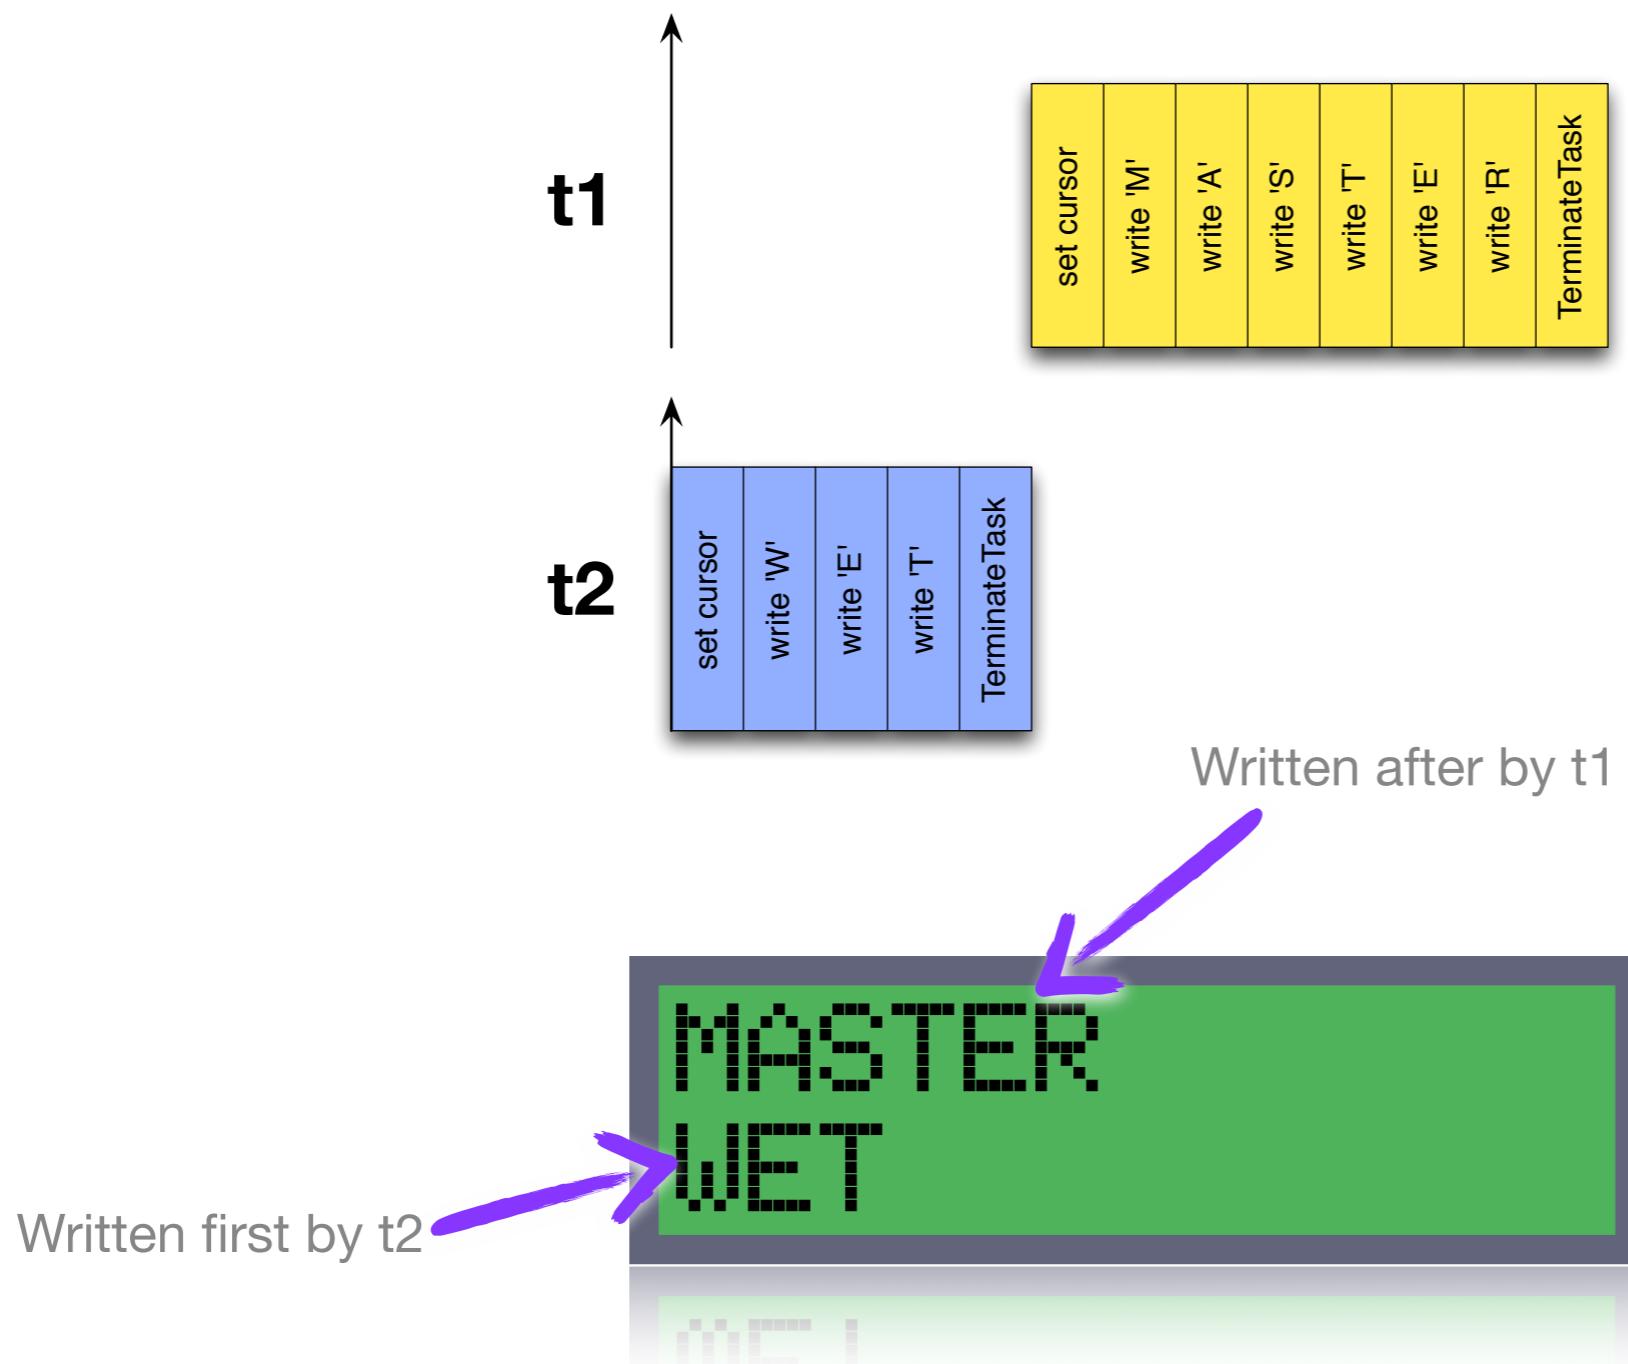
```
TASK(t2)
{
    lcd.setCursor(0,1);
    lcd.print("WET");
    TerminateTask();
}
```

Actually lcd.print writes chars one by one after getting

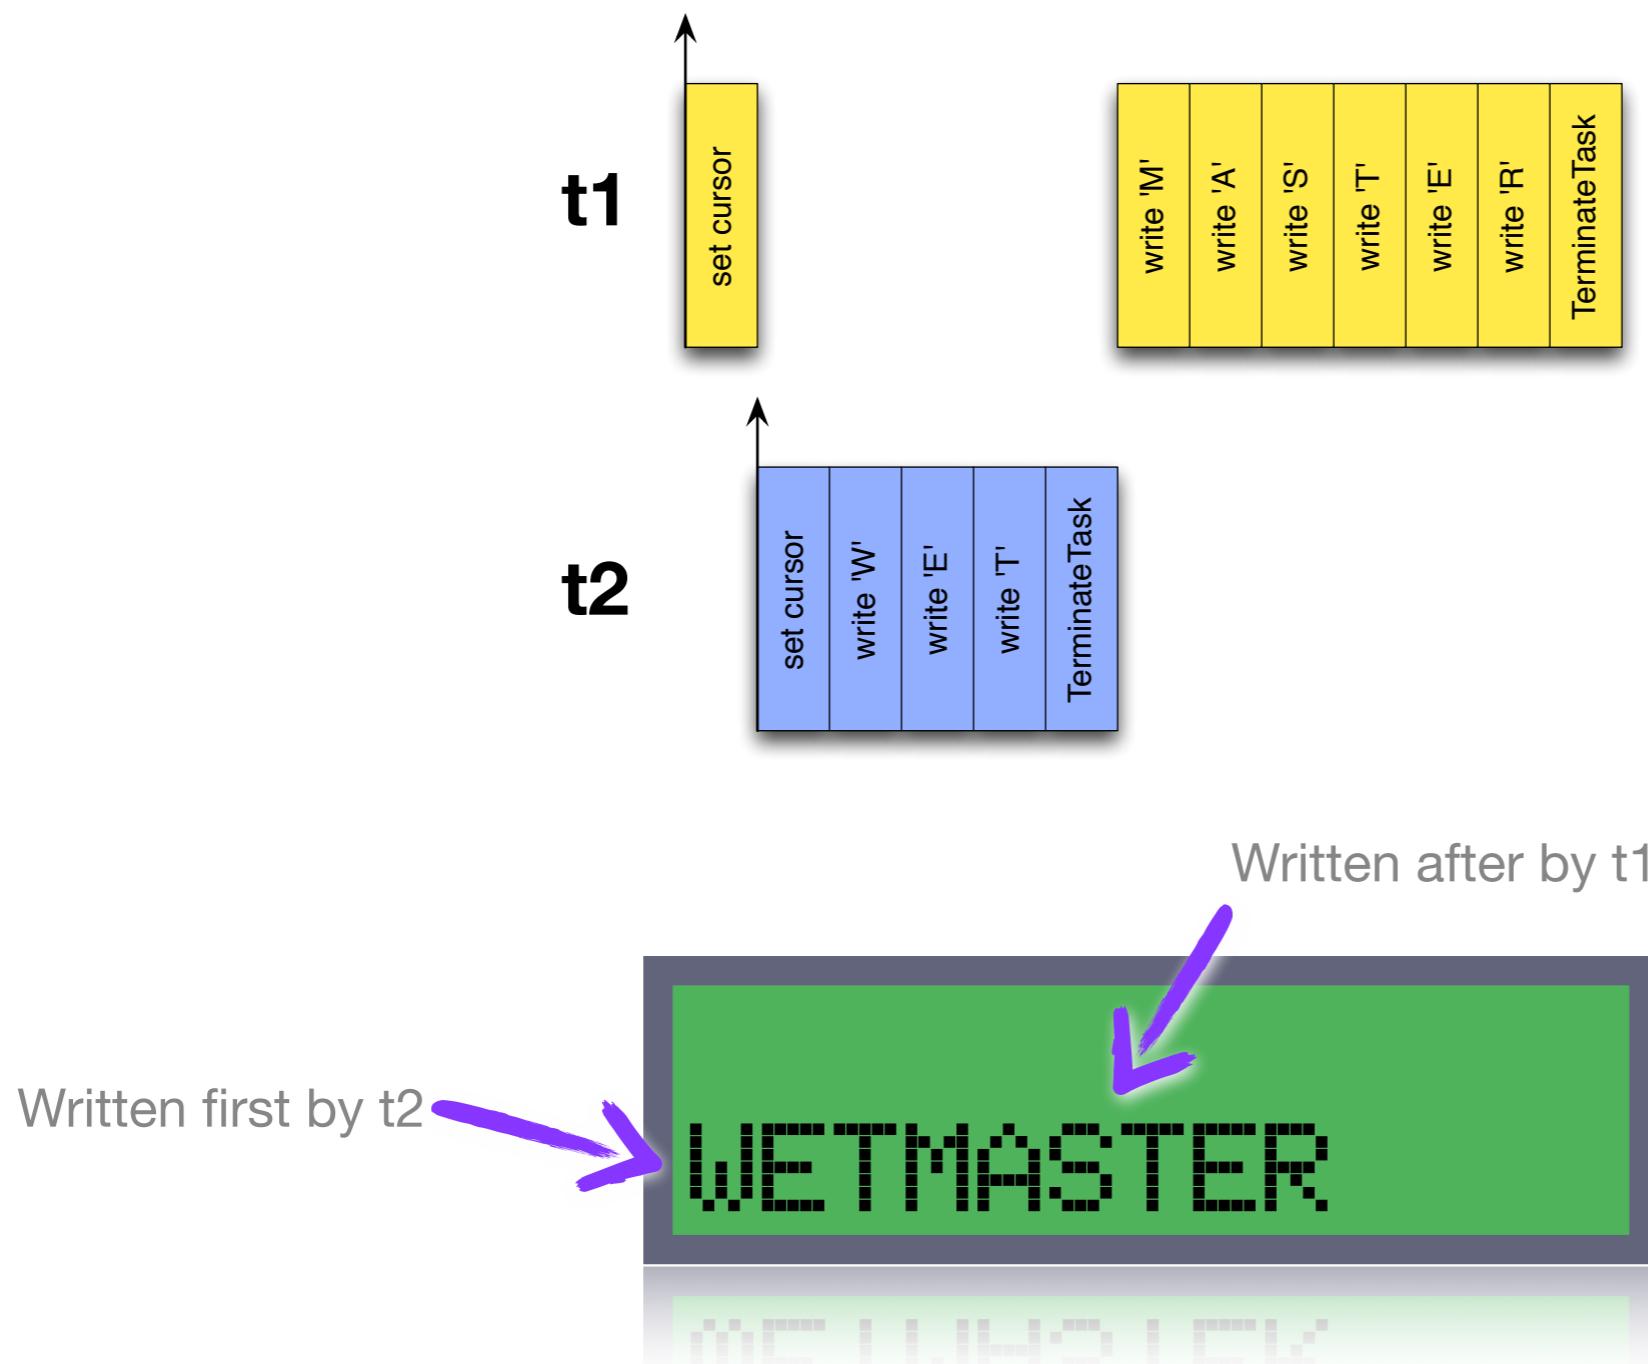
- the length of the string in the size variable
- a pointer to the string in the buffer variable

```
while (size--) {
    if (write(*buffer++)) n++;
    else break;
}
```

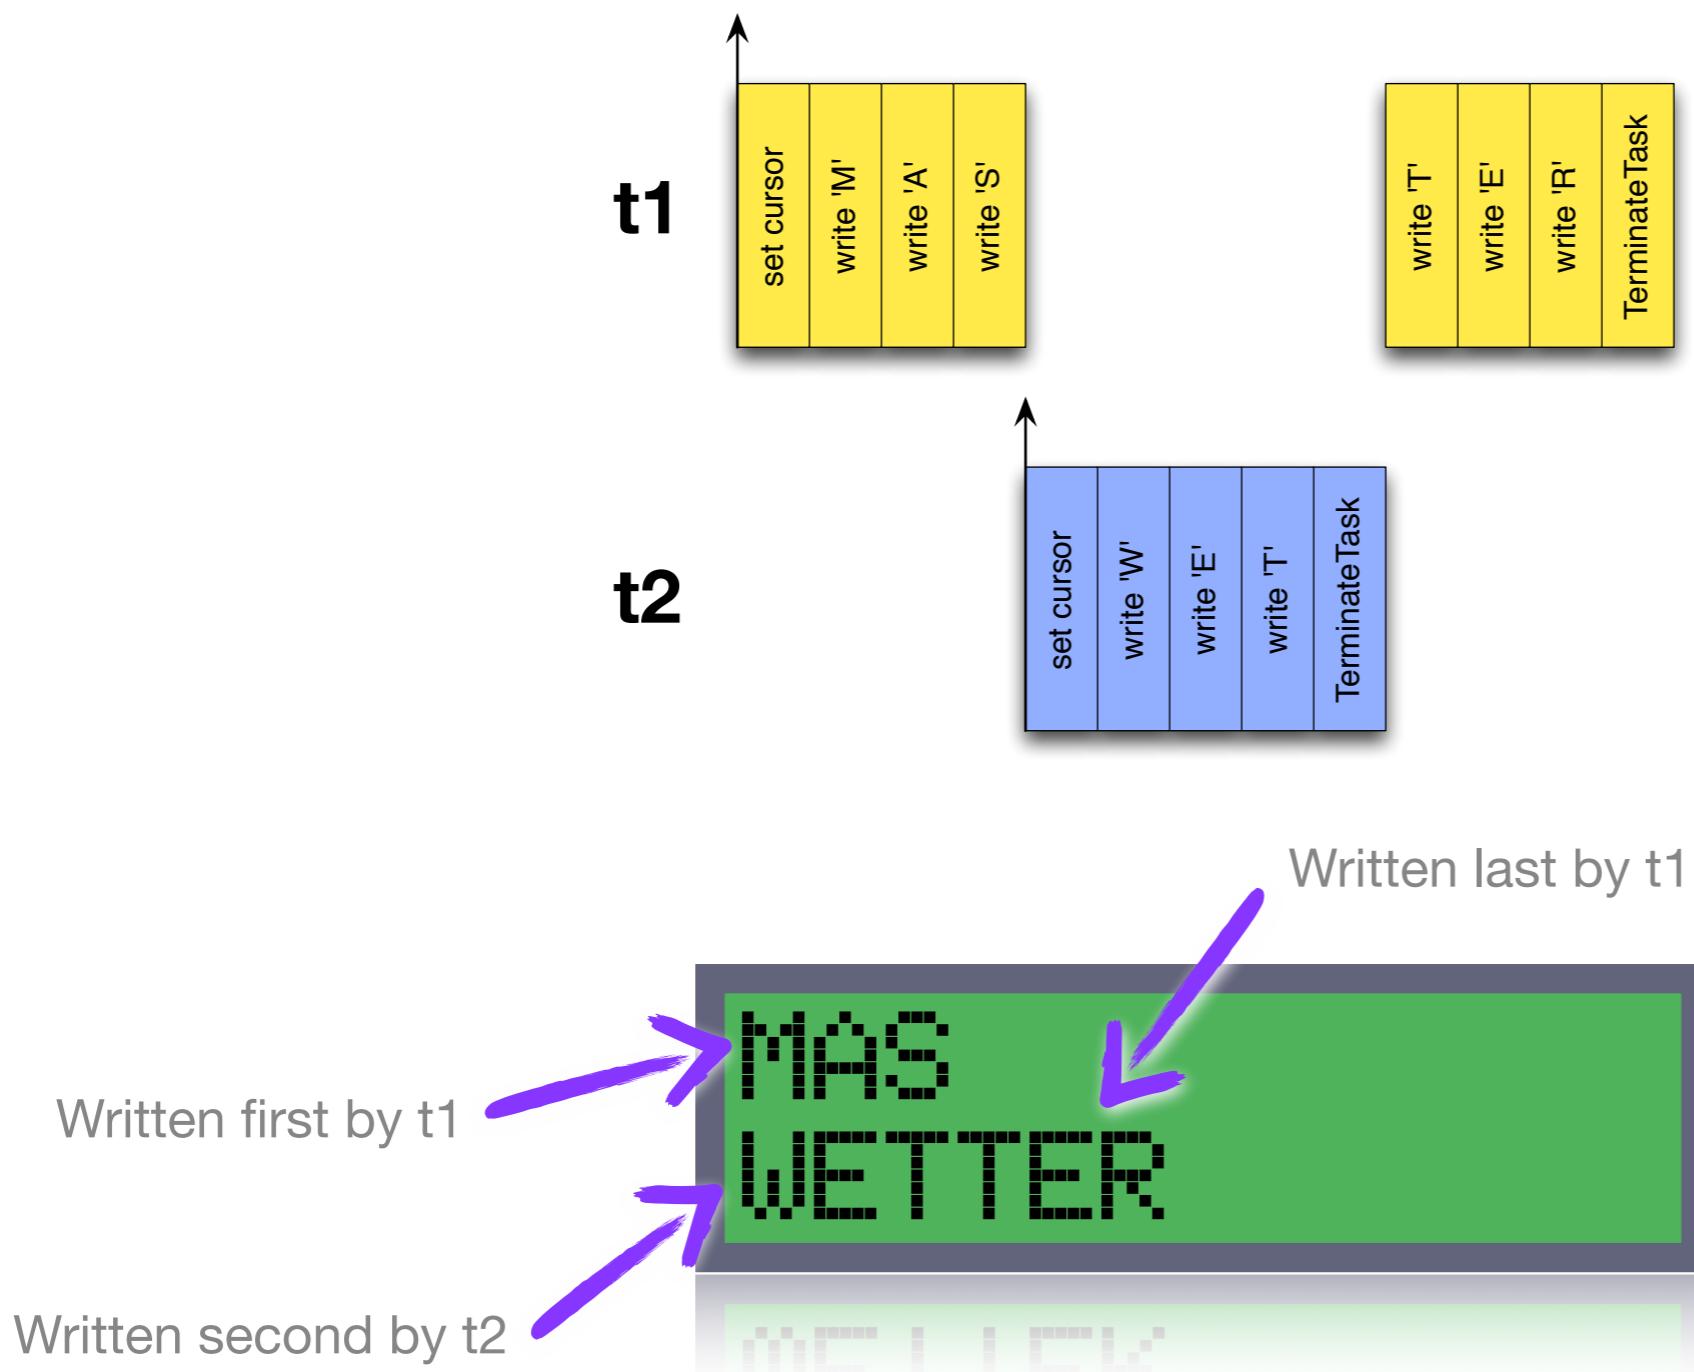
Resource sharing / LCD example (3)



Resource sharing / LCD example (4)



Resource sharing / LCD example (5)



Resource sharing / variable example (1)

```
volatile int i = 0;
```

```
TASK(t1)
```

```
{  
    i++;  
    TerminateTask();  
}
```

```
TASK(t2)
```

```
{  
    i++;  
    TerminateTask();  
}
```

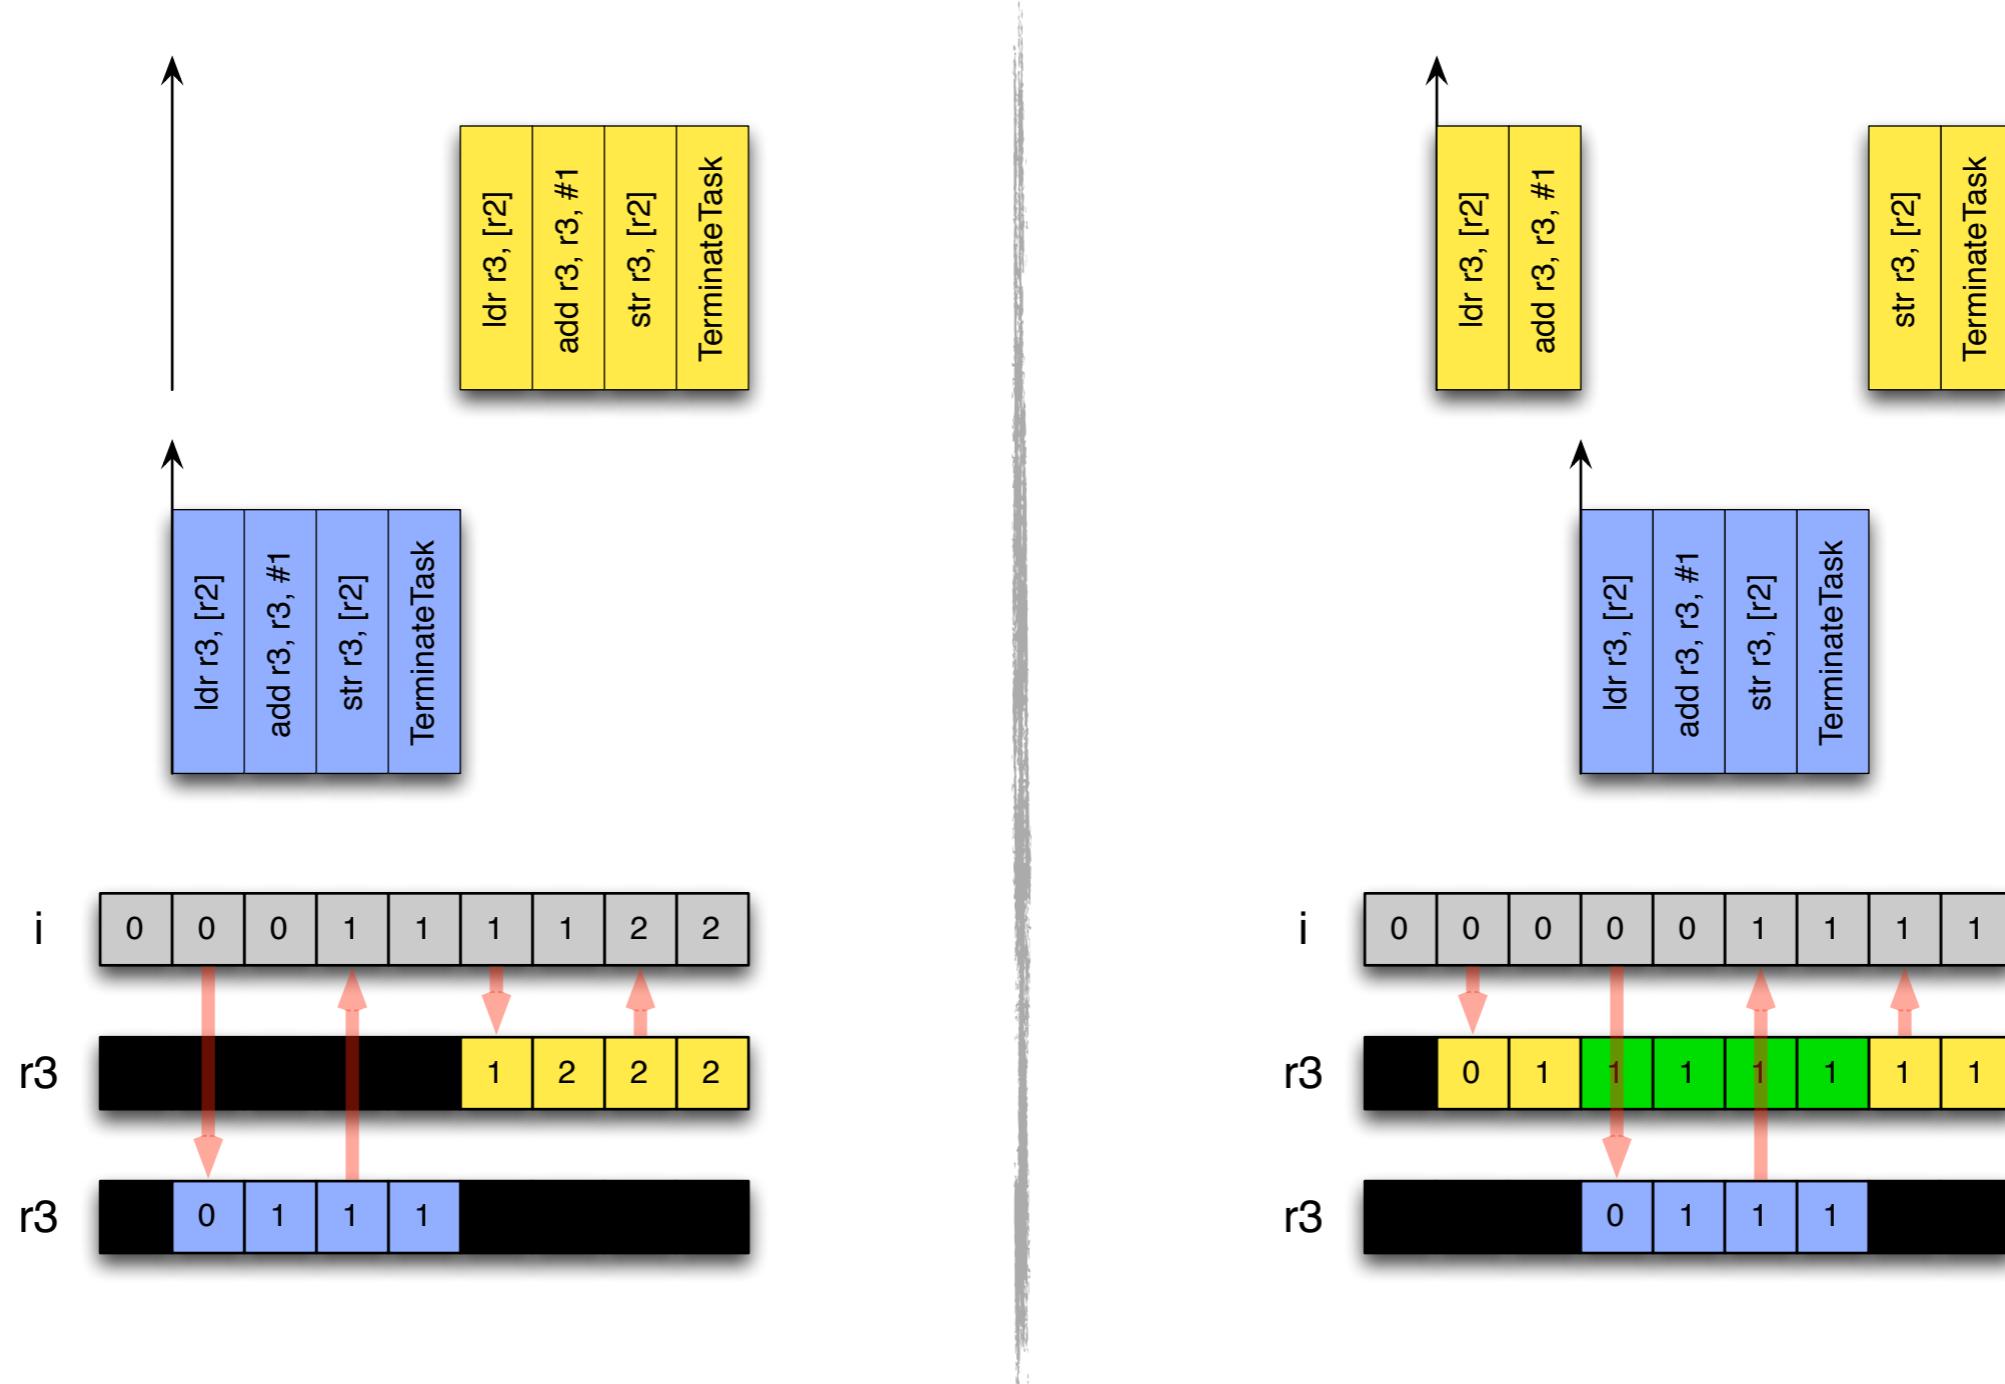
```
ldr r3, [r2]  
add r3, r3, #1  
str r3, [r2]
```

Actually i++ is not a single instruction of the processor

On ARM CORTEX,
assuming r2 contains
a pointer to i

prio(t2) > prio(t1)

Resource sharing / variable example (2)

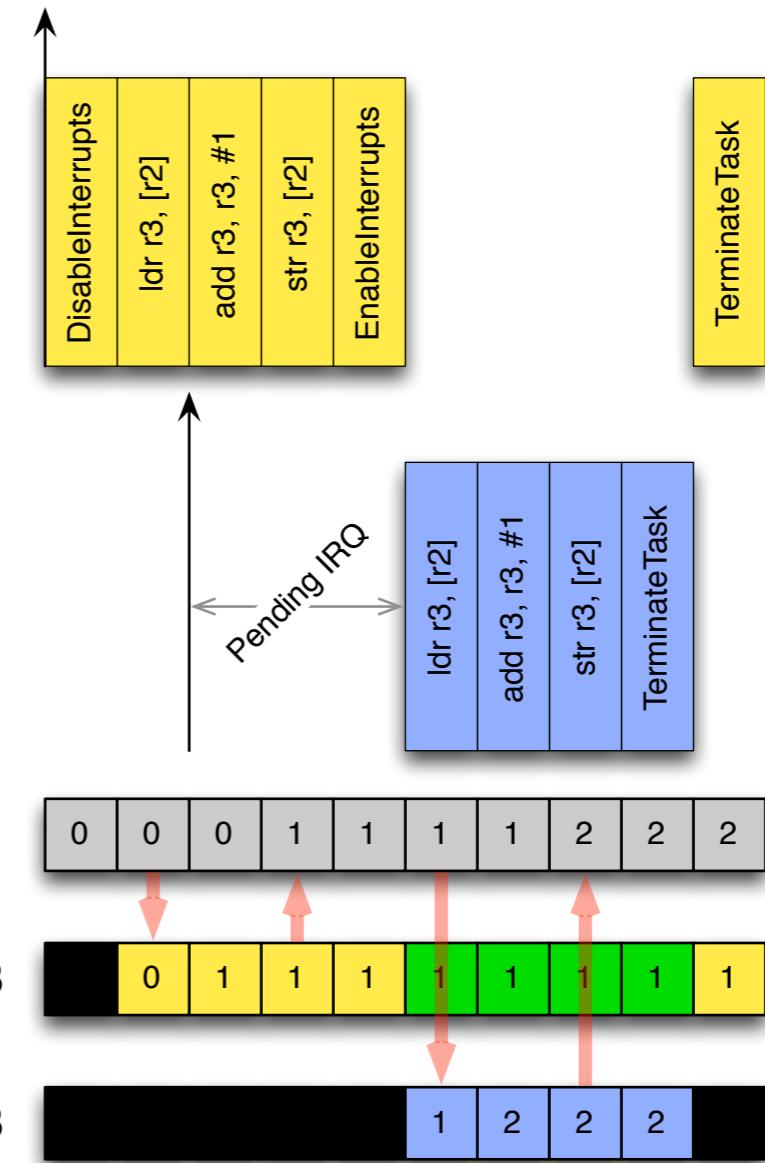


Resource sharing / no interrupts

```
volatile int i = 0;

TASK(t1)
{
    DisableInterrupts();
    i++;
    EnableInterrupts();
    TerminateTask();
}

TASK(t2)
{
    i++;
    TerminateTask();
}
```



Resource sharing / no interrupts

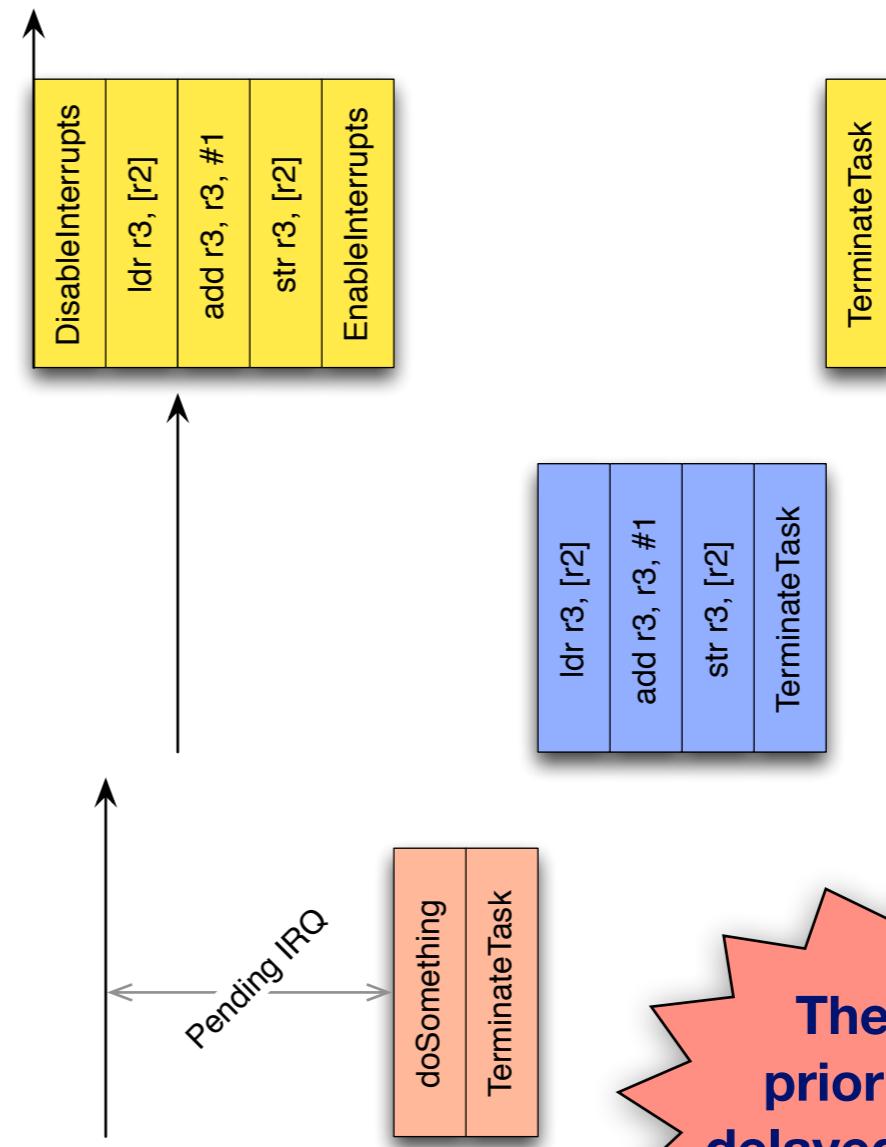
```
volatile int i = 0;

TASK(t1)
{
    DisableInterrupts();
    i++;
    EnableInterrupts();
    TerminateTask();
}

TASK(t2)
{
    i++;
    TerminateTask();
}

TASK(unrelated)
{
    doSomething();
    TerminateTask();
}
```

prio(unrelated) > prio(t2) > prio(t1)



The highest priority task is delayed whereas it does not use i

Resource sharing / how ?

- What we want is a mutual exclusion mechanism such that 2 or more tasks cannot execute code accessing/modifying a shared resource (device, variable) at the same time.
 - A **critical section** is an abstract portion of code delimited by proper services calls which reference the critical section (using an ID for instance).
 - more than one critical section, each with an id, can be defined for each of the shared resource.
 - When a task is within a critical section, another task cannot enter the same critical section.
 - Viewed from the other task, all the code executing within the critical section is **atomic**.
 - This rely on **atomic hardware mechanism**: Since service calls are executing with interrupts masked, it is ok (at least in monocore).

A simple mutex

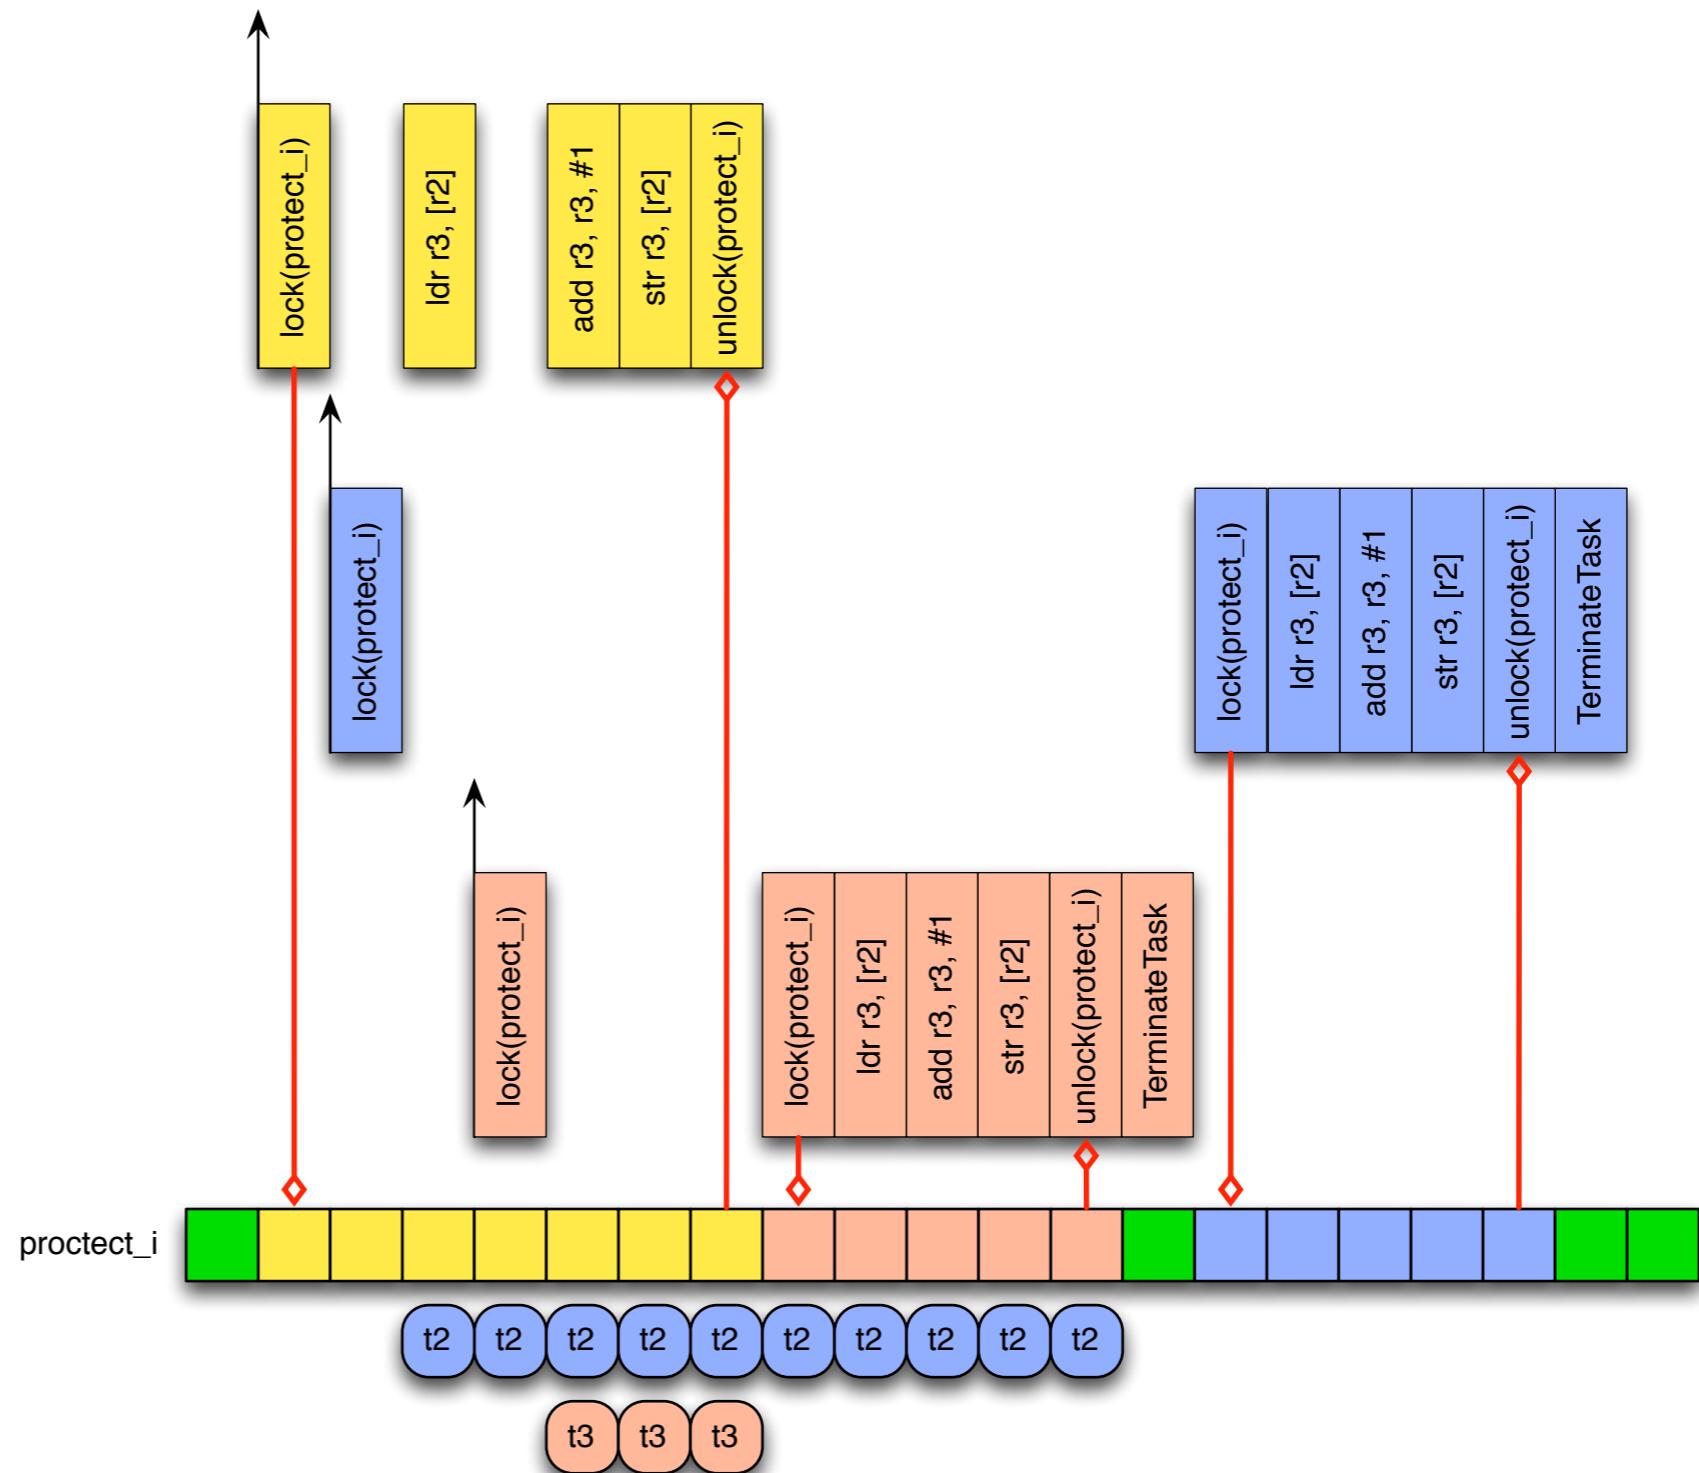
- Let's say we have a special datatype called a mutex (for MUTual EXclusion)
 - `mutex m;` defines a mutex called **m**
 - `lock(m);` locks m, it should be used at the start of a critical section
 - `unlock(m);` unlocks m, it should be used at the end of a critical section
 - any task calling `lock` on a mutex already locked is put in the **waiting state**
 - when a mutex is unlocked, one or more tasks waiting for the mutex are put back in the **ready state** and a rescheduling occurs

A simple mutex : example

```
mutex protect_i;  
  
TASK(t1)  
{  
    lock(protect_i);  
    i++;  
    unlock(protect_i);  
    TerminateTask();  
}
```

```
    TASK(t2)
    {
        lock(protect_i);
        i++;
        unlock(protect_i);
        TerminateTask();
    }
```

```
    TASK(t3)
    {
        lock(protect_i);
        i++;
        unlock(protect_i);
        TerminateTask();
    }
```



$\text{prio}(t_3) > \text{prio}(t_2) > \text{prio}(t_1)$

Semaphore

- Proposed by Edger Dijkstra
- It allows to protect access to shared resources
- This mechanism, available in some OS Posix compliant for instance) offers 3 functions:
 - Init() : initialize the semaphore;
 - P() to test the semaphore (Probieren);
 - V() to increment the semaphore (Verhogen).
- Works like mutex but with a counter. Identical to a mutex if the counter is initialized to 1.

Semaphore

- A counter is associated to the semaphore.
 - A call to P() is used to enter the critical section:
 - If the counter is > 0 , it is decremented and the resource may be taken.
 - If the counter is $= 0$, the task which called P() is put in the waiting state until the counter became > 0 . At that time the task will be awoken and the counter will be decremented again.
 - A call to V() is used release a resource:
 - The counter is incremented and a task which is waiting may be put in ready state.

Semaphore

- The counter associated to the semaphore may have a non-binary value
- example: access to a buffer:
 - Here are 2 functions to read and write the buffer. These functions may be called by concurrent tasks.
 - The S1 semaphore which has the initial value of its counter equal to 5 allows up to 5 writes. After that, a task which try to write is put in the waiting state until a task does a read.

```
void init() {  
    semInit(S1, 5);  
}  
  
void WriteBuffer(int data) {  
    P(S1);  
    //buffer write  
}  
  
void ReadBuffer(int *data) {  
    V(S1);  
    //buffer read  
}
```

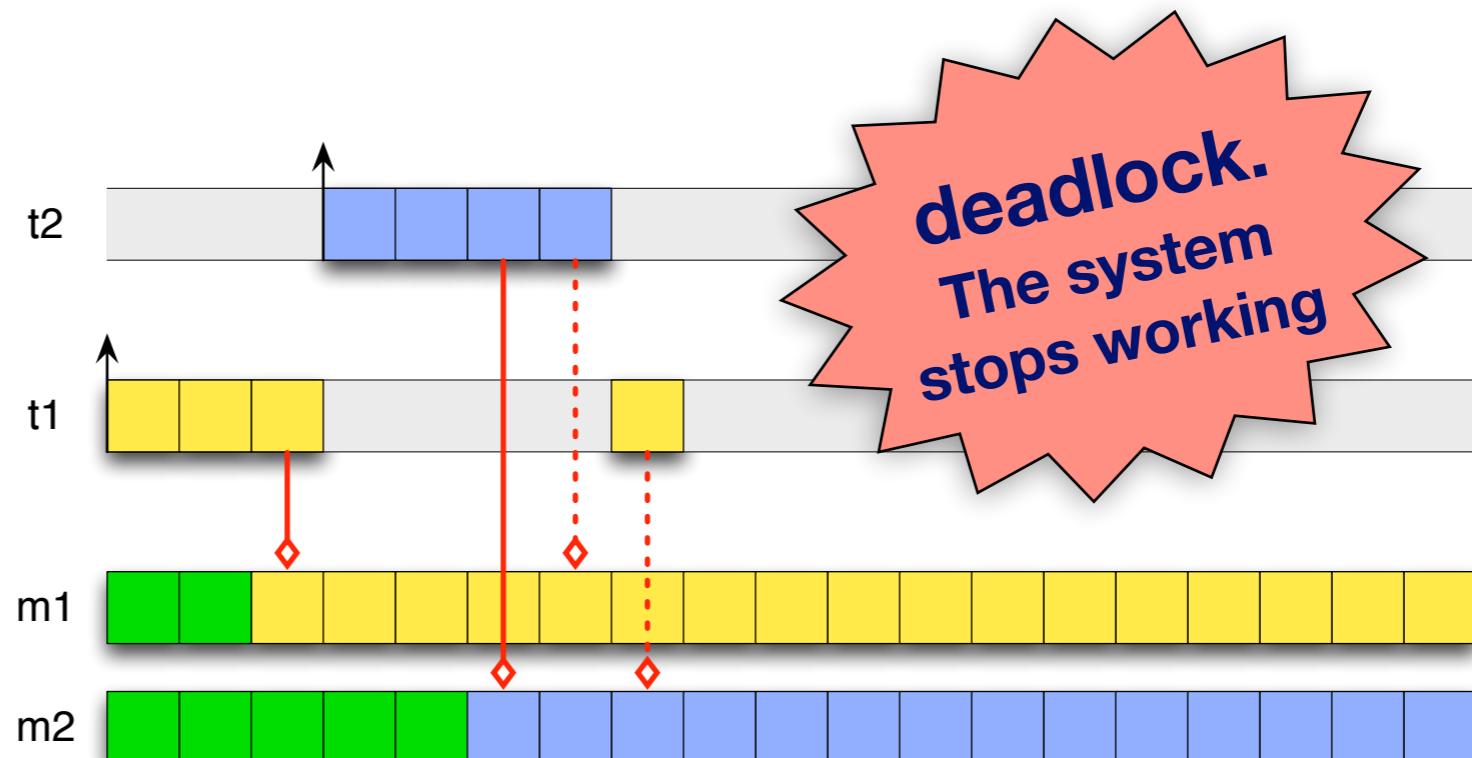
Problem: deadlock

- Classical synchronization mechanisms (semaphore, mutex) may have deadlocks:
 - example: 2 tasks take 2 mutex in reverse order, $\text{prio}(t2) > \text{prio}(t1)$

```
mutex m1;
mutex m2;
volatile int i;
volatile int j;

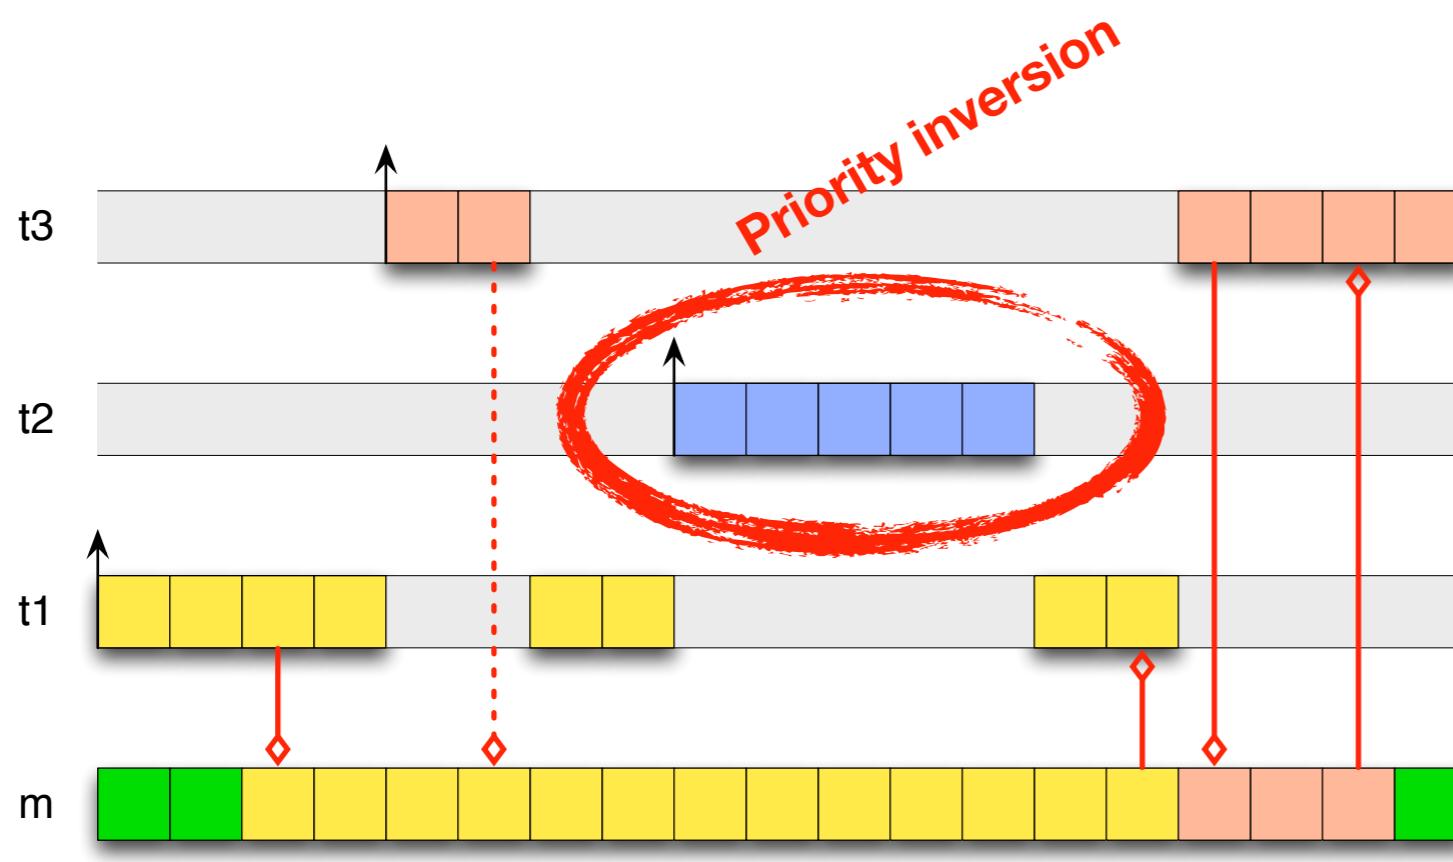
TASK(t1)
{
    lock(m1);
    lock(m2);
    i++;
    j--;
    unlock(m2);
    unlock(m1);
    TerminateTask();
}

TASK(t2)
{
    lock(m2);
    lock(m1);
    i--;
    j++;
    unlock(m1);
    unlock(m2);
    TerminateTask();
}
```

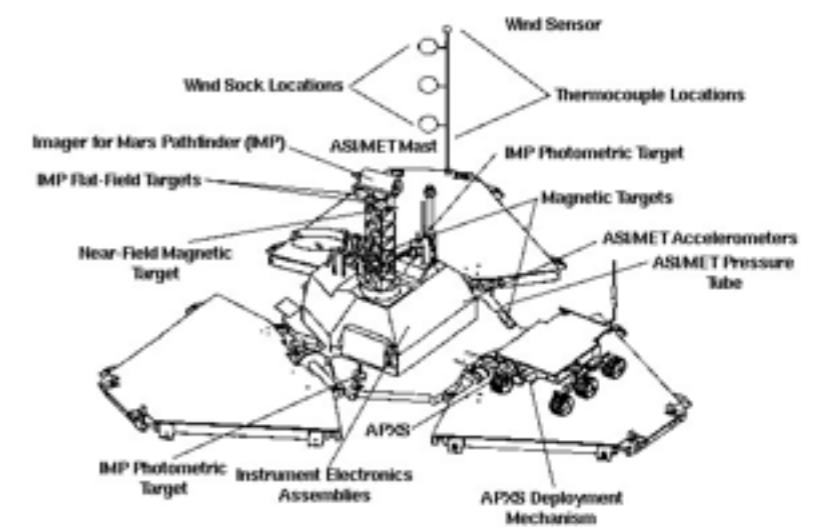


Problem: priority inversion

- Classical synchronization mechanisms (semaphore, mutex) may have the priority inversion problem:
 - A task with a lower priority may delay a higher priority task.
 - The following example shows 3 preemptable tasks.
 $\text{prio}(t3) > \text{prio}(t2) > \text{prio}(t1)$:



**Caused the failure of
Mars Pathfinder Probe**



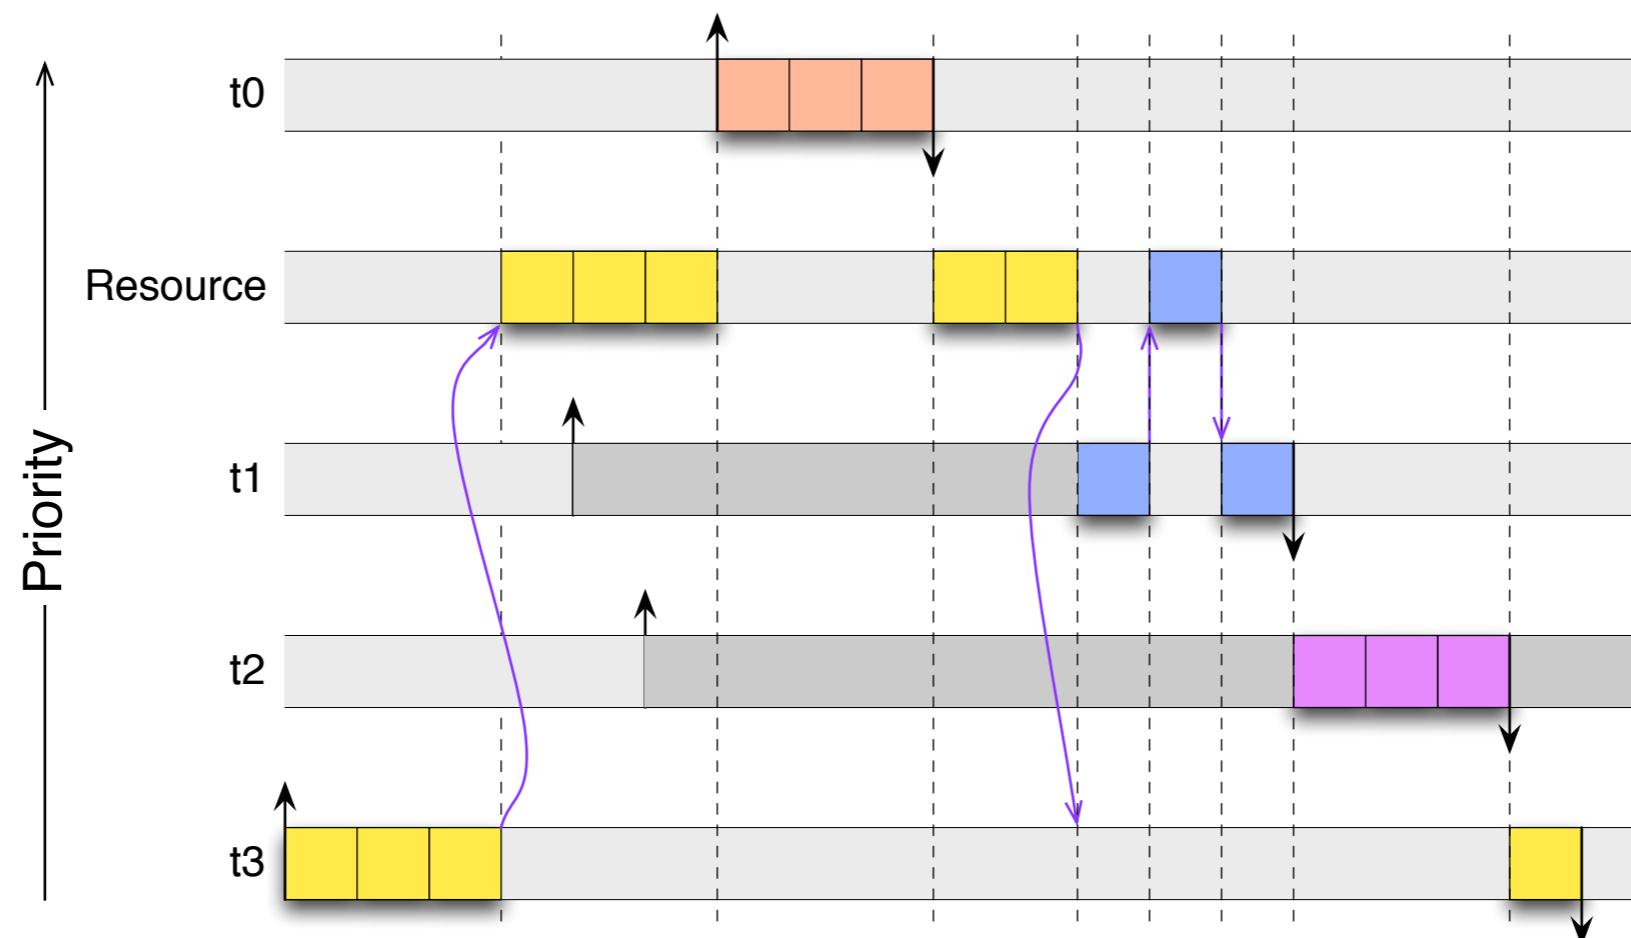
IPCP: Immediate Priority Ceiling Protocol

- Priority based protocol
 - works with fixed priority scheduling
 - the mutex (usually called a Resource) has a priority
 - the priority is the max of the priorities of tasks using the Resource
 - when a task takes the resource, **its priority is raised to the priority of the Resource** (if needed) ⇒ other tasks likely to take the Resource **cannot preempt** the task that holds the Resource
 - when the task release the Resource, **its priority is lowered to the previous priority** (if needed) and a **rescheduling occurs**.
- Deadlocks are not possible because resources are in strict ordering (the priority)
- Priority inversion is not possible because the priority changes



IPCP example

- 4 tasks
 - $\text{prio}(t_0) > \text{prio}(\text{Resource}) \geq \text{prio}(t_1) > \text{prio}(t_2) > \text{prio}(t_3)$
 - t_1 and t_3 share a Resource



Events

- Events are a synchronization scheme.
 - An event is a flag, a boolean variable, used to signal something happened
 - It is a N producers / 1 consumer model (private events) or N producers / M consumers model (public events).
 - A set of events is called an event mask
- Not designed for critical sections
- A task may
 - set an event
 - clear an event
 - wait for an event
 - get the current events

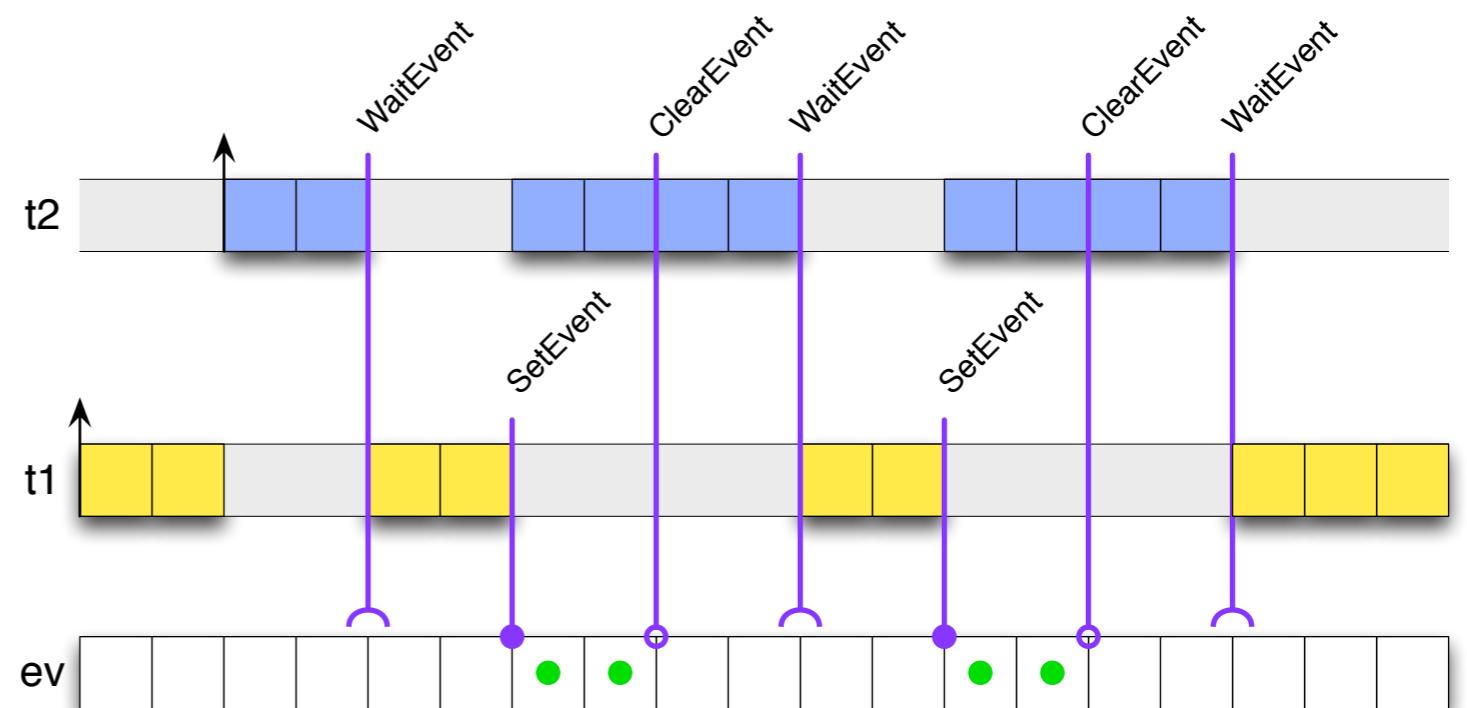
Events : example

- $\text{prio}(t2) > \text{prio}(t1)$

```
Event ev;
```

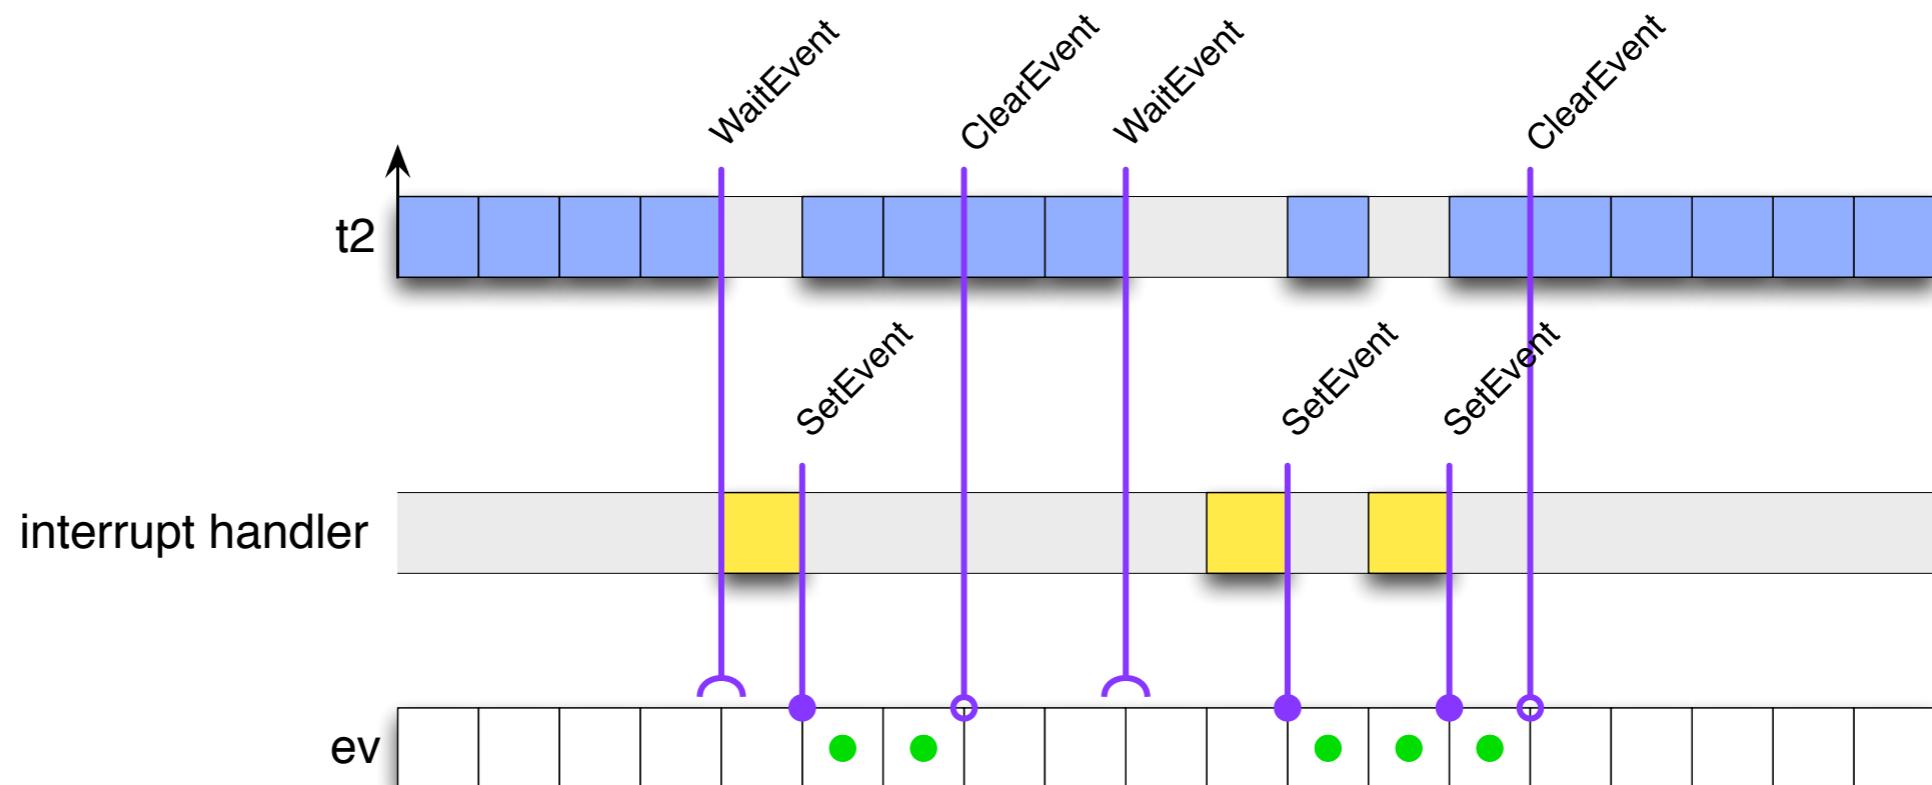
```
TASK(t1)
{
    while(1) {
        doSomeWork();
        SetEvent(t2, ev);
    }
    TerminateTask();
}
```

```
TASK(t2)
{
    while(1) {
        WaitEvent(ev);
        ClearEvent(ev);
        doSomeOtherWork();
    }
    TerminateTask();
}
```



Events may be lost

- If a second event is received while the previous one has not been taken into account, the second event is lost.
- Fixed priority scheduling solve the problem in monoprocessor system, not in multiprocessor system or if event is set by interrupt handler.



Events : implementation

- To implement this feature, an event corresponds to a binary mask: 0x01, 0x02, 0x04, ...
- An event vector is associated to 1 or more bytes. Each event is represented by one bit in this vector
- So each task owns:
 - a vector of the events set
 - a vector of the event it waits for
- Operation:
 - Event X signaling : `ev_set |= mask_X;`
 - Is event X arrived ? : `ev_set & mask_X;`
 - Wait for event X : `ev_wait | mask_X;`
 - Clear event X : `ev_set &= ~mask_X;`

Communication : shared data and messages

Shared data

- A set of global data shared among 2 or more tasks
 - simplistic and primitive approach
 - concurrency problem (we saw that in the previous section)
- Global variables are evil and should be avoided
 - System wide visibility:
 - implicit coupling of every pieces of software
 - how to be sure they are protected everywhere ?
 - Hidden bugs difficult to find.

Messages (1)

- A better way to move data from one task to another
- Messages can be of **queued** kind:
 - A queue is a FIFO handling N messages
 - If a write occurs on a queue that is full, the message is lost
 - Of course, N can be 1
- Messages can be of blackboard kind:
 - A blackboard handles one message only
 - Writing a message to a blackboard overwrites the previous message
 - Reading a message from a blackboard gets the last written message
 - What if one reads a message from a blackboard that was never written ?

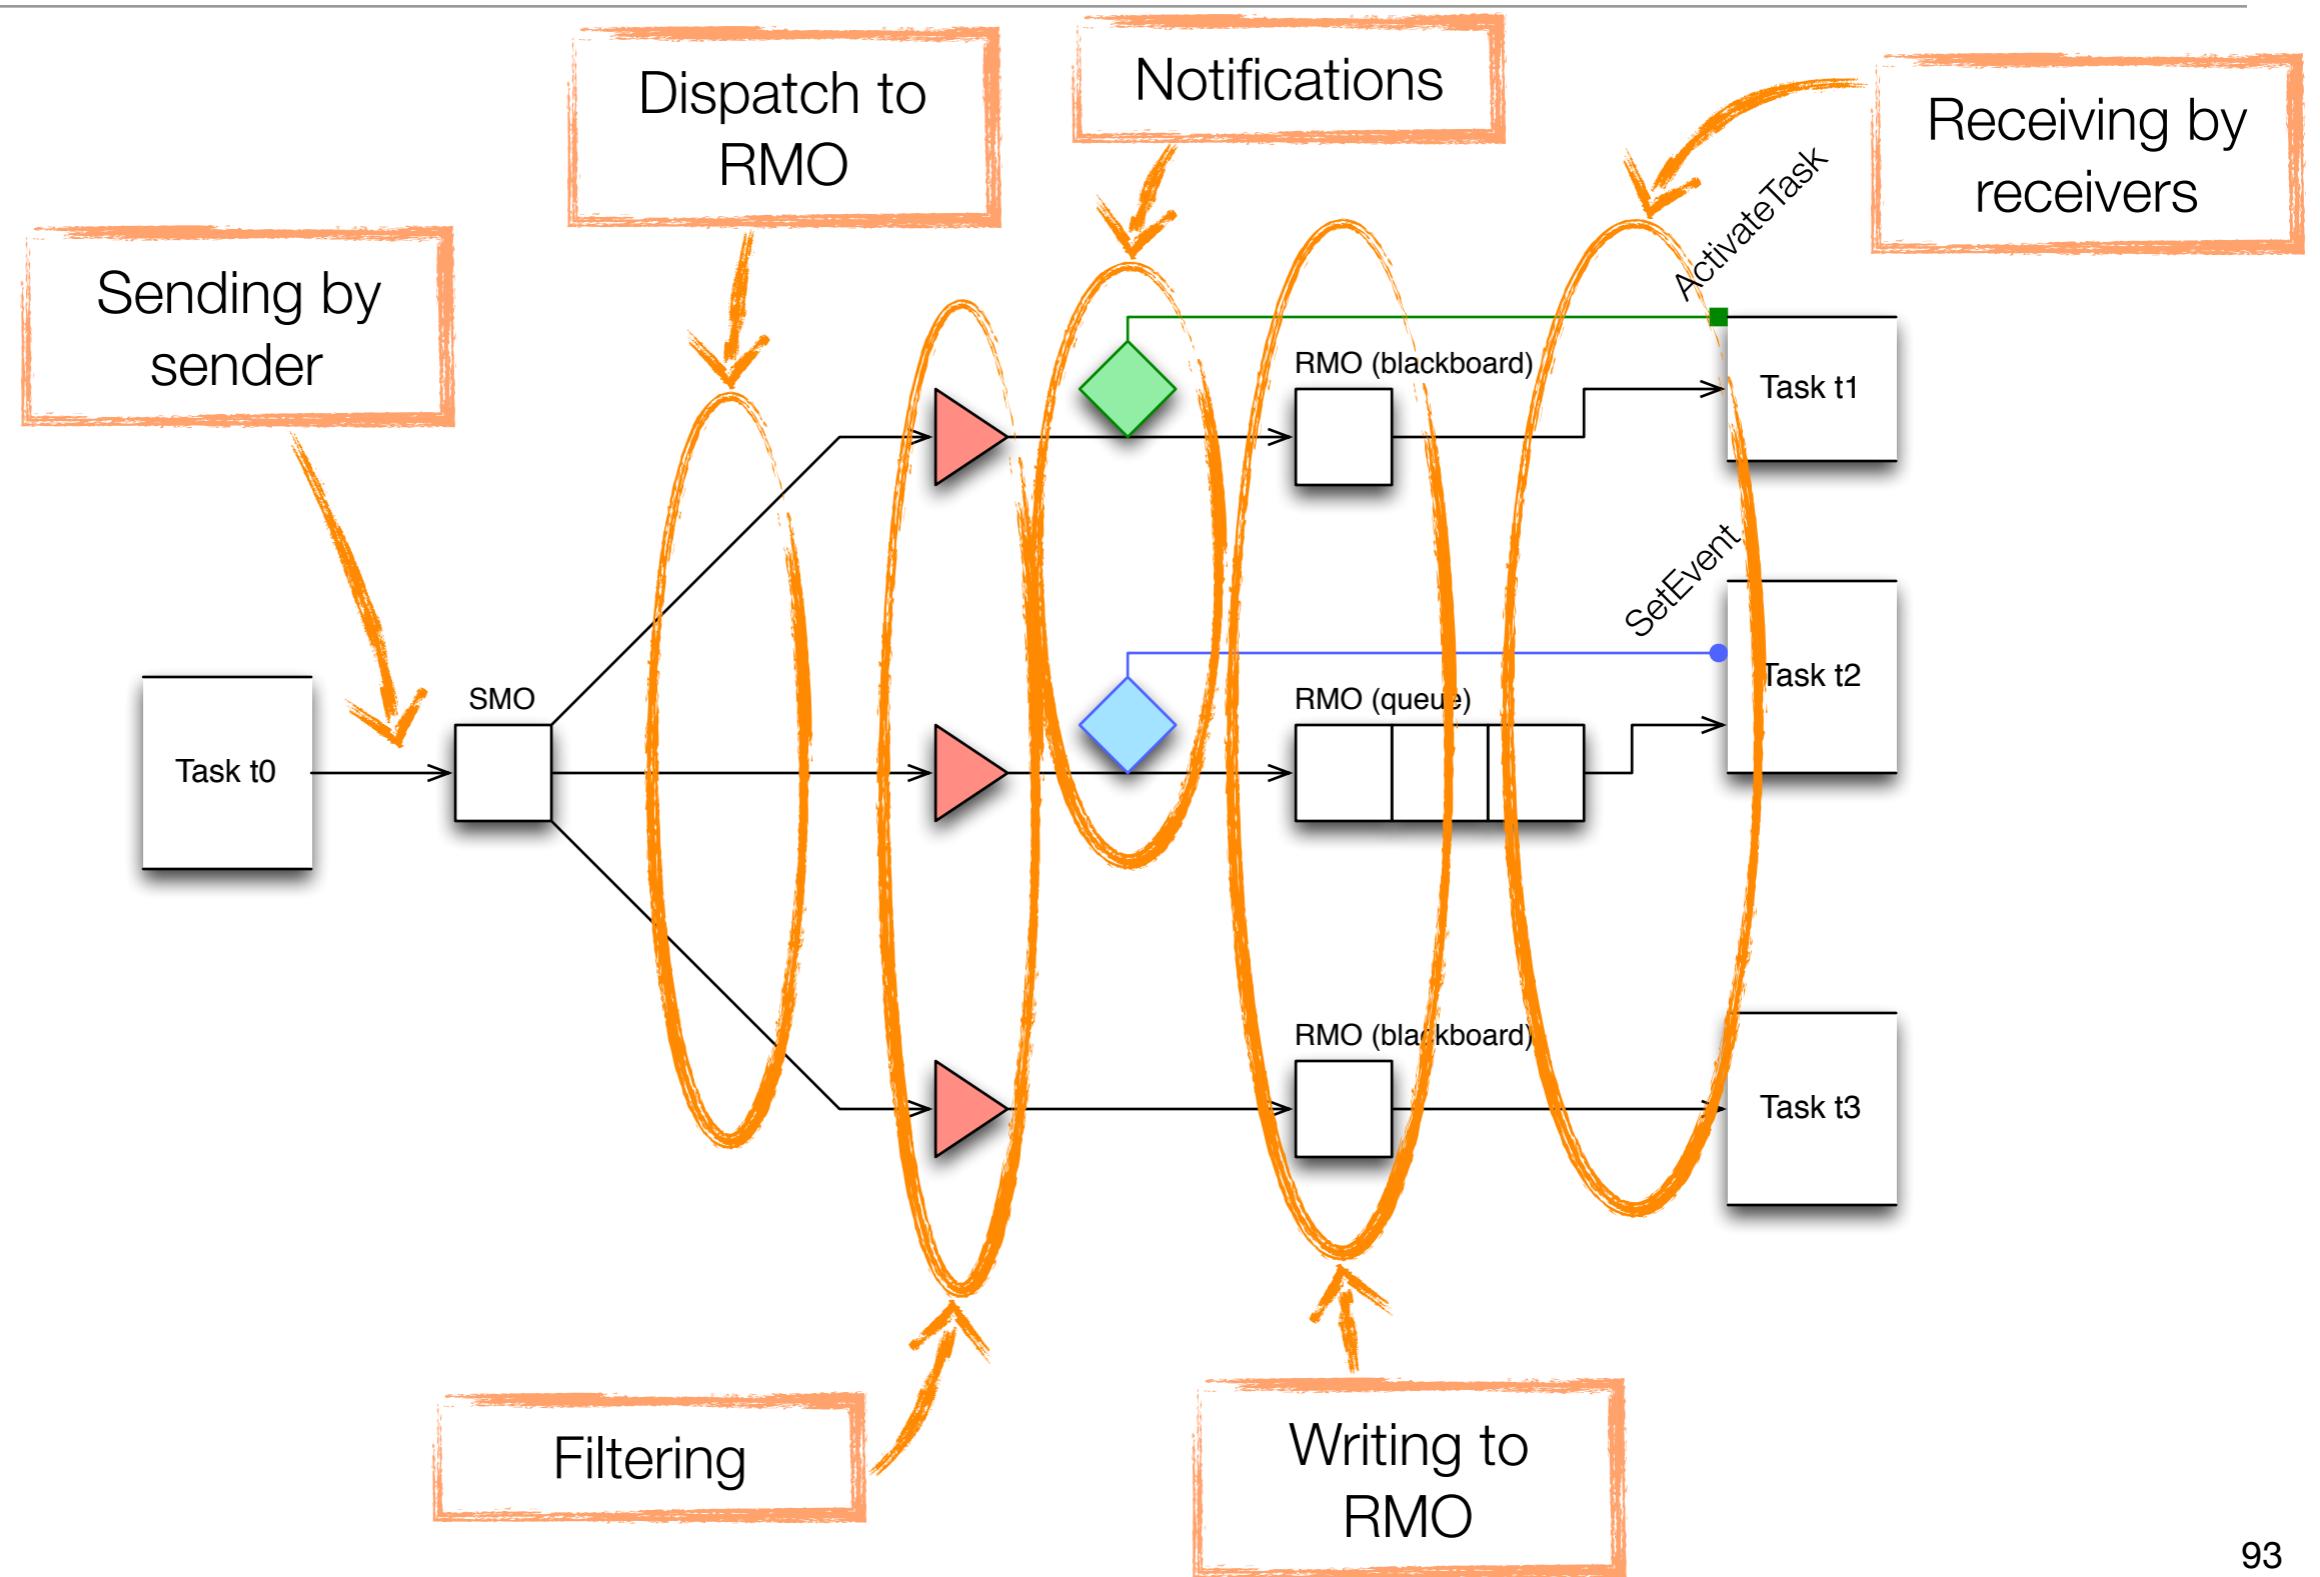
Messages (2)

- Concurrent accesses to a message is managed by the OS
 - **No need to use a mutex**, it is implicitly provides
- Synchronization may be included:
 - If a task try to read an empty queue, it is put in the WAITING state
 - If a message is received in an empty queue and a task is WAITING, then the task is put in the READY state
 - An event can be sent to a task when a message is received
 - A task may be activated when a message is received

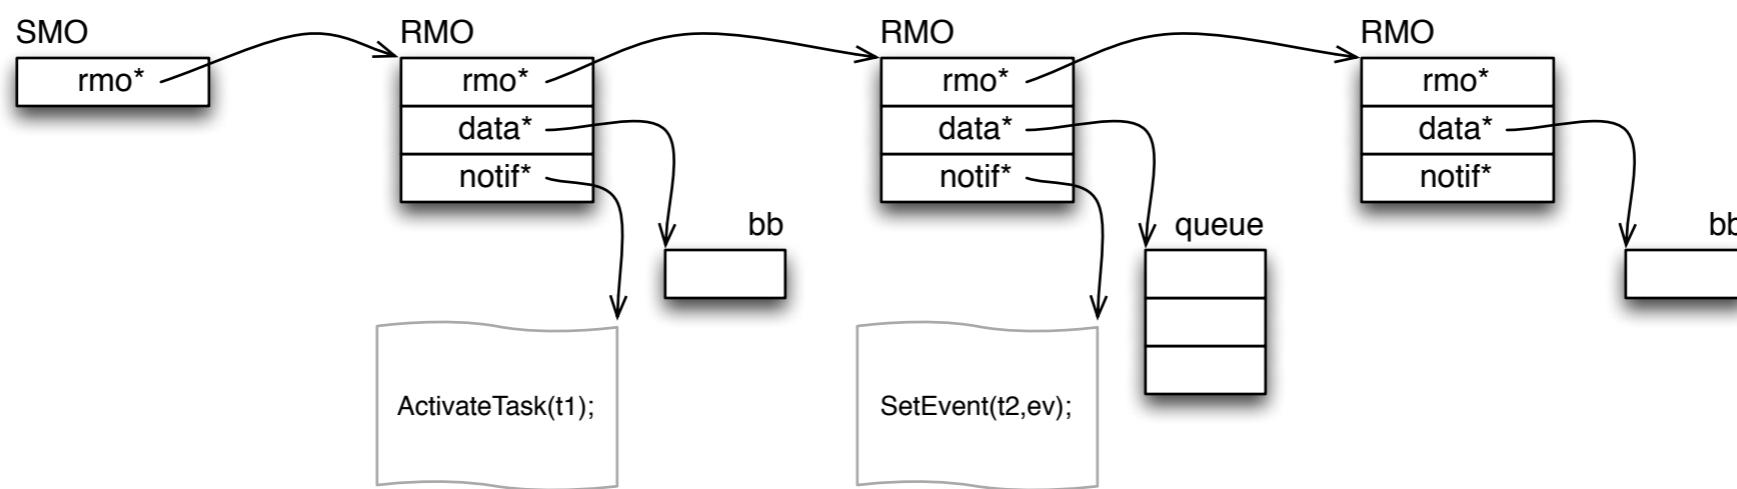
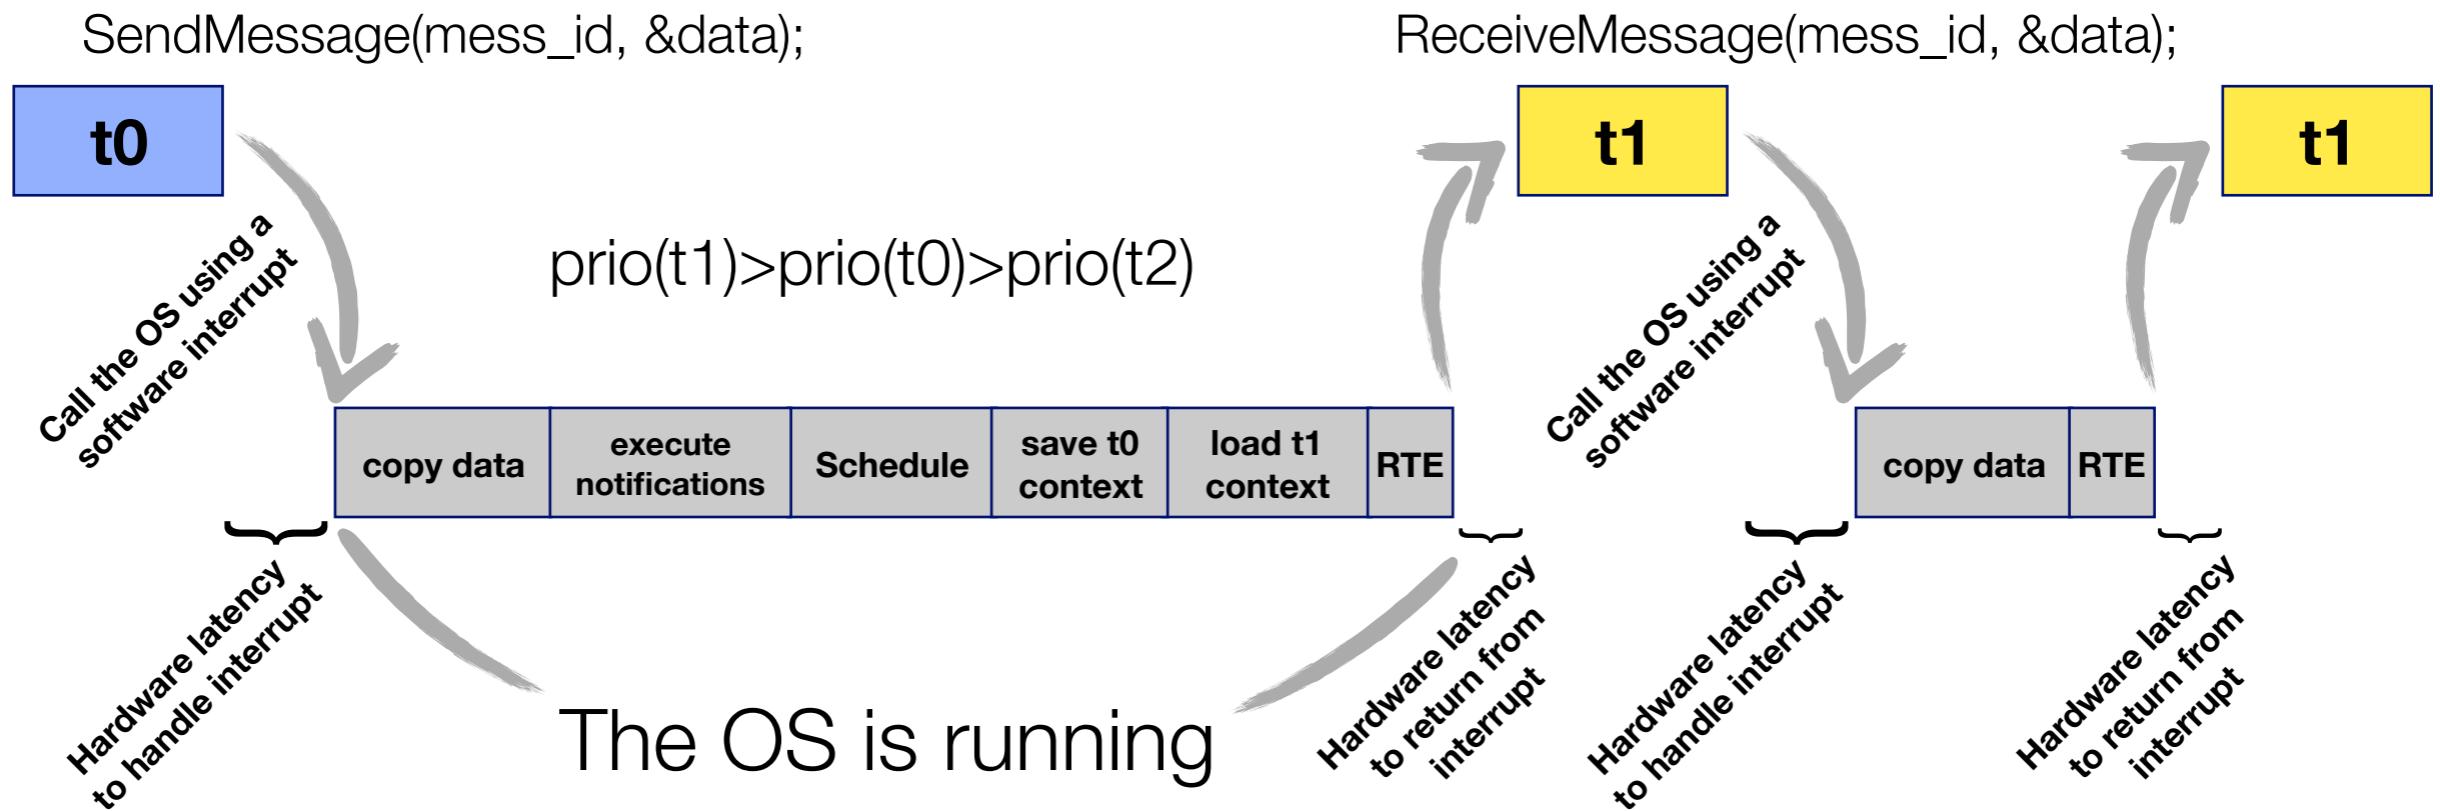
Example: OSEK/VDX messaging aka OSEK/COM

- OSEK/VDX messaging is design for communication within the same micro controller (internal communication) and for communication across a communication network (external communication).
We will see the internal communication.
- Messages objects:
 - Sending message objects (SMO) to send messages
 - Receiving message objects (RMO), linked to sending message objects, to receive messages
- Filters:
 - Filters are associated to RMO and allows to **receive or discard** the message according to rules
- Notifications:
 - Notifications are associated to RMO and allows to **set an event or activate a task**.

Example: OSEK/VDX messaging

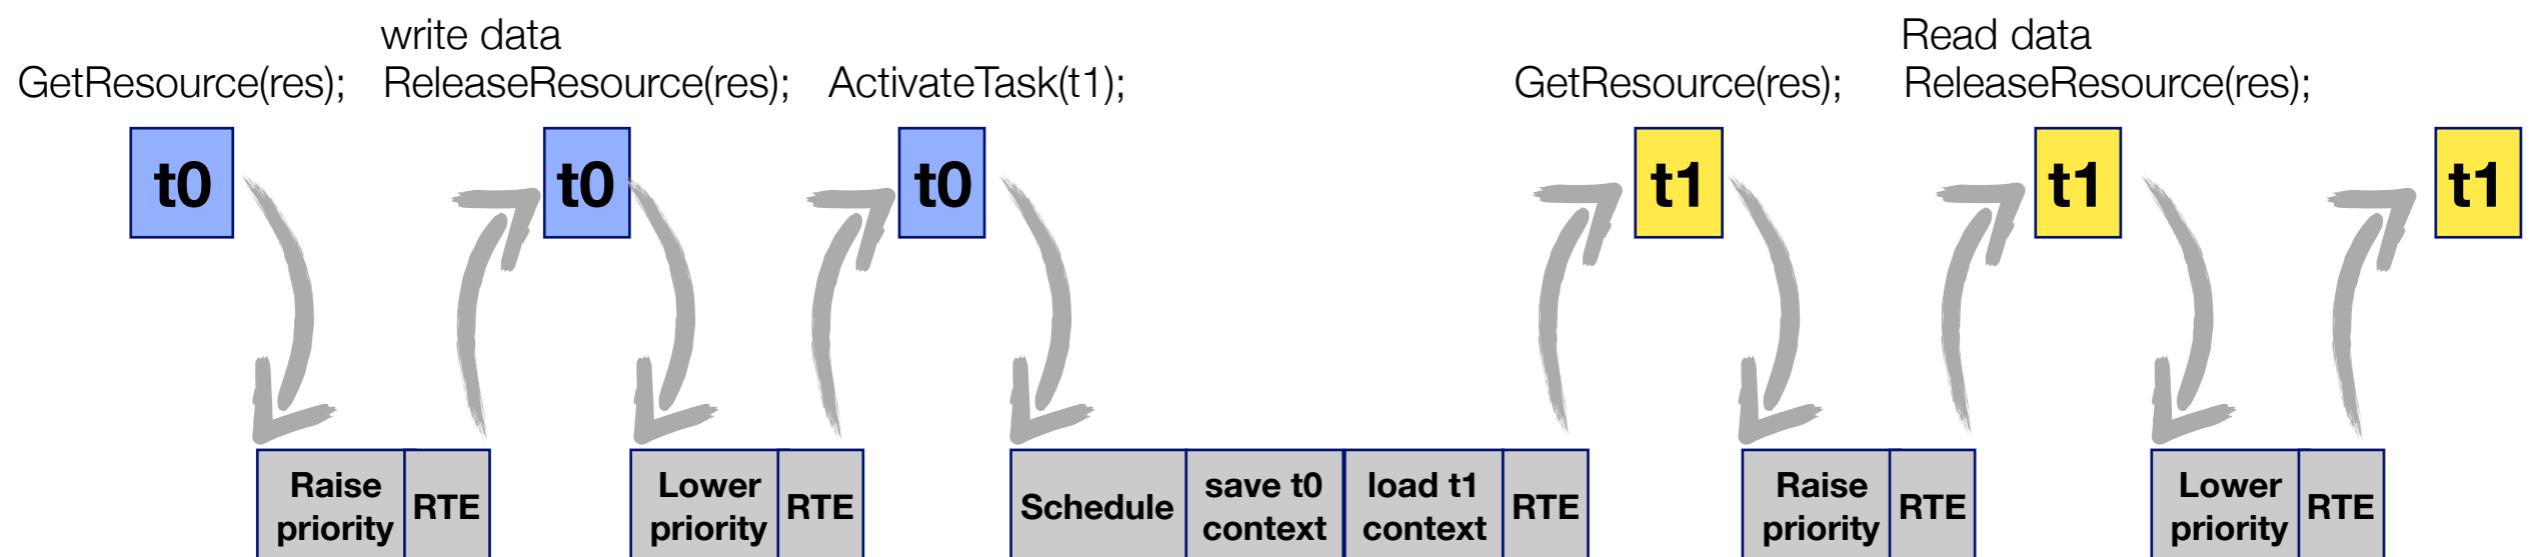


How it works



Comparison with a shared variable

$\text{prio}(t1) > \text{prio}(t0)$



Message filtering (1)

- It is not always useful to receive all the messages...
 - when 2 values in a row are the same
 - when the value is greater than a threshold. The message is used by a different task. For instance a control system where an alternate law is used according to the value
 - when the sender has a period x times smaller than the period of the receiver. In this case 1 message every x is received
- For these cases (and for others) OSEK/COM allows to set a filter for each receiving message.
- The message is delivered if the filter result is TRUE

Message filtering (2)

Filter	Sub-attribute	Result
ALWAYS		TRUE
NEVER		FALSE
MASKEDNEWEQUALSX	MASK, X	new & MASK == X
MASKEDNEWDIFFERSX	MASK, X	new & MASK != X
NEWISEQUAL		new == old
NEWISDIFFERENT		new != old
MASKEDNEWEQUALSMASKEDOLD	MASK	new & MASK == old & MASK
MASKEDNEWDIFFERSMASKEDOLD	MASK	new & MASK != old & MASK
NEWISWITHIN	MIN,MAX	(MIN <= new) && (new <= MAX)
NEWISOUTSIDE	MIN, MAX	(new < MIN) (new > MAX)
NEWISGREATER		new > old
NEWISLESSOREQUAL		new <= old
NEWISLESS		new < old
NEWISGREATEROREQUAL		new >= old
ONEEVERYN	PERIOD, OFFSET	msg_number % PERIOD == OFFSET